

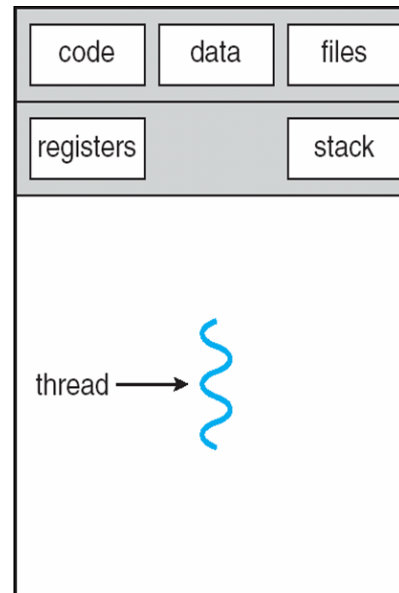
Threads

CS3026 Operating Systems

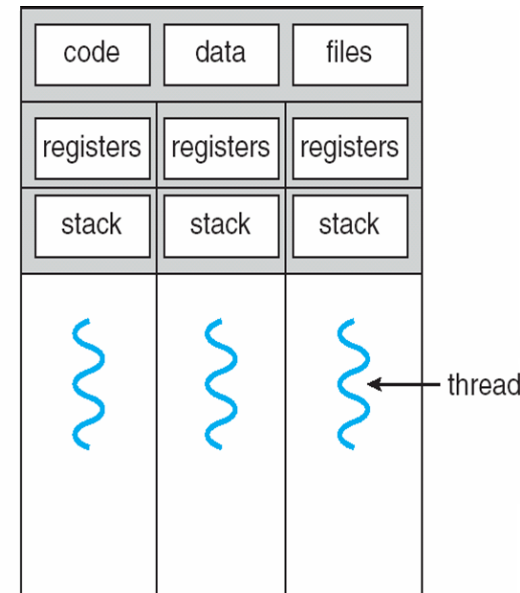
Lecture 06

Multithreading

- Multithreading is the ability of an operating system to support multiple threads of execution within a single process
- Processes have at least one thread of control
 - Is the CPU context, when process is dispatched for execution



single-threaded process



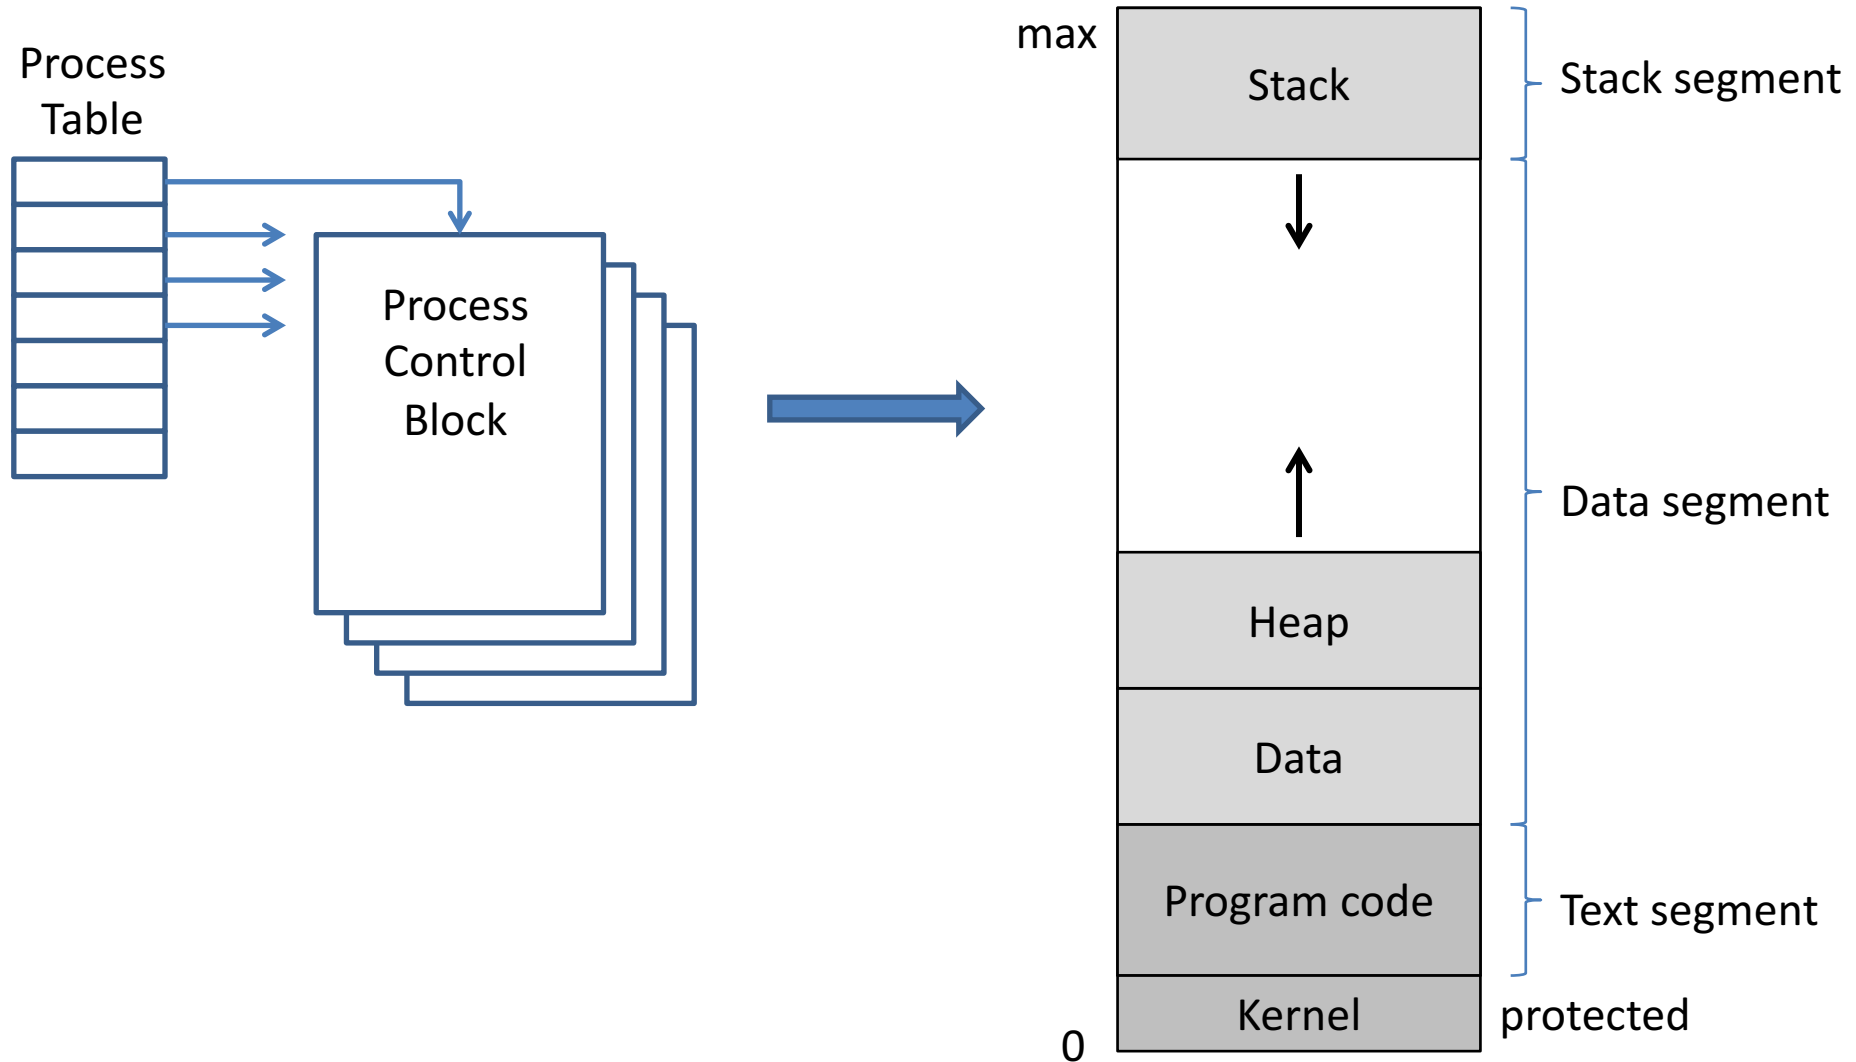
multithreaded process

- Multiple threads run in the same address space, share the same memory areas
 - The creation of a thread only creates a new thread control structure, not a separate process image

Multithreading

- Our process model so far: we defined a process as the **Unit of resource ownership** as well as the **Unit of dispatching**
- We want to separate these two concerns
 - **Resource ownership:**
 - Process remains unit of resource ownership
 - **Program Execution / Dispatching:**
 - A process can have multiple Threads of execution
 - Threads (lightweight processes) become the unit of dispatching
 - CPU may have multiple cores, true parallel execution of threads

Process in Unix




Process – Unit of Resource Ownership

- Unit of **resource ownership** and **protection**
 - *Resource ownership*:
 - Process image, virtual address space
 - Resources (I/O devices, I/O channels, files, main memory)
 - Protection
 - Processors, other processes
 - Operating system protects process to prevent unwanted interference between processes
 - memory, files, I/O resources

Threads – Units of Dispatch

- Thread is defined as the **unit of dispatching**:
 - Represents a single thread of execution within a process
- Threads are also called “lightweight processes”
 - Operating system may be able to manage multiple threads of execution directly (e.g. Linux tasks)
- A thread is provided with its own register context and stack space
- Multiple threads run in the same address space, share the same memory areas
 - The creation of a thread only creates a new thread control structure, not a separate process image

Threads in Unix



- It depends on the actual Kernel implementation how threads are managed
- Are threads implemented at User level only or also at kernel level?

Threads

- All threads share the same address space
 - Share global variables
- All threads share the same open files, child processes, signals, etc.
- There is no protection between threads
 - As they share the same address space they may overwrite each others data
- As a process is owned by one user, all threads are owned by one user

Threads vs Processes: Advantages

- Advantages of Threads
 - Much faster to create a thread than a process
 - Spawning a new thread only involves allocating a new stack and a new thread control block
 - 10-times faster than process creation in Unix
 - Less time to terminate a thread
 - Much faster to switch between threads than to switch between processes
 - Threads share data easily
 - Thread communication very efficient, no need to call kernel routines, as all threads live in same process context

Threads vs Processes: Disadvantages

- Disadvantages
 - Processes are more flexible
 - They don't have to run on the same processor
 - No protection between threads
 - Share same memory, may interfere with each other
 - If threads are implemented as user threads instead of kernel threads
 - If one thread blocks, all threads in process block


Thread Management

Thread Management

- Threads are described by the following:
 - Thread execution state
 - running, ready, blocked
 - Thread Control Block
 - A saved thread context when not running (each thread has a separate program counter)
 - An execution stack
 - Some per-thread static storage for local variables
 - Access to memory and resources of its process, shared with all other threads of that process

Thread States

- Threads have now also three basic states
 - Running: CPU executes thread
 - Ready: thread control block is placed in Ready queue
 - Blocked: thread awaits event
- If one thread blocks
 - Is the whole process with all other threads blocked?
 - Or is only this single thread blocked?



Thread Operations

- There are four basic operations for managing threads
 - Spawn / create
 - A thread is created and provided with its own register context and stack space, it can spawn further threads
 - Block:
 - if a thread waits for an event, it will block
 - If the kernel manages threads: the processor may switch to another thread in the same or a different process
 - Unblock:
 - When the event occurs, for which the thread is waiting, it will be queued for execution
 - Finish:
 - When a thread completes, its register context and stacks are de-allocated

User and Kernel Threads

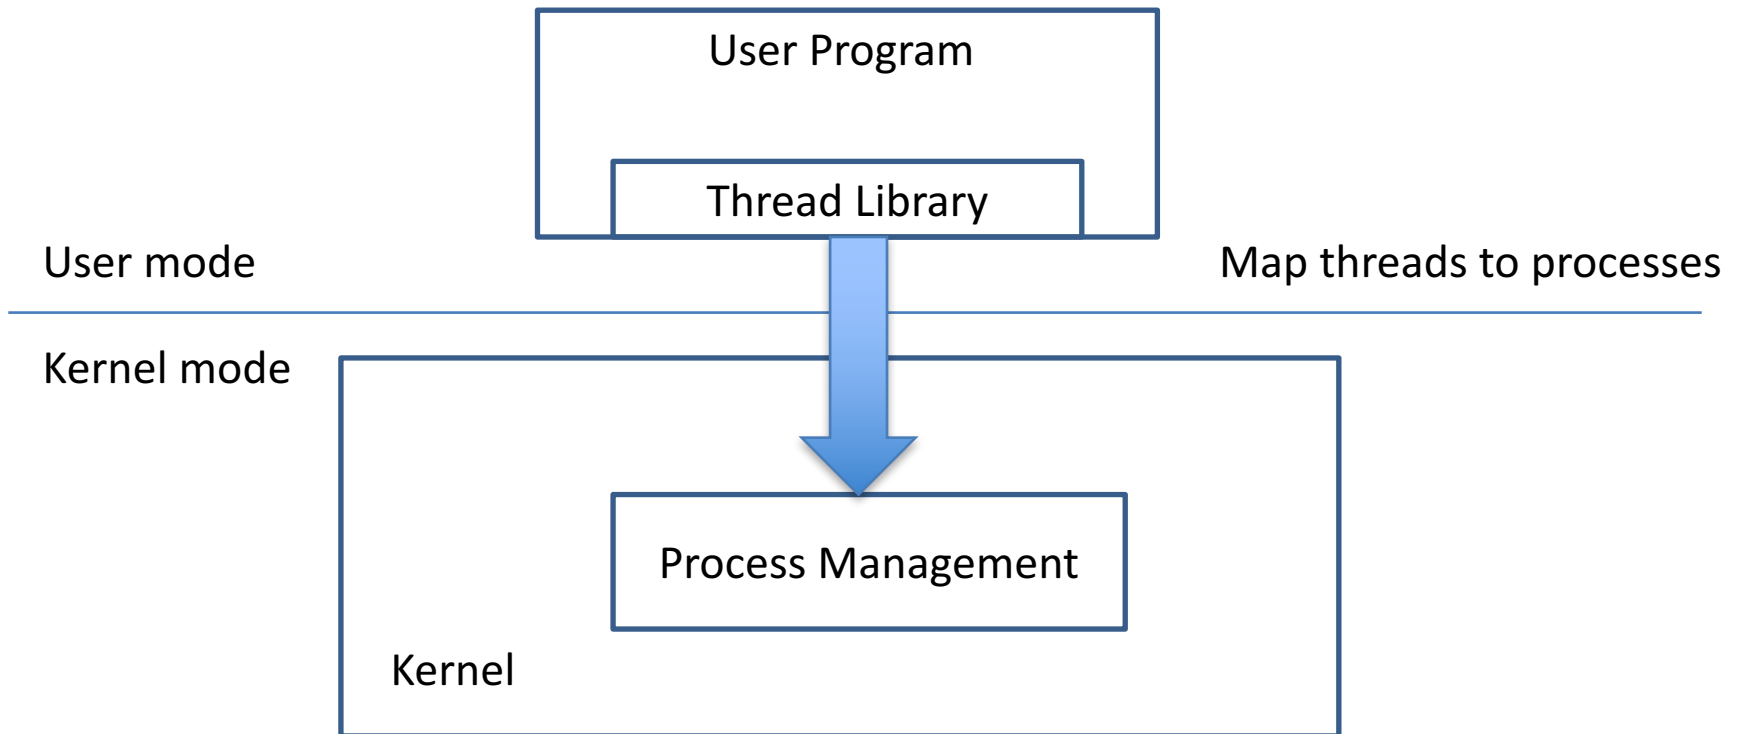
Thread Management

- If one thread blocks
 - Is the whole process with all other threads blocked?
 - Or is only this single thread blocked?
- This depends on the implementation of the kernel
 - The kernel may only know processes, threads have to be managed by the program itself
 - The kernel knows threads, threads in a program can be directly managed by kernel

Thread Implementation


- Three categories of threads
 - User-level threads
 - Kernel-level threads
 - Mixed user-kernel-level threads
- Characterised by the extent of the kernel being involved in their management

User Threads



- Operating System unaware of threads, completely controlled by User program


Kernel Threads




- Operating System manages threads, better mapping of threads to multiple CPU cores

Thread Implementation


- Pure User Level Threads
- Pure Kernel Level Threads
- Combined User / Kernel level threads



Pure User-Level
ULT



Pure Kernel-Level
KLT




CombinedLevel
ULT/KLT

User-Level Threads

- User-Level Threads
 - Kernel not aware of the existence of threads
 - Process uses thread library functions to manage its threads
- Benefit
 - Light thread switching in user mode
 - No mode switch necessary (no call of kernel functions)
 - We can implement our own thread scheduling
- Also called “green threads” on some systems (e.g. Solaris)

User-Level Threads: Disadvantage

- Process is still the **Unit of Dispatch**, not a thread:
 - Kernel doesn't know threads
- Disadvantage:
 - Blocking of one thread blocks entire process, including all other threads in it
 - Only one thread can access the kernel at a time, as the process is the unit of execution known by kernel
 - No Distribution in Multi-processor systems:
 - All threads run on the same processor in a multiprocessor system
 - Threads cannot run in parallel utilising different processors, as the process is dispatched on one processor




Kernel-level Threads

- Thread is **Unit of dispatch**
 - Kernel is aware of the existence of threads
 - Kernel manages each thread separately
- Benefit
 - Fine-grain scheduling by kernel on thread basis
 - If a thread blocks (e.g. waiting for I/O), another one can be scheduled by kernel without blocking the whole process
 - Threads can be distributed to multiple processors and run in parallel
- Example Systems: Windows XP/7/8, Solaris, Linux, Mac OSX

Kernel-Level Threads: Disadvantage

- A context switch between threads of the same process involves kernel
 - Mode switch: we have to switch CPU into Kernel mode
- Mode switch is costly
 - 2 mode switches for each thread context switch, is as costly as process switch




Hybrid Implementations

User-Kernel-Level Threads

- Try to combine advantages of both user-level and kernel-level threads
 - User-level: light-weight thread switching
 - Kernel-level: allows dispatch at thread level (same or different process), when one thread blocks
 - true parallelism of threads in multiprocessor systems possible
- Basic technique: Mapping of user-level threads onto a limited set of kernel threads
- Different hybrid Multithreading Models:
 - Many-to-one
 - One-to-one
 - Many-to-many


Many-to-One Model

- All user-level threads of one process mapped to a single kernel-level thread / process
- These are the User-level threads as discussed before
- Thread management in user space
 - Efficient
 - Application can run its own scheduler implementation
- One thread can access the kernel at a time
 - Limited concurrency, limited parallelism
- Examples
 - “Green threads” (e.g. Solaris)
 - Gnu Portable Threads



One-to-One Model


- Each user-level thread mapped to a kernel thread
- These are the Kernel-level threads as discussed before
- One blocking thread does not block other threads
- Multiple threads access kernel concurrently



- Problem
 - Kernel may restrict the number of threads created
- Example systems
 - Windows, Linux, Solaris 9 (and later), Mac OSX


Many-to-Many Model

- Many user-level threads are multiplexed (mapped dynamically) to a smaller or equal number of kernel threads
 - Thread pool, no fixed binding between a user and a kernel thread
- The number of kernel threads is specific to a particular application or computer system
 - Application may be allocated more kernel threads on a multiprocessor architecture as on a single-processor architecture
- No restriction on user-level threads
 - Applications can be designed with as many user-level threads as needed
 - Threads are then mapped dynamically onto a smaller set of currently available kernel threads for execution



Two-level Model

- Is a variant of the Many-to-Many model, allows a fixed relationship between a user thread and a kernel thread
- Was used in older Unix-like systems
 - IRIX, HP-UX, True64 Unix, Solaris 8



Threading Issues – Thread Pools

- Threads come with some overhead
- Unlimited thread creation may exhaust memory and CPU
- Solution
 - Thread pool: create a number of threads at system startup and put them in a pool, from where they will be allocated
 - When an application needs to spawn a thread, an allocated thread is taken from the pool and adapted to the application's needs
- Advantage
 - Usually faster to service a request with already instantiated thread then creating a new one
 - Allows number of threads in applications to be bound by thread pool size
- Number of pre-allocated threads in pool may depend on
 - Number of CPUs, memory size
 - Expected number of concurrent requests

Threading Issues – fork() and exec()

- Semantics of fork() and exec() changes in a multithreaded program
 - Remember:
 - fork() creates an identical copy of the calling process
 - In case of a multithreaded program
 - Should the new process duplicate all threads?
 - Or should the new process be created with only one thread?
 - If after fork(), the new process calls exec() to start a new program within the created process image, only one thread may be sufficient
 - Solution: some Unix systems implement two versions of fork()

Threading Issues – Signal Handling

- Signals are used in Unix systems to notify a process about the occurrence of an event
 - E.g.: CTRL-C terminates a program
- All signals follow the same pattern
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Once delivered, a signal must be handled
- In multithreaded systems, there are 4 options
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process

Thread Programming

Thread Programming

- POSIX standard threads: pthreads
- Describes an API for creating and managing threads
- There is at least one thread that is created by executing `main()`
- Other threads are spawned / created from this initial thread

POSIX Thread Programming

- Thread creation

```
pthread_create ( thread, attr, start_routine, arg )
```

- Returns a new thread ID with parameter “thread”
- Executes the routine specified by “start_routine” with argument specified by “arg”

- Thread termination

```
pthread_exit ( status )
```

- Terminates the thread, sends “status” to any thread waiting by calling pthread_join()

POSIX Thread Programming

- Thread synchronisation

```
pthread_join ( threadid, status)
```

- Blocks the calling thread until the thread specified by “threadid” terminates
- The argument “status” passes on the return status of pthread_exit(), called by the thread specified by “threadid”

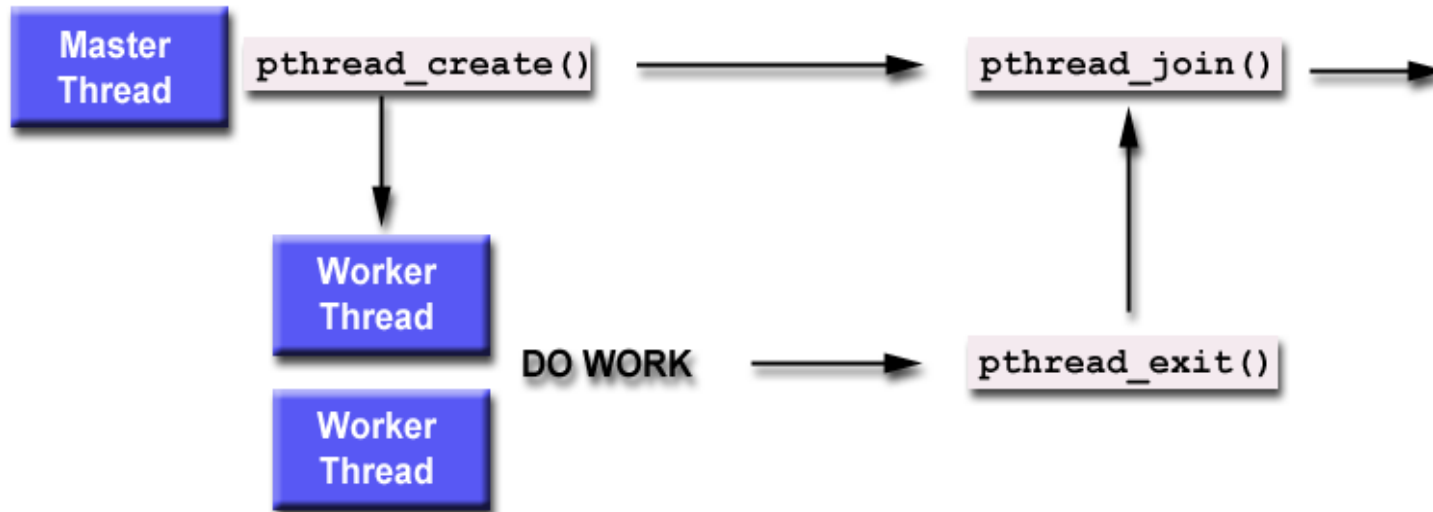
- Thread yield

```
pthread_yield ( )
```

- Calling thread gives up the CPU and enters the Ready queue

Thread Programming

Wait for completion with join()



SMP Support – Processor Affinity

- Processor Affinity: Relationship between a thread and a processor
 - No affinity
 - Threads of any process may run on any processor
 - No guarantee that they are rescheduled on the same processor after interruption
 - Soft affinity:
 - Dispatcher tries to re-assign threads to the same processor after interruption
 - Helps to reuse data still held on processor cache, less cleanup and reload necessary
 - Hard affinity
 - Application restricts execution of a thread to a particular processor / processor core

Linux Threads

- No distinction between processes and threads
- User-level threads are mapped into kernel-level processes (called “tasks”)
- A new process is created by copying the attributes of the current process
- A clone() system call exists
 - Cloned processes share resources, but have separate stacks

Mac OSX Grand Central Dispatch

- Thread pool
- Mac OSX can operate even more fine-grained:
 - Blocks of program code are **Units of dispatch**
 - A software designer can designate portions of program code, called blocks, that may be dispatched independently and run concurrently
 - These are extensions to a programming language
- Concurrency depends on the number of CPU cores available and the thread capacity of a system

Summary

- User-level threads
 - Created and managed by a thread library that runs in the user space of a process
 - No mode switch required to switch from one thread to another
 - Only a single user-level thread within a process can execute at a time, no execution in parallel on a multiprocessor system
 - If one thread blocks, the entire process is blocked
- Kernel-level threads
 - Threads within a process are maintained by the kernel
 - A mode switch is required to switch from one thread to another
 - Multiple threads within the same process can execute in parallel on a multiprocessor system
 - Blocking of one thread does not block the entire process
- Process: resource ownership
- Thread: program execution / dispatching