

Memory Management

CS3026 Operating Systems

Lecture 08

Memory Management

- Functional Requirements:
 - Relocation
 - Process images may be positioned at arbitrary locations in memory and may be relocated
 - Important for Virtual memory management
 - Partitioning
 - Create memory partitions and allocate them to processes
 - Security and Isolation
 - Protect segments of memory, isolate memory areas of processes
 - Needs hardware support
 - Sharing
 - Allow shared access to memory by different processes

Memory Management

- Non-Functional Requirements
 - Performance
 - Minimal overhead of memory management
 - Fast allocation
 - Avoid thrashing
 - Fairness
 - Avoid starvation of processes
 - Tune Working set according to process needs

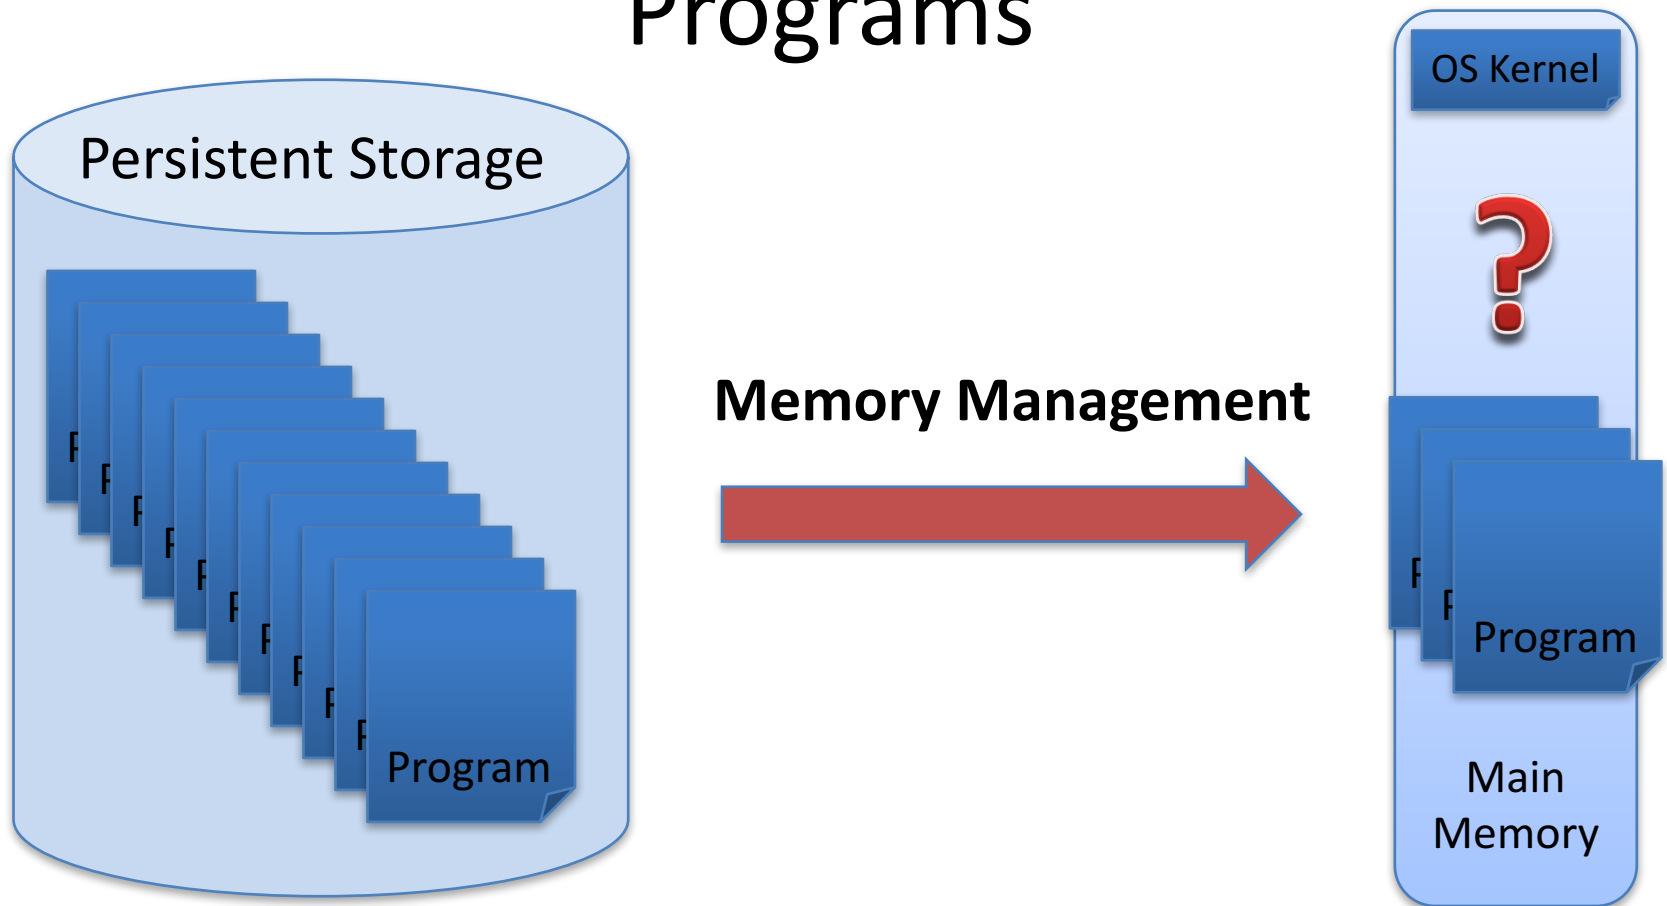
Meeting the Requirements

- Modern systems use virtual memory
 - Supported by hardware and software
 - Programs are not restricted in size by actual physical memory
 - Number of processes executing on system is not limited by physical memory
 - Based on paging and segmentation

Meeting the Requirements

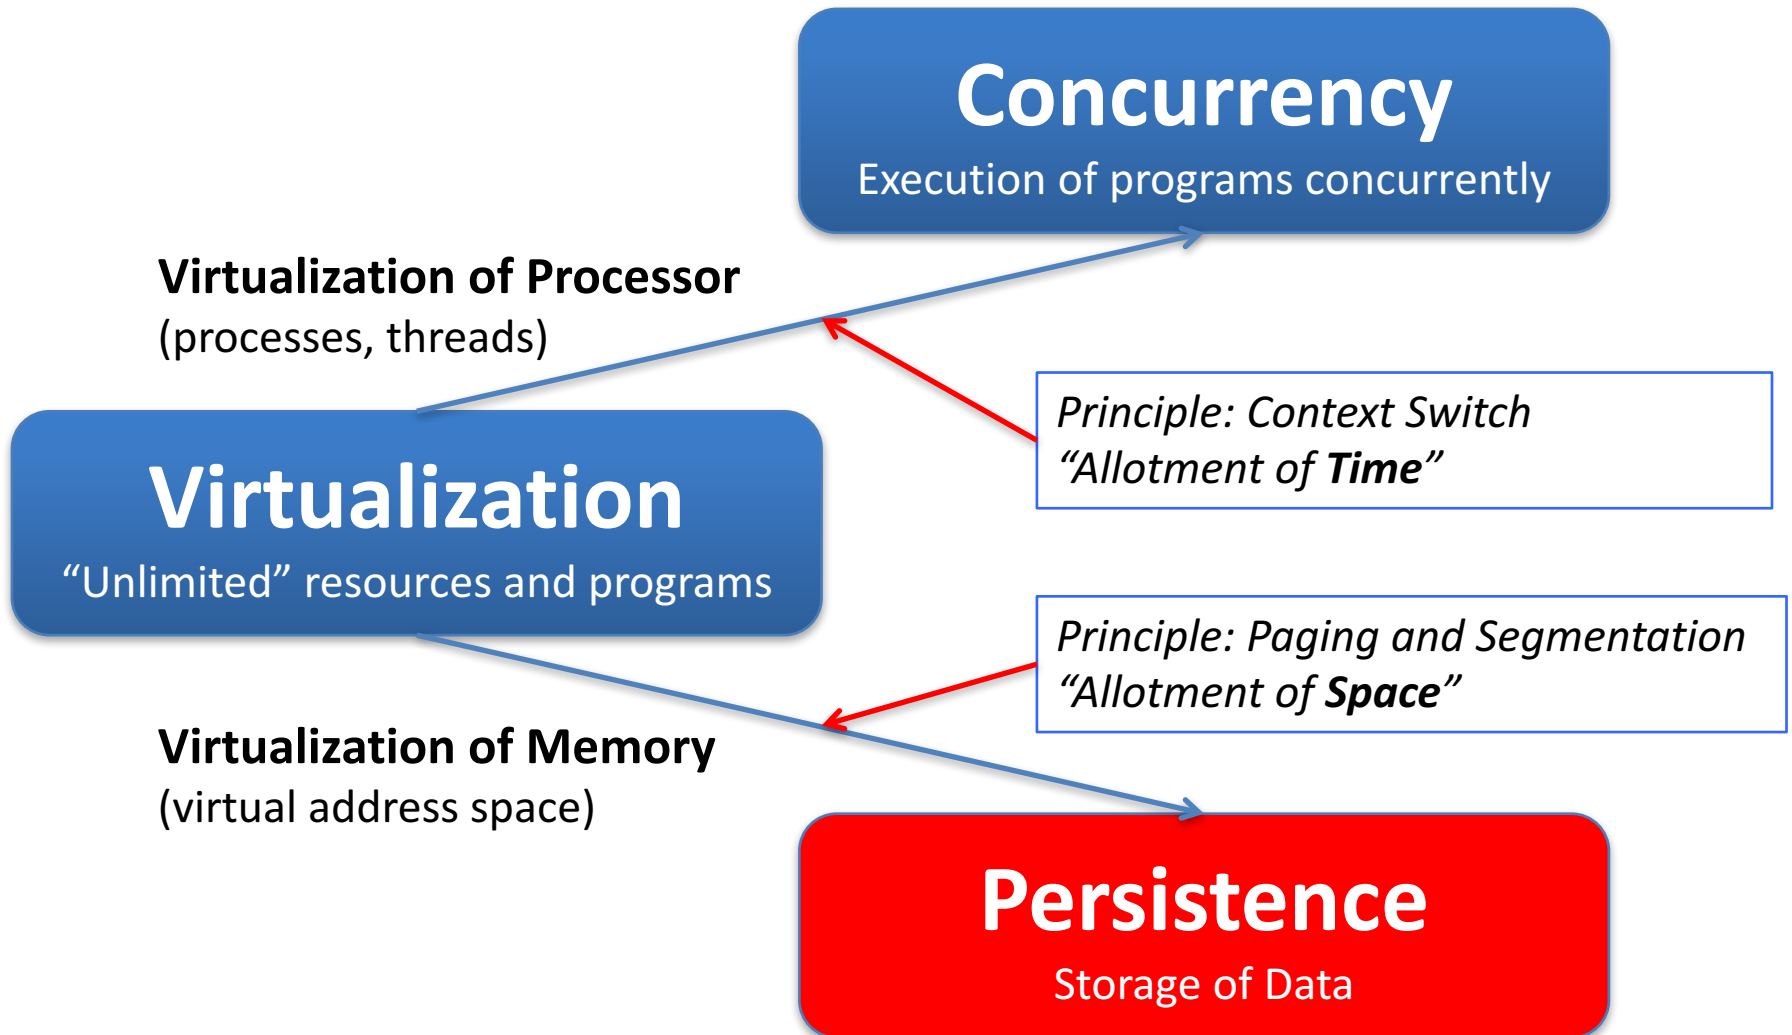
- Logical Organisation
 - User programs operate within a virtual address space
 - User programs have a modular structure, each process is divided into segments (code, data, stack), mechanisms for protecting segments (read-only, execute-only) and sharing among processes
- Physical Organisation
 - Virtual memory based on a two-level hierarchy between physical memory and disk space
 - CPU can only access data in registers and physical memory
- Address Translation
 - MMU Memory Management Unit translates logical address into physical address

Multitasking – Concurrent Execution of Programs



- We need a memory management strategy to allow concurrent execution of programs

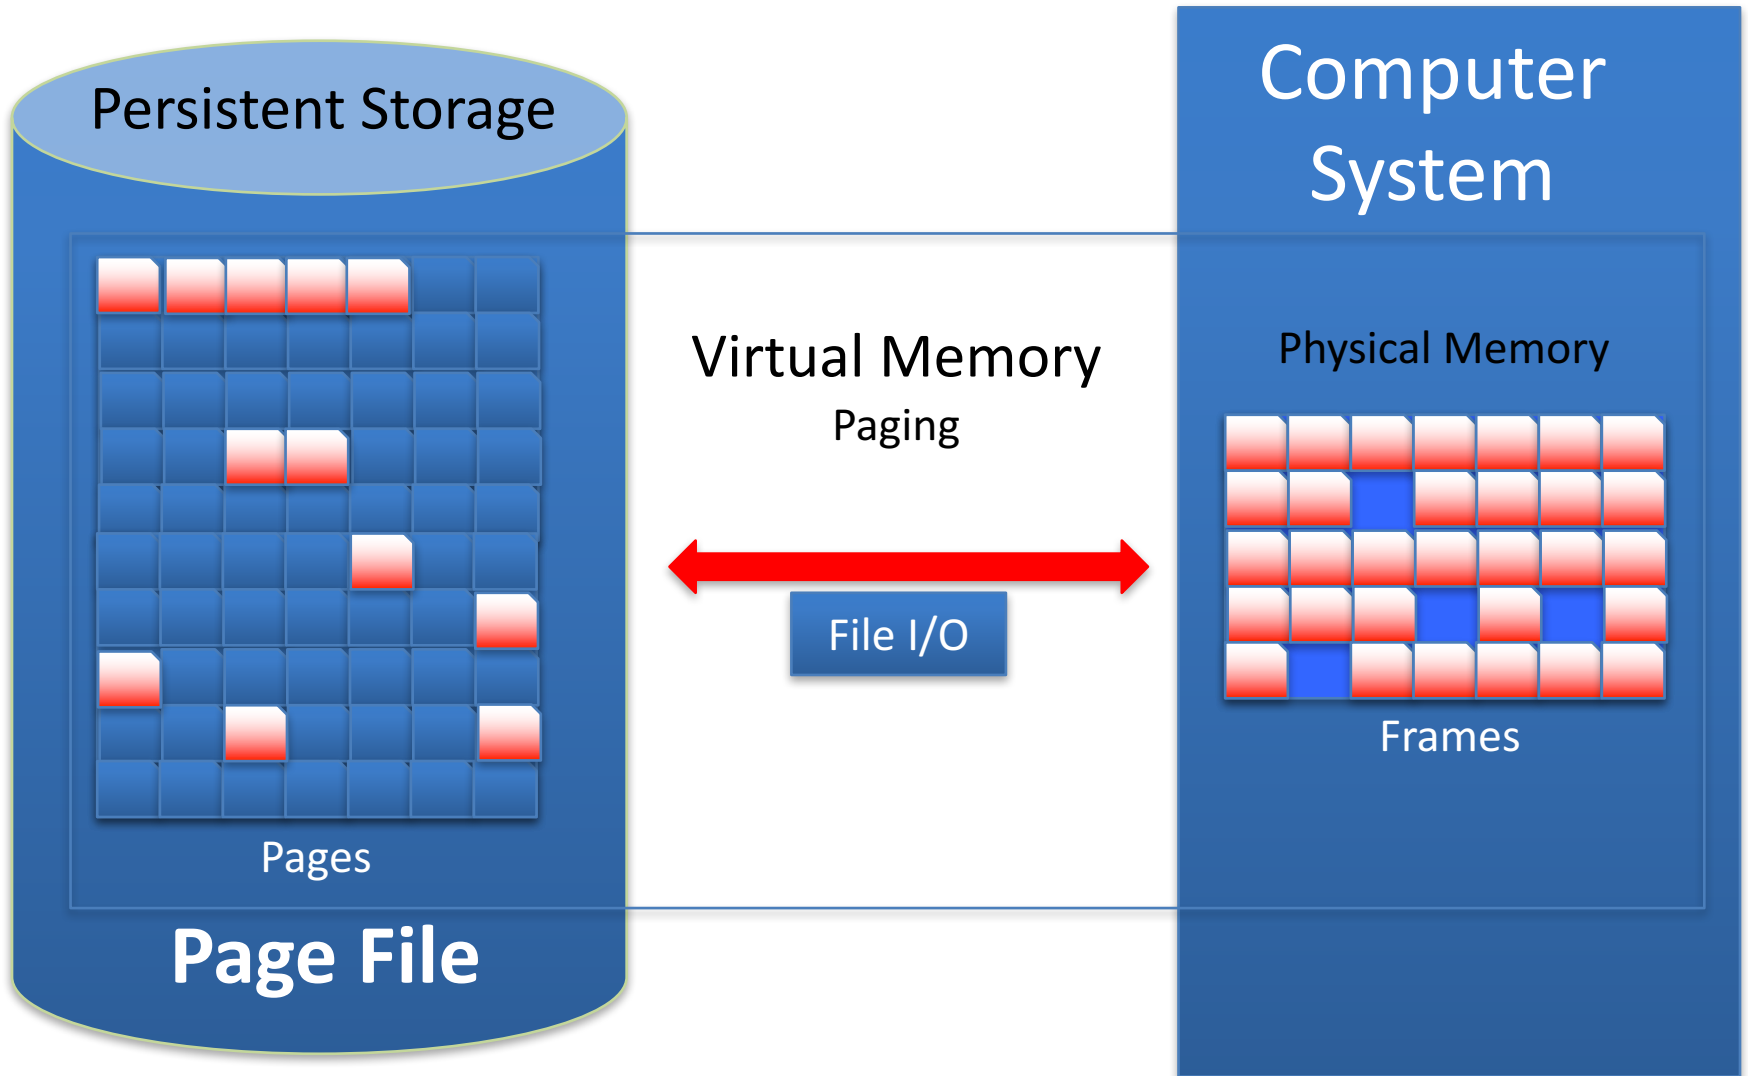
Core Concepts



Virtual Memory Management

Virtualization

Memory Management

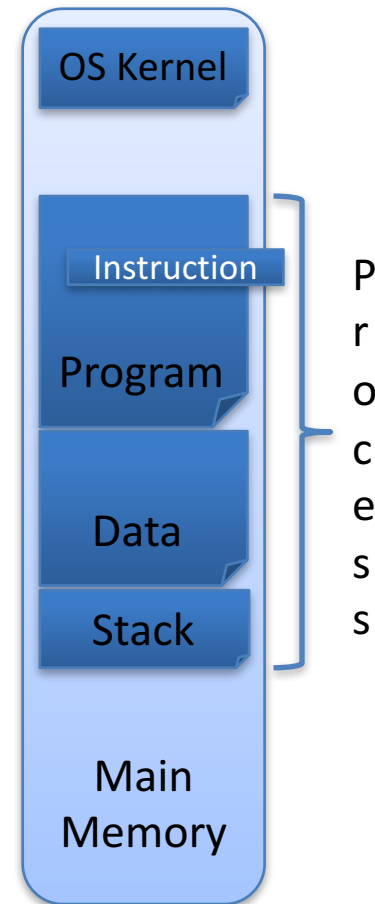


Paging

- Organise both virtual address space of process image and physical memory into fixed-sized blocks:
 - **Frames**: Physical memory is divided into equal fixed-sized blocks, called “frames” or “page frames”
 - **Pages**: Virtual address space of a process is divided into equal fixed-sized blocks called “pages”
 - Any page can be assigned to any free frame
- The operating system allocates a small amount of frames to a process
- Page table
 - Maps pages to frames

Virtual Memory Management

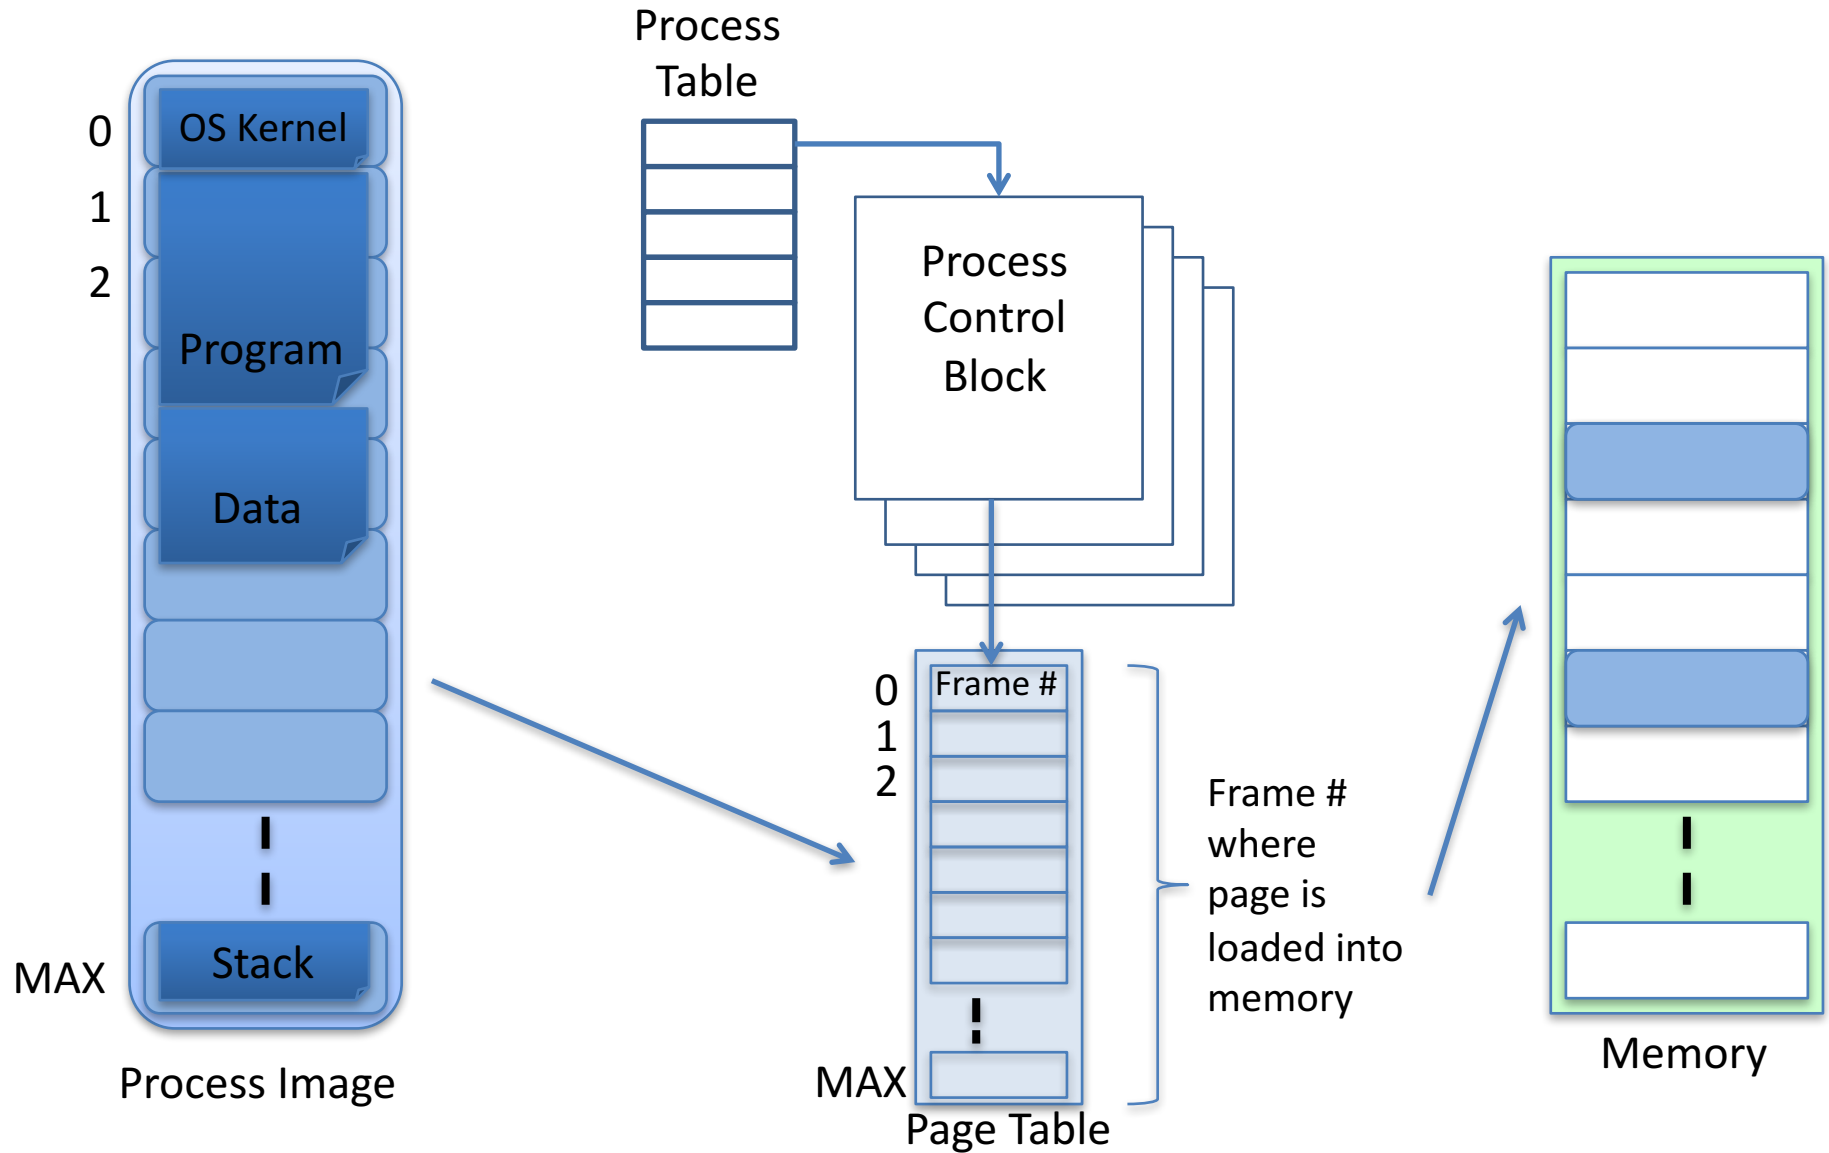
- Non-contiguous allocation of memory for processes
 - A process may be broken up into a number of pieces – pages that don't need to occupy a contiguous area of main memory
 - Segments may start at different physical locations
- Processes only partially loaded into physical memory
 - Not all pages of a process image have to be in physical memory at the same time
 - Process is only allowed to load a limited set of pages – “working set”
- All memory references are logical addresses that are dynamically translated into physical addresses at run time



Paging and Virtual Memory

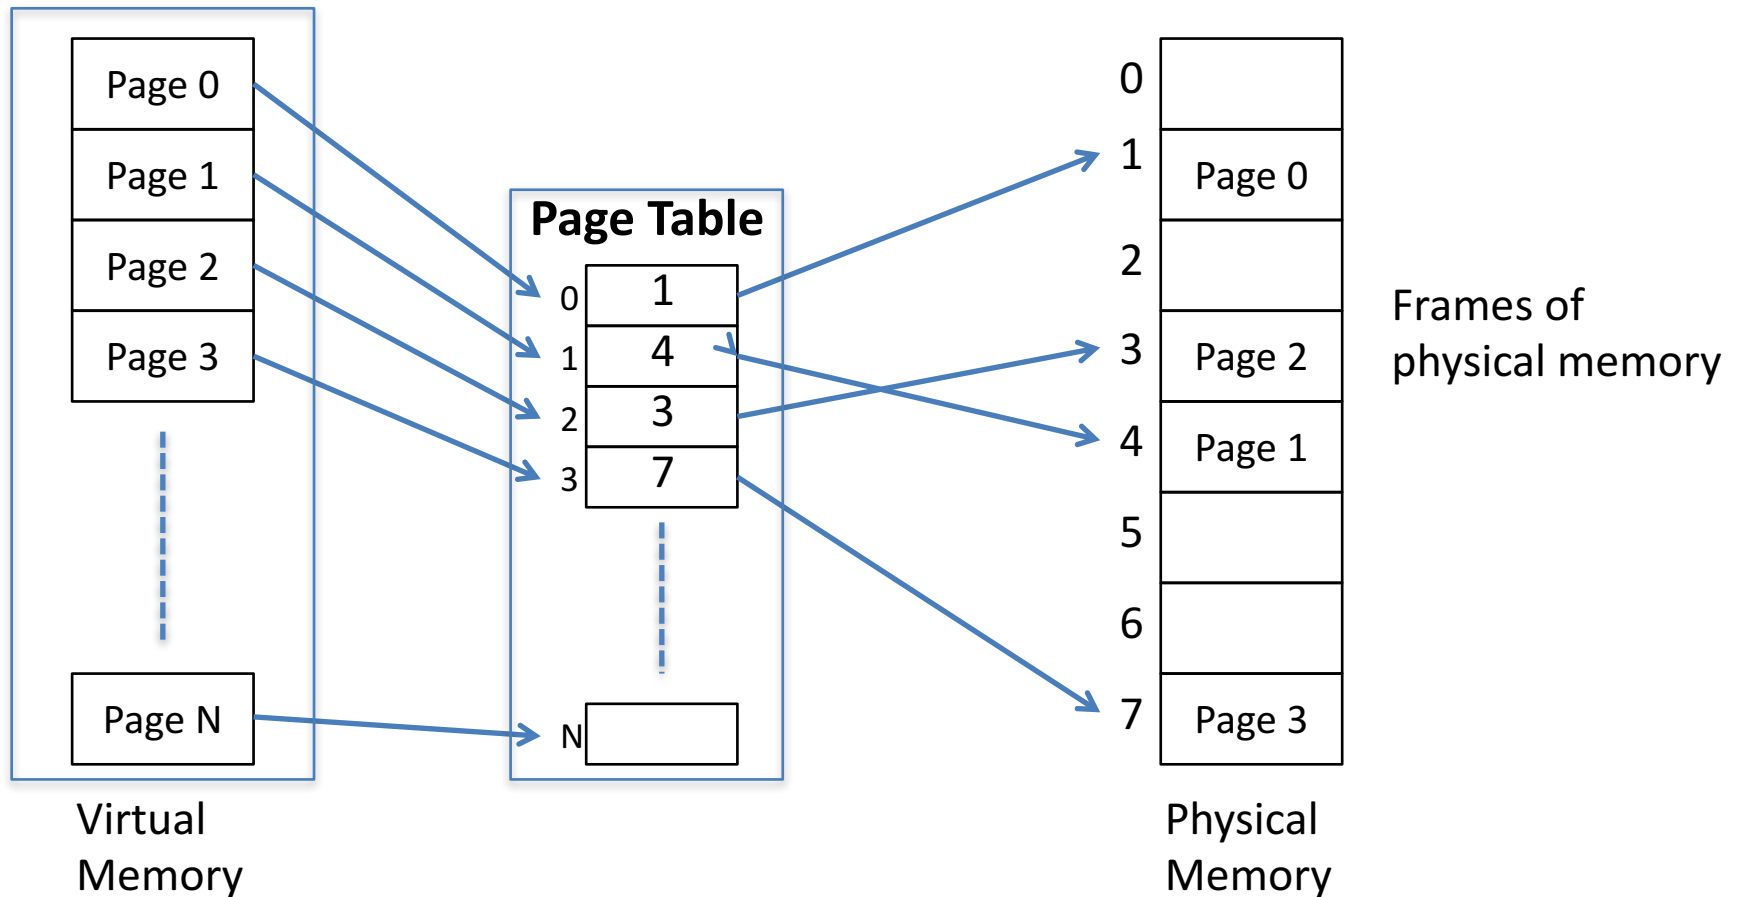
- Helps to implement virtual memory
 - Page file held on hard disk, is also organised in blocks of the same size as a page
 - Pages of a process can easily be swapped in and out of memory
 - We can implement a process image only being partially loaded, pages are loaded on demand
- No time-consuming search for free best-fit memory fragments
- No external memory fragmentation
- Paging is supported by modern hardware

Page Table of a process



Paging

- Allocation of pages to frames in physical memory can be arbitrary and non-contiguous



Page Table Management

- Operating system manages one page table per process
- A pointer to the page table is stored in the process control block
- Context switch (starts a new process)
 - the stored page table of the process has to be referenced correctly
 - We can load a “page table pointer” into one of the processor registers and use it for address translation

Paging: Address Translation

From logical to physical address

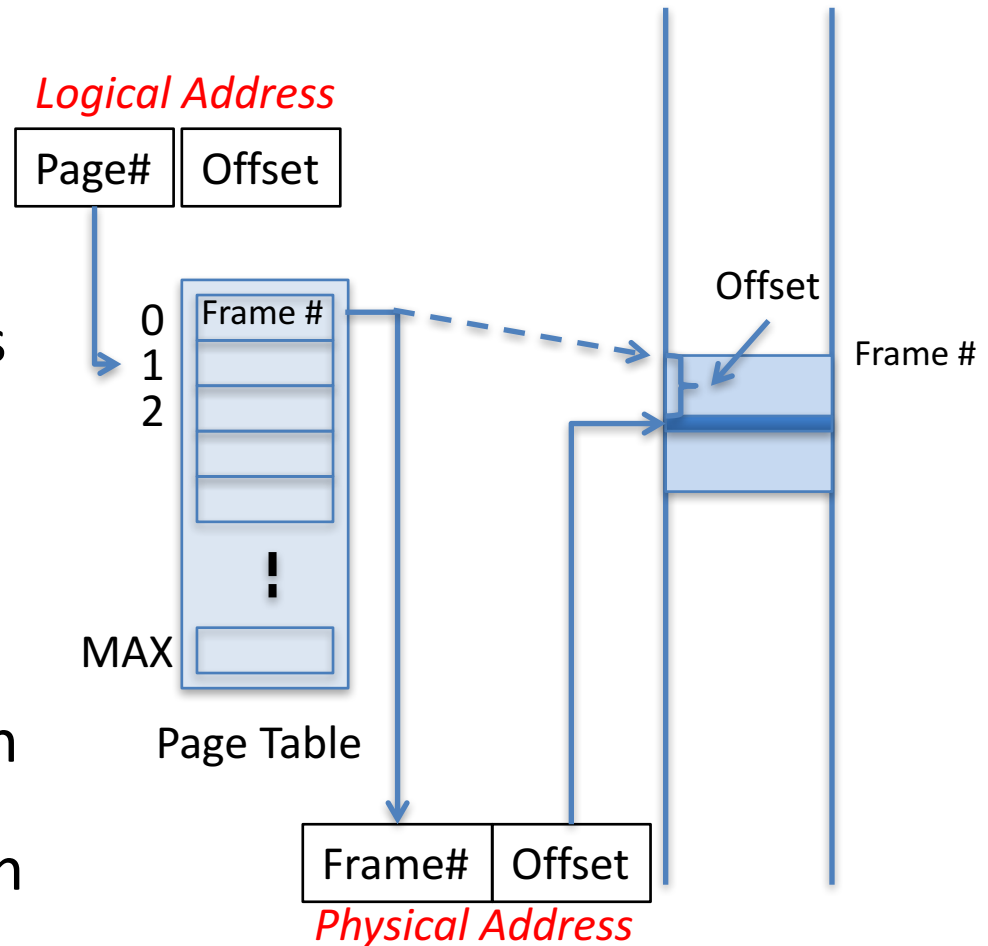
Addressing Virtual Memory

- Pages are typically of sizes between 512 bytes and 16MB
- Consider a 32-bit logical address
- Consider 4kB per page:
 - we need 12 bits ($2^{12} = 4096$) to address a location within the page
 - 20 bits for the page number: we can have $2^{20} = 1$ Mio pages
 - **Page table must have 1 Mio entries !!**
- Address space: 4GB of virtual memory

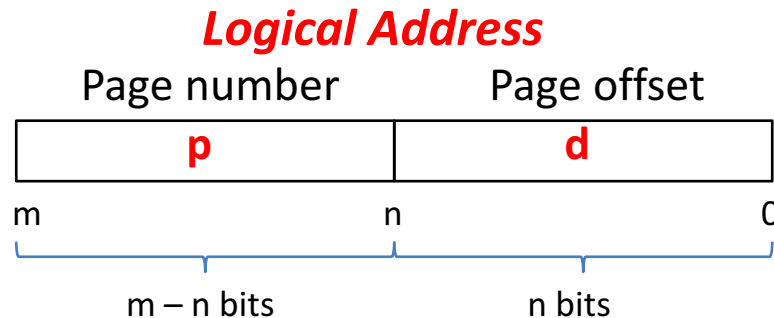
Bytes	Exponent				
1,024	2^{10}	1kb	1024bytes		
1,048,576	2^{20}	1MB	1024kb		1024 x 1024
1,073,741,824	2^{30}	1GB	1024MB		1024 x 1024 x 1024
4,294,967,296	2^{32}	4GB	4 x 1024MB		4 x 1024 x 1024 x 1024
1,099,511,627,776	2^{40}	1TB	1024GB		1024 x 1024 x 1024 x 1024
17,592,186,044,416	2^{44}	16TB	16 x 1024GB		16 x 1024 x 1024 x 1024 x 1024
1,125,899,906,842,620	2^{50}	1PB	1024TB		1024 x 1024 x 1024 x 1024 x 1024
1,152,921,504,606,850,000	2^{60}	1EB	1024PB		1024 x 1024 x 1024 x 1024 x 1024 x 1024
18,446,744,073,709,600,000	2^{64}	16EB			16 x 1024 x 1024 x 1024 x 1024 x 1024 x 1024

Paging: Address Translation

- A program in execution uses a logical address for addressing memory locations
- This logical address points to a particular location in virtual memory
 - Location within a virtual memory **page**
- Therefore, logical address has to contain information about the virtual memory **page** and the **offset** within this page



Paging: Address Translation



Logical address space: 2^m

Size of a page: 2^n

Number of pages: 2^{m-n}

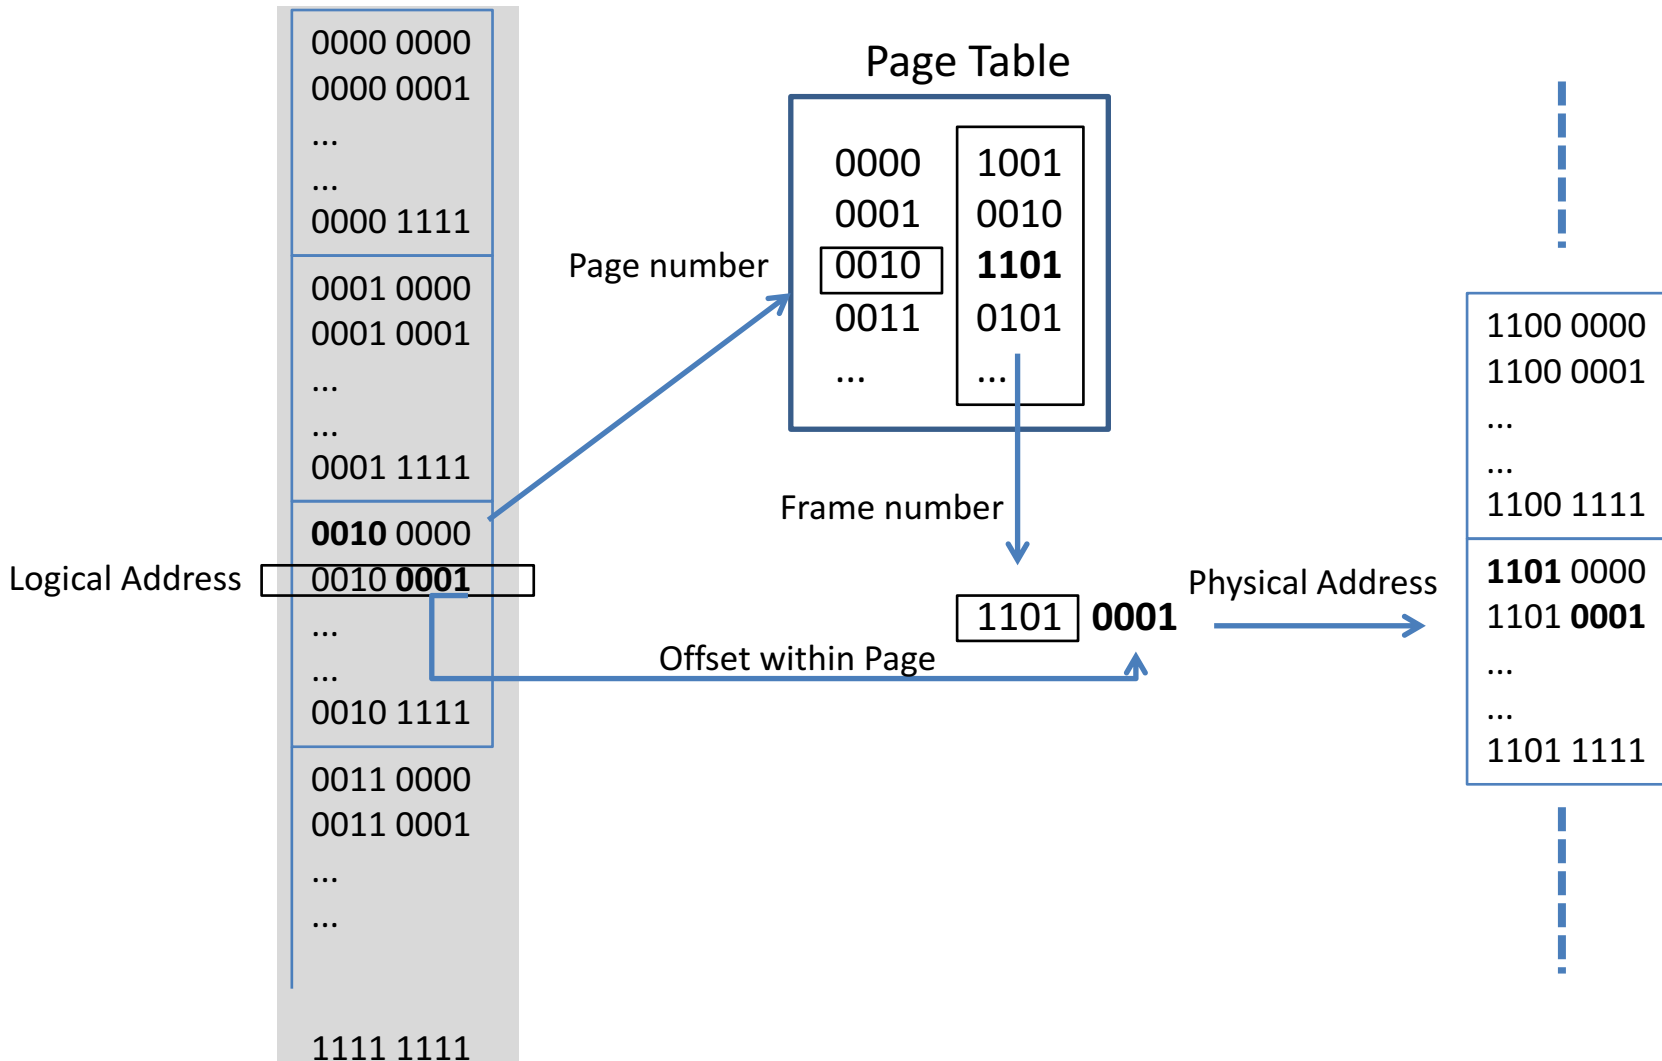
Size of page table

- Logical addresses have two parts
 - Page number (**p**)
 - Used as an index into the page table
 - An entry in this table contains the base address where a particular page is located in physical memory
 - Page offset (**d**)
 - Combined with the base address, defines the physical memory address within the space in memory occupied by the page
 - The high-order $m-n$ bits encode the page number (index into page table), lower n bits encode the page offset

Paging: Address Translation

Virtual Address Space

Physical Memory



Paging: Address Translation

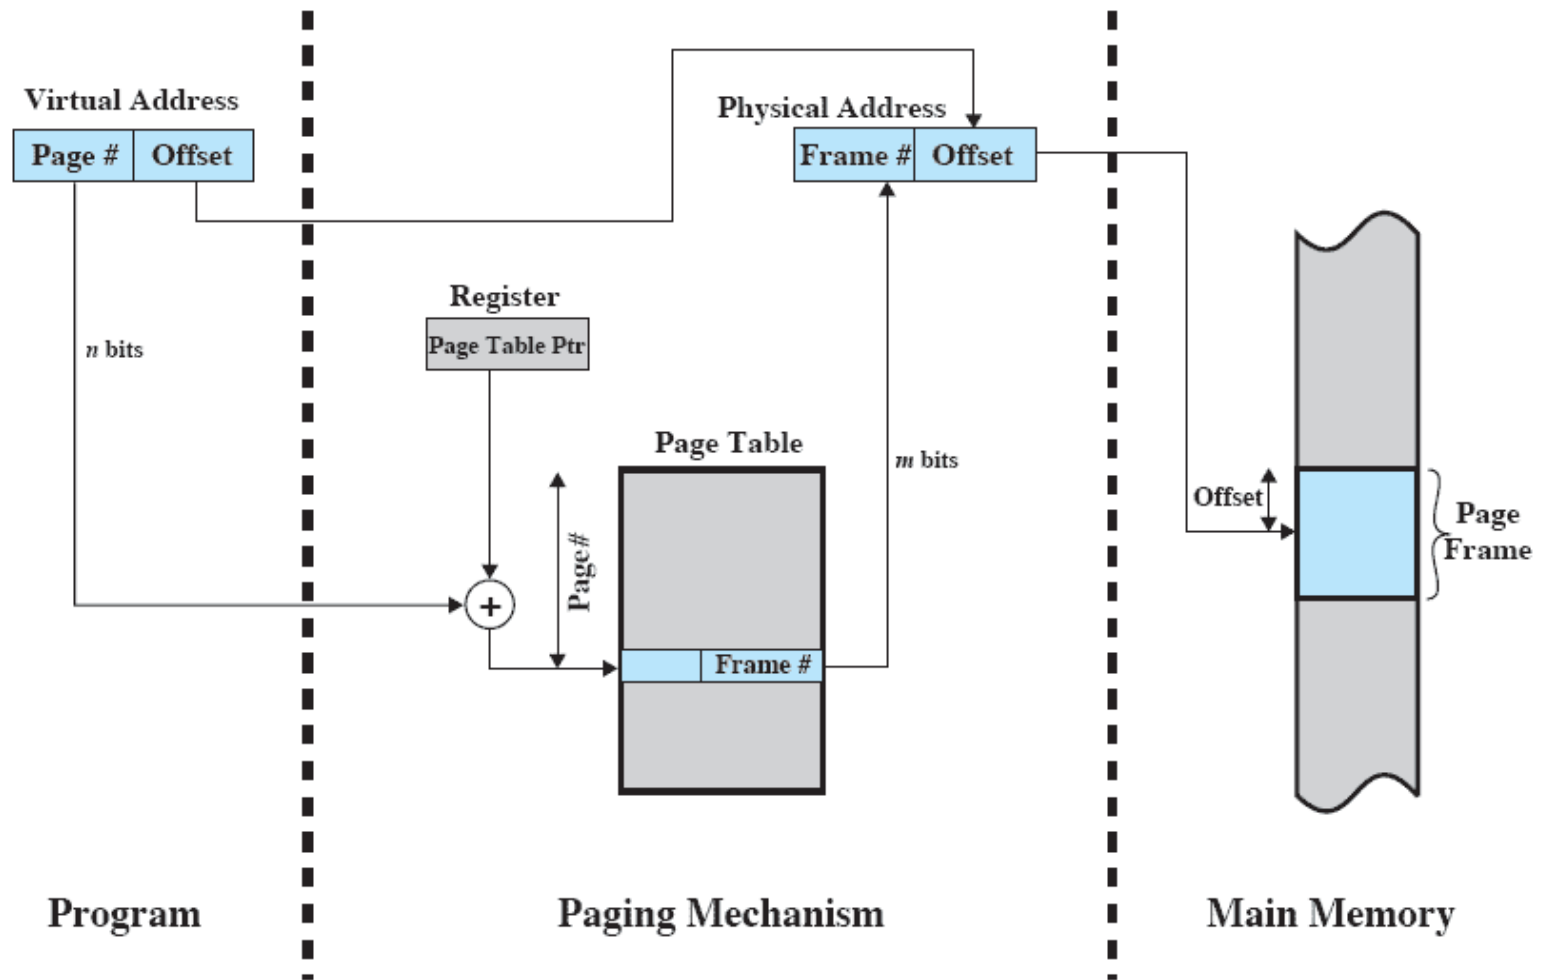
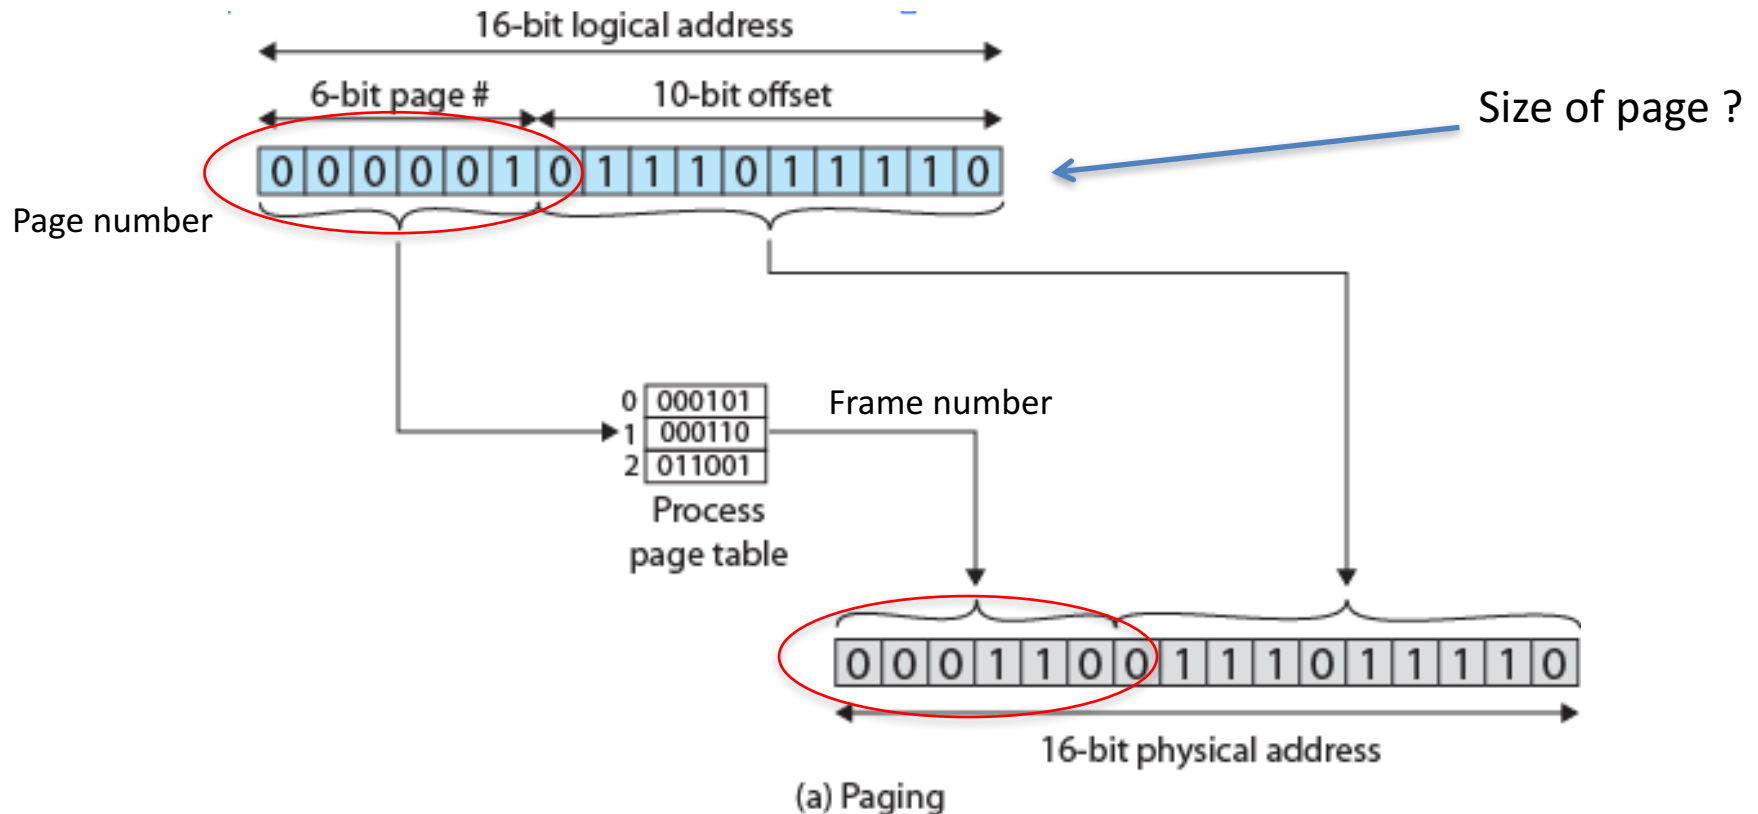


Figure 8.3 Address Translation in a Paging System


Address Translation



- The physical address of memory location is generated by retrieving the physical base address of a page from the page table (index into page table contained in logical address) and combining it with the offset

Example


- Let's assume:
 - **Virtual memory consists of 32 pages**
 - A page is of size 2kB
 - Actual physical memory only consists of 8 frames
- How many bits needed for **logical address**?
 - Note: 1kB = 1024 bytes = 2^{10} bytes, 10 bits needed for addressing each byte within a page of size 1kB
- Addressing virtual memory (logical address space):
 - How many bits?
 - Page size: 2kB = 2048 bytes = 2×1024 bytes = $2 \times 2^{10} = 2^{11}$ bytes
 - We need 11 bits for the offset within a page
 - We have $32 = 2^5$ pages
 - We need 5 bits for the index in the page table (page number)



Logical Address: 16 bit

Example

- Let's assume:
 - Virtual memory consists of 32 pages
 - A page is of size 2kB
 - **Actual physical memory only consists of 8 frames**
- How many bits needed for **physical address**?
 - Note: 1kB = 1024 bytes = 2^{10} bytes, 10 bits needed for addressing each byte within a page of size 1kB
- Addressing physical memory:
 - How many bits for the physical address?
 - We need 11 bits for the offset within a frame (same size as page)
 - We have $8 = 2^3$ frames
 - We need 3 bits for addressing a frame (content of page table)



Physical Address: 14 bit

Addressing Virtual Memory

- Pages are typically of sizes between 512 bytes and 16MB
- Consider a 32-bit logical address
- Consider 4kB per page:
 - we need 12 bits ($2^{12} = 4096$) to address a location within the page
 - 20 bits for the page number: we can have $2^{20} = 1$ Mio pages
 - **Page table must have 1 Mio entries !!**
- Address space: 4GB of virtual memory

Bytes	Exponent				
1,024	2^{10}	1kb	1024bytes		
1,048,576	2^{20}	1MB	1024kb		1024 x 1024
1,073,741,824	2^{30}	1GB	1024MB		1024 x 1024 x 1024
4,294,967,296	2^{32}	4GB	4 x 1024MB		4 x 1024 x 1024 x 1024
1,099,511,627,776	2^{40}	1TB	1024GB		1024 x 1024 x 1024 x 1024
17,592,186,044,416	2^{44}	16TB	16 x 1024GB		16 x 1024 x 1024 x 1024 x 1024
1,125,899,906,842,620	2^{50}	1PB	1024TB		1024 x 1024 x 1024 x 1024 x 1024
1,152,921,504,606,850,000	2^{60}	1EB	1024PB		1024 x 1024 x 1024 x 1024 x 1024 x 1024
18,446,744,073,709,600,000	2^{64}	16EB			16 x 1024 x 1024 x 1024 x 1024 x 1024 x 1024

Fragmentation

- There is no external fragmentation
 - Any free frame can be allocated to a process that needs it
 - This is due to the fixed page size the correspondence between page size and frame size
- There is internal fragmentation
 - As page / frame sizes are fixed, there may be some waste
 - The last frame allocated to a process may not be needed in its completeness
 - Worst case: process may need n frames plus one additional byte! A complete additional frame has to be allocated just for this additional byte of required memory
 - Expected internal fragmentation: on average one-half page per process

Control Bits

- Page table entries are annotated with control bits:
 - Present or valid bit: indicates whether a page is loaded into memory or not
 - Modify bit: indicates whether a page has been modified since it was loaded
 - If not, no need to write it out when it is replaced
 - Other bits may indicate read/write protection, whether a page is shared etc.

Frame #	valid-invalid bit
	v
	v
	v
	v
	i
....	
	i
	i

page table

Resident Set

- Is the number of pages currently loaded into physical memory
- Is restricted by the number of frames the operating system allocates to a process
- Allocation problem:
 - How many frames per process ??

Demand Paging

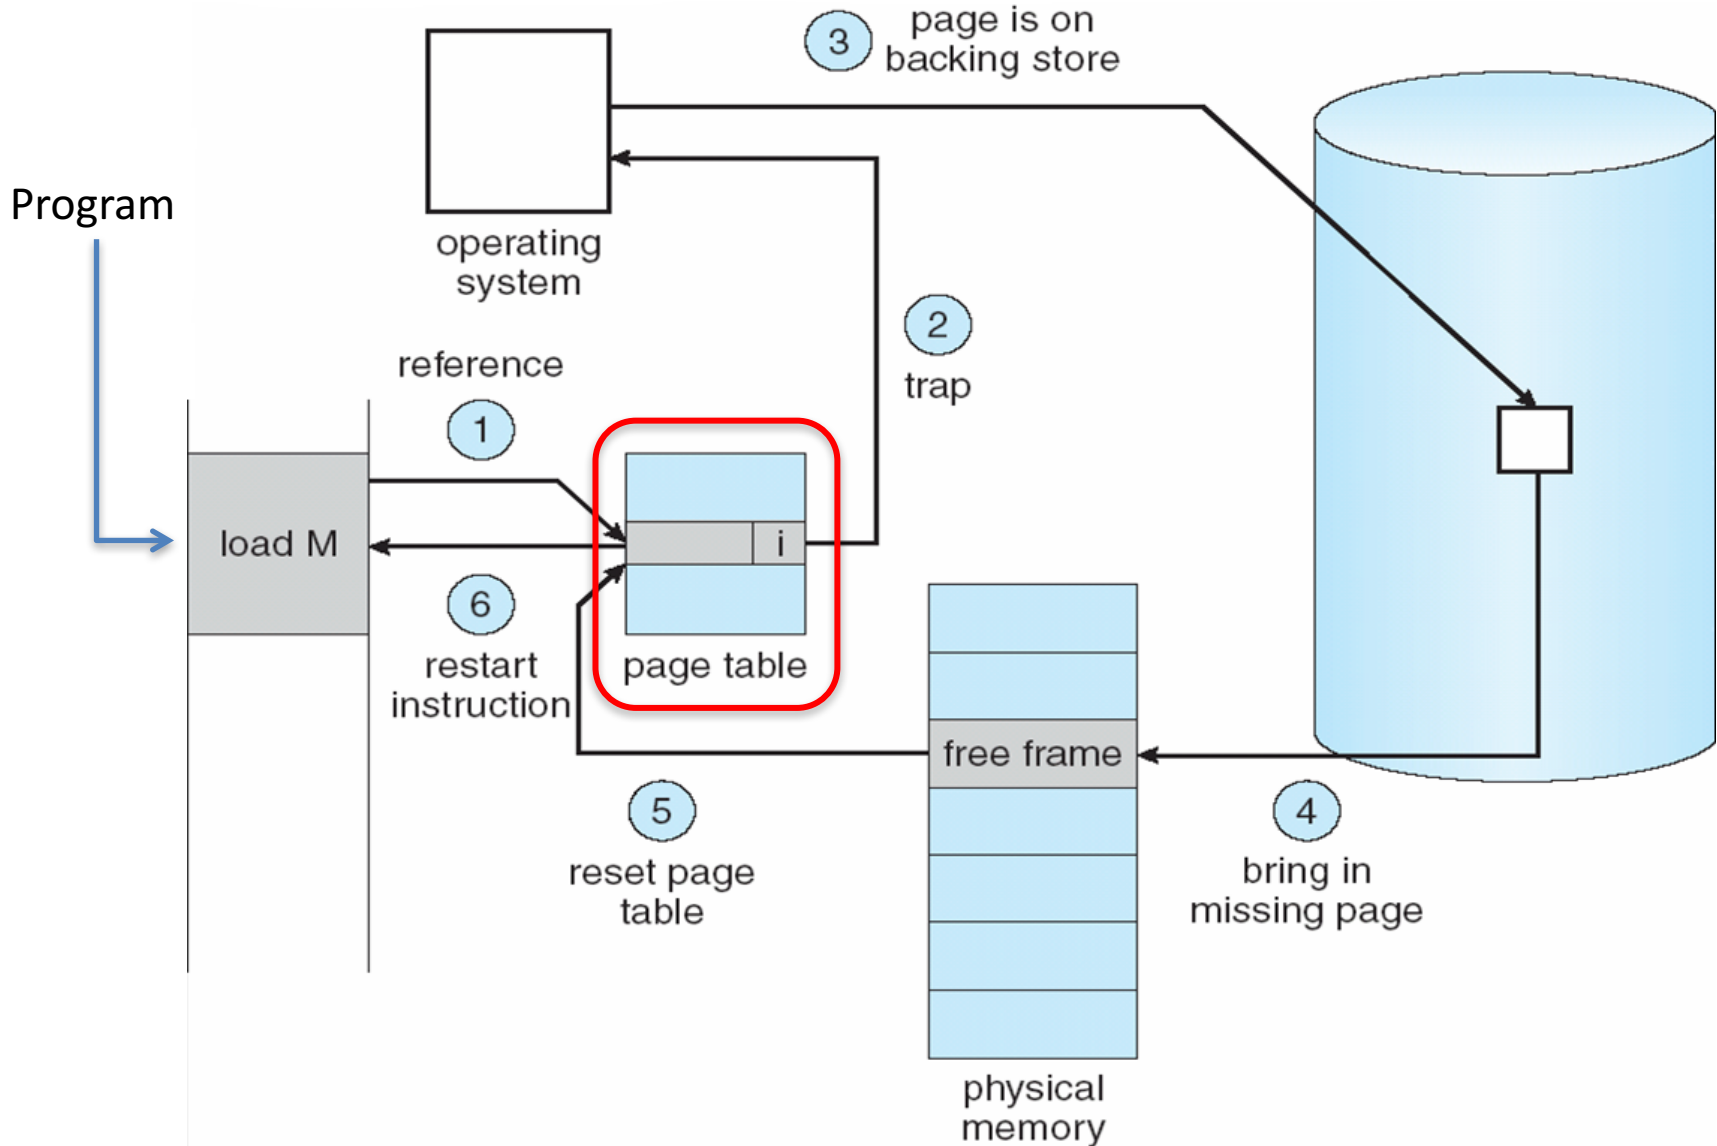
Demand Paging

- Virtualisation of memory management
 - Not all pages of a process must be in physical memory
 - A large part of the process image can be held on secondary storage
 - A page load mechanism is needed to load pages on demand
 - Only those pages that are actually needed, are loaded into memory
 - When program accesses memory, always check first whether page is in memory
 - If page is not loaded – “page fault”, leads to process interruption and I/O operations

Page Fault and Demand Paging

- If processor encounters a logical address outside the resident set of the currently executing process:
 - Process is interrupted, put in blocked state
 - I/O request is issued to read in that part of the process that contains referenced memory location
 - Other processes are scheduled for execution during this I/O activity (context switch)
 - I/O interrupt routine indicates end of read action to load needed parts of blocked process into physical memory
 - Process is put back into ready queue, waits for being rescheduled

Page Fault Handling



Principle of Locality

- Despite the elaborate and time-consuming procedure to handle page faults, virtual memory management is efficient
- Principle of Locality

“References to memory locations within a program tend to cluster”

- if there is an access to a memory location, the next access is, in all likelihood, very near to the previous one
- Loading one page may satisfy subsequent memory access, as all required memory locations may be contained within this page, no loading of page necessary

Principle of Locality

- We can observe that only a few pieces of a process image will be needed during short time periods
 - This is the current memory “locality” of a process
 - It is therefore possible to make informed guesses about which pieces will be needed in the near future
- This principle indicates that virtual memory management can be managed efficiently

Segmentation

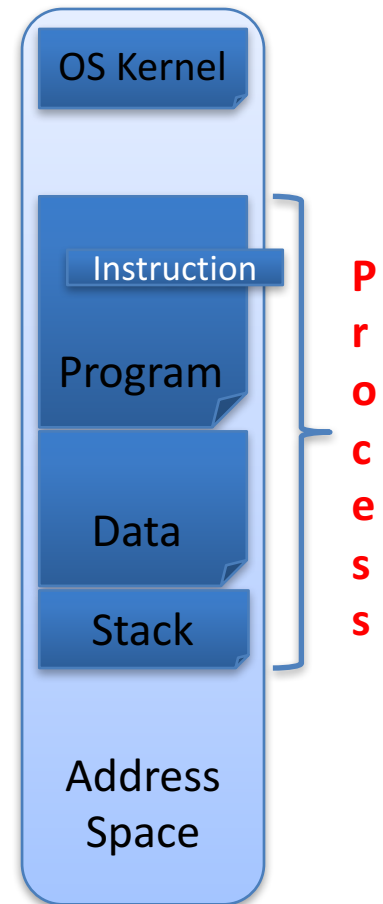
Segmentation

- **Each segment is a separate address space**
 - Has a base address where the segment starts, memory within a segment is addressed relative to this base address
 - Allows parts of a program code to be altered and recompiled independently
 - Shared libraries
- Segments allow Programs to be divided into logical units of variable length
 - Process image regarded as a collection of segments
 - Code segment
 - Data segment
 - Stack, Heap, shared libraries etc.
- Segments are visible to programmer
- In programs, a “virtual address” to a memory location consists of two parts: starting location of the segment and offset within the segment

Virtual Address


Segment *i*

Offset



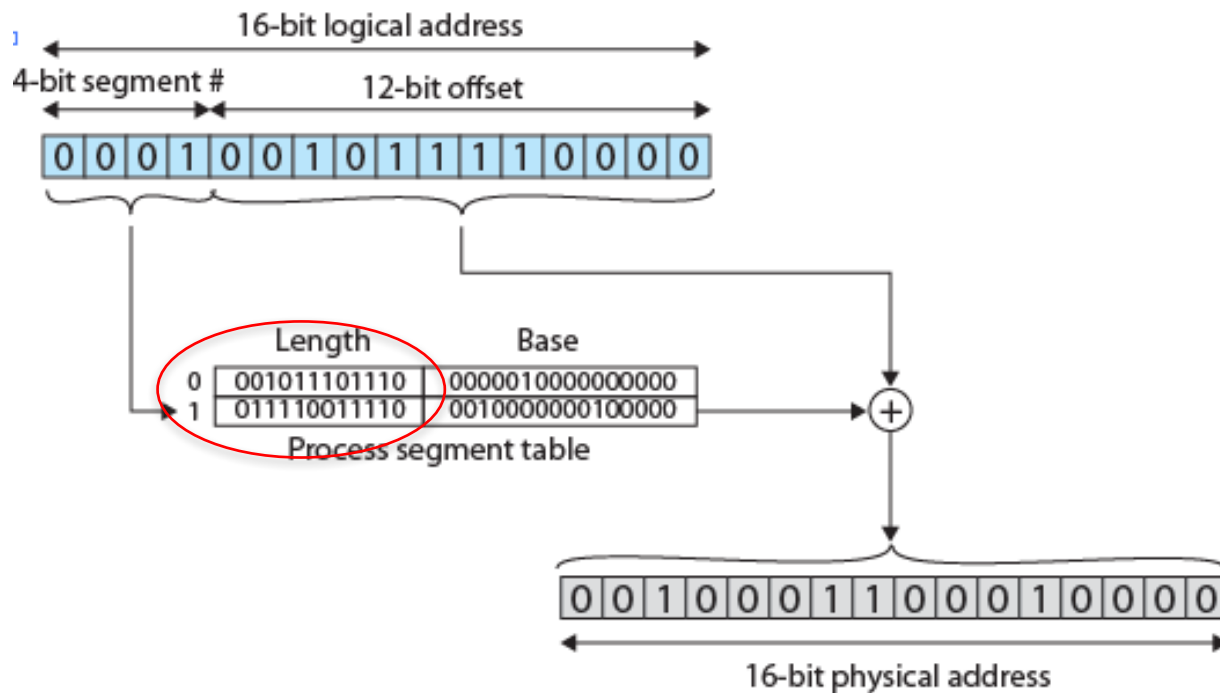
Segmentation

- Each process has a segment table
 - Each segment has a base address
 - Segments are of variable size
 - Segment table entry has to specify the **length** of a segment
 - Held in the process control block
- Entry in the segment table:
 - Base address
 - Length of a segment



Segmentation

- Each segment is a separate logical address space
 - Programmer may view memory as consisting of multiple address spaces
- Address translation via process segment table



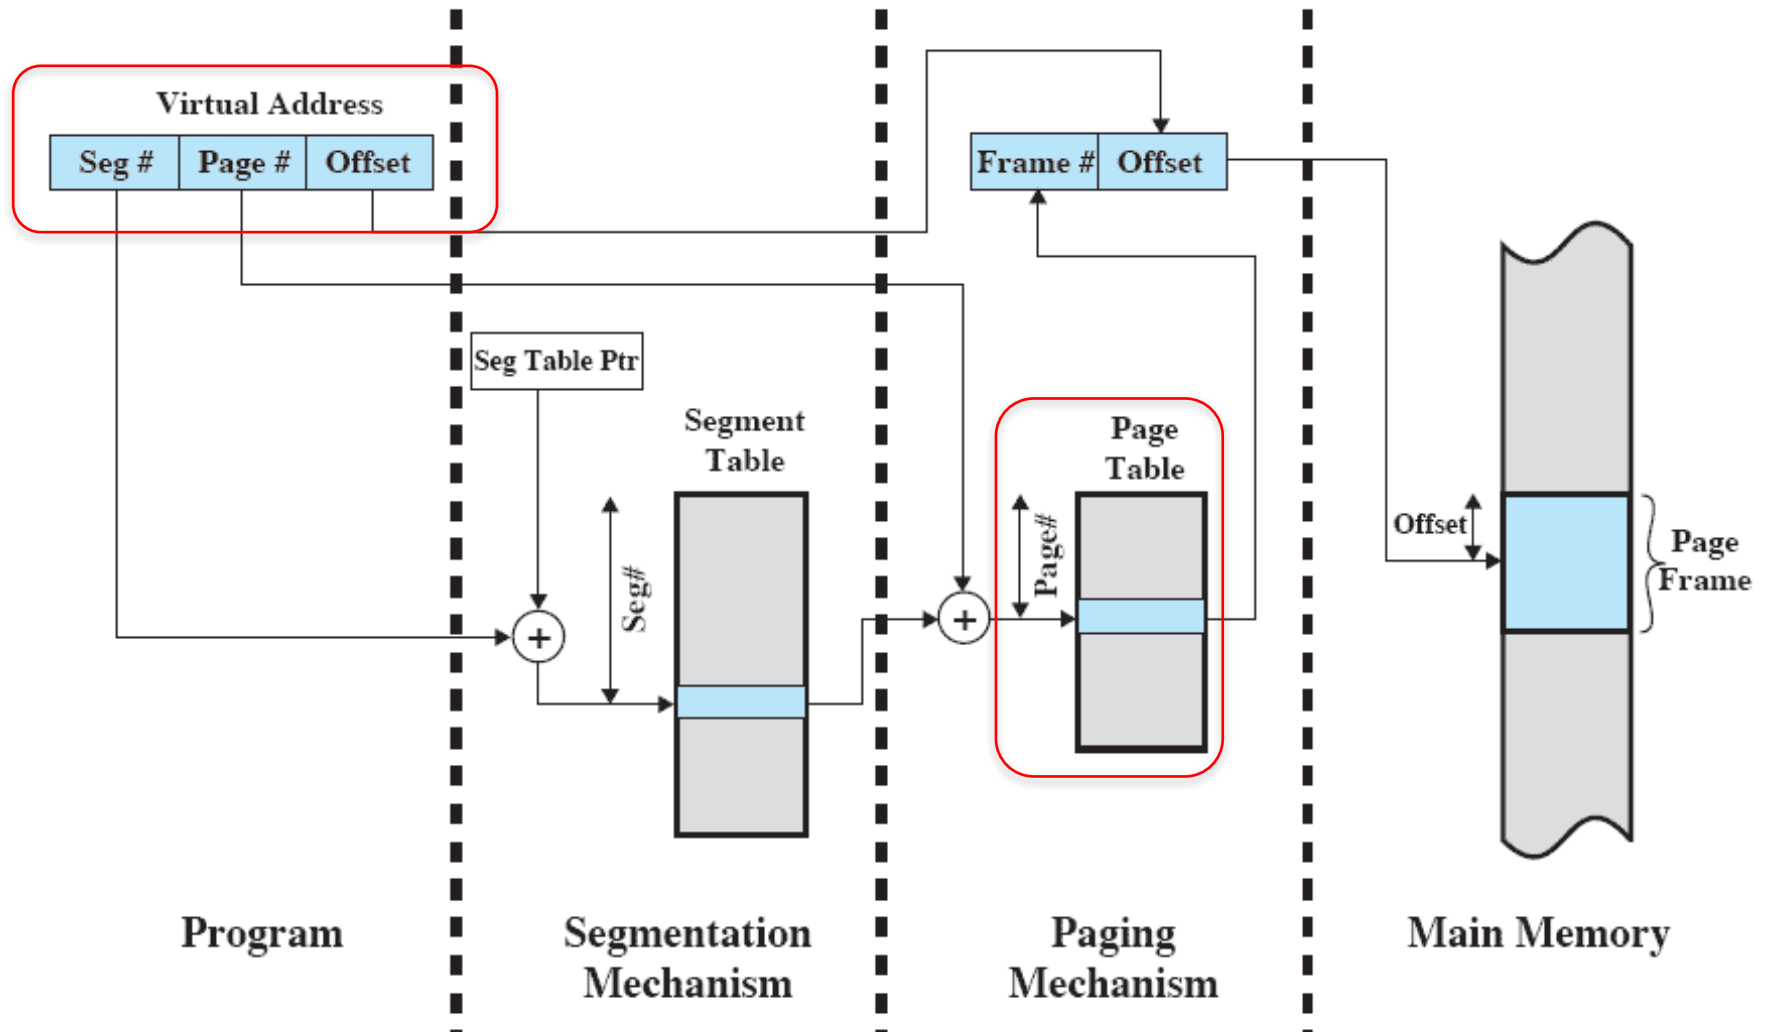
(b) Segmentation

Segmentation

- With segmentation, program may occupy non-contiguous memory
 - Each segment in separate memory area
 - Better utilisation of memory
 - Smaller contiguous segments of memory needed
- Segmentation is a dynamic partitioning scheme
 - No internal fragmentation, but external fragmentation

Combined Segmentation and Paging

- Each segment is broken into fixed-size pages



Segmentation and Paging

- Virtual address consists of at least 3 elements
 - Segment number
 - Page number
 - Offset within page
- Each segment is a separate virtual address space
 - Consists of a set of fixed-sized pages

Protection and Sharing

- Each segment can have special protection (read / write / execute policies)
 - Each segment is specified in physical memory by a base address and its length, each memory access has to be checked against these two parameters
 - E.g.: text segment is read-only
- Segments can be shared among processes