

Memory Management

CS3026 Operating Systems

Lecture 10

Page Replacement Policies


Page Replacement Algorithms

- The optimal policy (OPT) as a benchmark
- Algorithms used
 - Least recently used (LRU)
 - First-In-First-Out (FIFO)
 - Clock Algorithm

Optimal Policy OPT

- Assumption
 - 3 memory frames available
 - Process has 5 pages
 - When process is executed, the following sequence of page references occurs (called a page reference string):

Page reference string: 2 3 2 1 5 2 4 5 3 2 5 2



F = page fault occurring after the frame allocation is initially filled

OPT policy results in 3 replacement page faults (minus the initial ones to fill the empty frames, total page faults are 6)

LRU Least Recently Used Policy

Reference to page 5:

- **We replace page 3 in Frame 2**

Reference to pages 4:

- **We replace page 1 in Frame 3**



Page 1 least recently used
Page 3 least recently used

Page ref.	2	3	2	1	5	2	4	5	3	2	5	2
Frame1	2	2	2	2	2	2	2					
Frame2		3	3	3	5	5	5					
Frame3				1	1	1	4					
Page fault					F		F					

Reference to page 3:

- **We replace page 2 in Frame 1**

Reference to page 2:

- **We replace page 4 in Frame 3**



Page ref.	2	3	2	1	5	2	4	5	3	2	5	2
Frame1	2	2	2	2	2	2	2	2	3	3	3	3
Frame2		3	3	3	5	5	5	5	5	5	5	5
Frame3				1	1	1	4	4	4	2	2	2
Page fault					F		F		F	F		


- With LRU, we replace page 2 with page 3, and immediately afterwards, we need page 2 again
- LRU is not able to detect this situation
- However, it is a strategy that comes close to OPT

FIFO and Second Chance Algorithm

First-In-First-Out (FIFO)

- First in First out
 - The first page that was loaded is also the first to be discarded (oldest page) – FIFO ordering of pages
- Problem
 - Oldest page may be the most frequently used
- Is the simplest replacement policy to implement
 - Treats memory frames allocated to a process as a circular buffer
 - Pages are replaced in a round-robin style
 - Strict ordering of pages due to age
- Does not indicate how heavily a page is actually used

FIFO Example

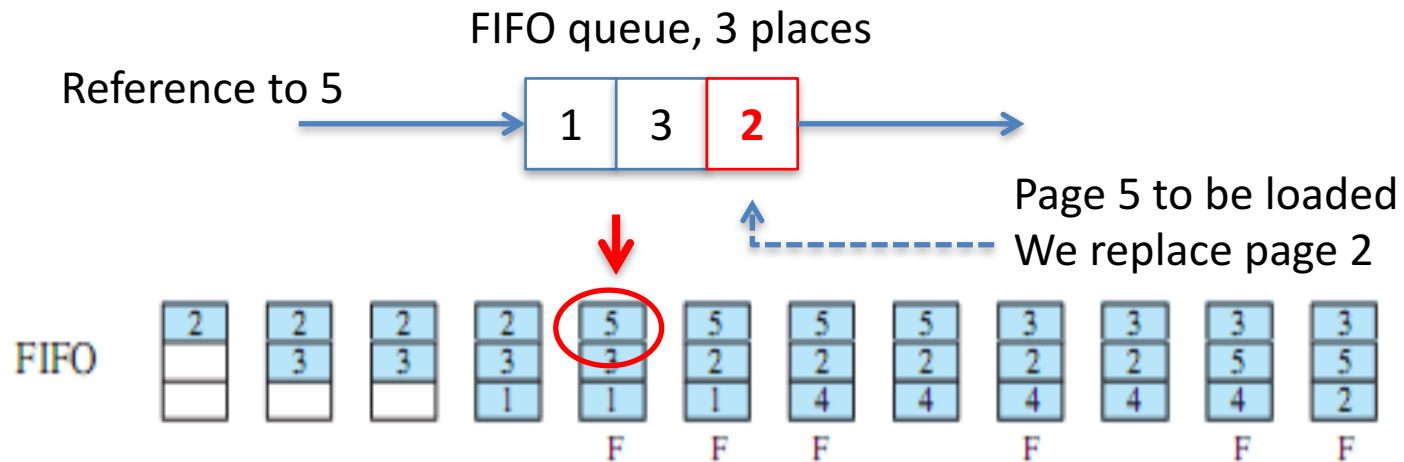


F = page fault occurring after the frame allocation is initially filled

- FIFO policy results in 6 replacement page faults (minus the initial ones to fill the empty frames, total page faults are 9)

FIFO Example

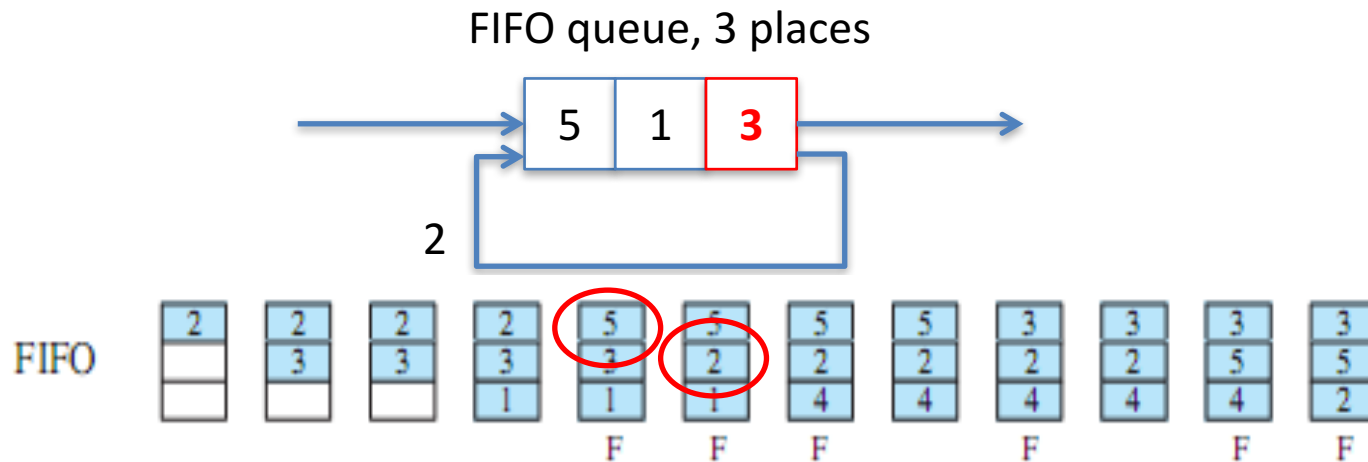
Page reference string: 2 3 2 1 5 2 4 5 3 2 5 2



- FIFO queue has 3 places (for 3 frames)
- References to pages 2, 3 and 1 will fill the available set of frames
- When page 5 is referenced, we replace the “oldest page”, which is 2
- We maintain a FIFO queue to record which page is the “oldest”

FIFO Example

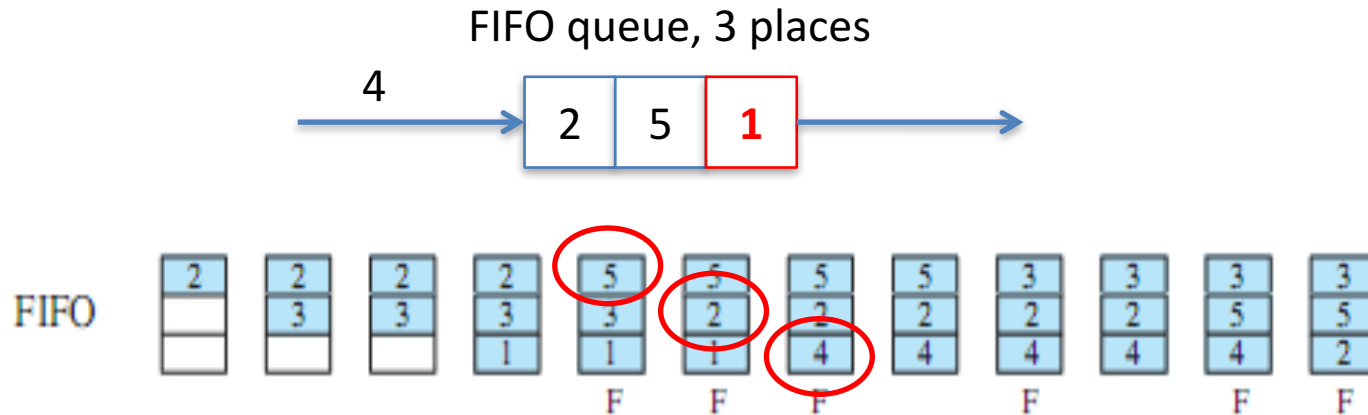
Page reference string: 2 3 2 1 5 2 4 5 3 2 5 2



- FIFO queue has 3 places (for 3 frames)
- After page 5, page 2 is referenced again immediately, therefore it has to be brought back into memory
- We replace page 3


FIFO Example

Page reference string: 2 3 2 1 5 2 4 5 3 2 5 2



- FIFO queue has 3 places (for 3 frames)
- When queue is full, oldest page discarded, new page added
- Problem: actual demand or “access frequency” for pages is not considered, may lead to frequent reload of the same page

FIFO Example



- FIFO policy results in 6 replacement page faults (minus the initial ones to fill the empty frames, total page faults are 9)

FIFO Page Replacement

- Benefit of FIFO queue
 - Sequence in queue shows which page is the oldest
 - Simple to implement
- Page replacement
 - Newly loaded page added to the end of the queue, tail of FIFO queue
 - Oldest page is the first in line, is removed when replaced
 - Gives us a simple criteria which page to replace

FIFO Page Replacement

- Problem
 - FIFO only records sequence of arrival, not the “importance” of a page
- What if the “oldest” page is the most frequently used (the “most important” page)?
 - FIFO would discard this page and frequently load again – danger of “thrashing”

Second Chance Algorithm

Second Chance Algorithm

- Is a modification of the FIFO replacement algorithm that tries to avoid replacing recently used pages
 - Tries to account for the “importance” of a page
- Uses additional information
 - Each page has reference bit: records whether a process accessed page recently
- A page is given a second chance, if it was accessed recently:
 - will be moved to the tail of the FIFO queue and becomes the youngest page, not replaced
 - Page is “overlooked” and remains in physical memory
 - Reference bit is reset by this “second chance” action


Second Chance Algorithm

- When oldest page selected for replacement, the reference bit is inspected
 - If value is 1: page was accessed recently, page gets a “second chance”:
 - **Clear reference bit, set to 0**
 - Check next-oldest page in queue
 - If value is 0: replace page
- A page is given a second chance:
- Worst case:
 - All reference bits == 1, degenerates into pure FIFO

Clock Policy

Implementation of Second-Chance Algorithm

- The set of frames is considered as laid out like a circular buffer (FIFO queue)
 - Is considered a FIFO strategy
 - Can be circulated in a round-robin fashion
- Each frame has a “use” bit
 - When a page is first loaded (a page fault that leads to its load), the “use” bit of the frame is set to 1
- Use bit is “refreshed”
 - Each subsequent reference to this page again sets /overwrites this “use” bit to 1



(a) State of buffer just prior to a page replacement

Clock Algorithm


Looking for a page to be replaced

- Initially, all frames in circular buffer are empty, use bit == 0
- Frame pointer: is regarded the “clock hand”, is called the “next frame” pointer
- Procedure for replacing a page / finding unused frame
 - Find page to be replaced (or unused frame):
 - Progress “next frame” pointer to first frame with use bit == 0
 - During this progression: reset use bit of each visited frame entries to 0
 - Load page into the first found frame with use bit == 0, set use bit to 1
 - Move position pointer to next frame
- When the circular buffer is filled with pages, the “next frame” pointer has made one full circle
 - it will be back at the first page loaded
 - points to the “oldest” page

Clock Policy

Find the right Page to be replaced

- Finding a page to replace
 - Advance “next frame” pointer from frame to frame, until a frame with “use” bit == 0 is found
 - As long as visited frames have “use” bit == 1, set “use” bit of visited frames to 0, give the pages situated there a **second chance**
 - Load page into first frame found with “use” bit == 0, set “use” bit to 1
 - Page is replaced, pointer is advanced to *next* frame



Load action: set “use” bit = 1
Visit action: reset “use” bit = 0
Load only, if “use” bit == 0

Clock Policy

Before and after Page Replacement


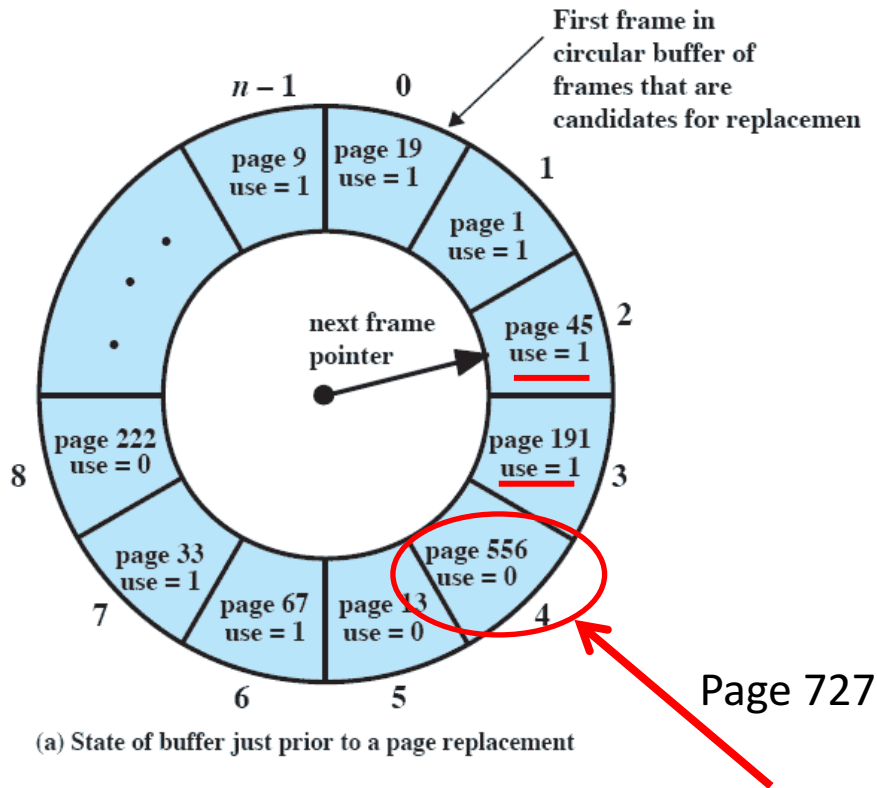


Figure 8.16 Example of Clock Policy Operation

Page-Buffering

- Avoid unnecessary read and write operations
- A replaced page is not lost, but assigned to one of two lists
 - Free page list
 - Add unmodified replaced page to this list
 - When page is needed again, its content is still held in memory, no I/O needed, can be reused immediately
 - Modified page list
 - Add replaced page that was modified, to this list
 - Pages are written to disk in clusters, saves I/O

Resident Set Management

How many Frames shall we allocate
to a process?

Resident Set

- Resident Set:

The portions of a process – the set of pages – that can be held in main memory at a particular time

- As long as the process makes memory references only to memory locations that are in the resident set, no page faults occur
- Resident set is restricted by the number of frames allocated to a process

Thrashing

- Thrashing is a situation where a process is repeatedly page faulting
 - The process does not have enough frames to support the set of pages needed for fully executing an instruction
 - It may have to replace a page that itself or another process may need again immediately
- A process is thrashing if it is spending more time on paging than execution

Thrashing



- Performance depends on a minimum number of frames allocated
 - If number of free frames decreases, page faults increase
 - Increase of page faults slows process execution
 - I/O operations
 - Instruction that leads to a page fault has to be restarted, effectively executing instructions twice
- We have to manage the size of the resident set to allocate the “right” amount of frames to a process

Principle of Locality

- Principle of Locality
 - A “locality” is a set of pages that are **actively** used together during a time period by a program
- A program may have **several different localities**
 - As a process executes, it moves from locality to locality
- If we want to prevent thrashing, we must provide a process with as many frames as needed to accommodate a “locality” – the “working set”
 - The resident set should be allowed to be of a size so that page faults are minimised

Localities

- A program may have **several different localities**
 - As a process executes, it moves from locality to locality
- Let's consider the following page reference string:



- First, pages 1,2,3 are referenced multiple times
- Then, the process switches to another locality, referencing pages 4,5,6,7

Accommodating the Locality



- Let's assume that we allocate **3 frames** as the resident set:

Page reference string: 1 2 3 1 2 3 4 5 6 7 4 5 6 7

Enough frames to
accommodate locality,
Reference to pages
1,2,3

Transition to another Locality
Frequent page faults

Locality with pages 4,5,6,7
Not enough frames
Frequent page faults



Working Set

- Working Set

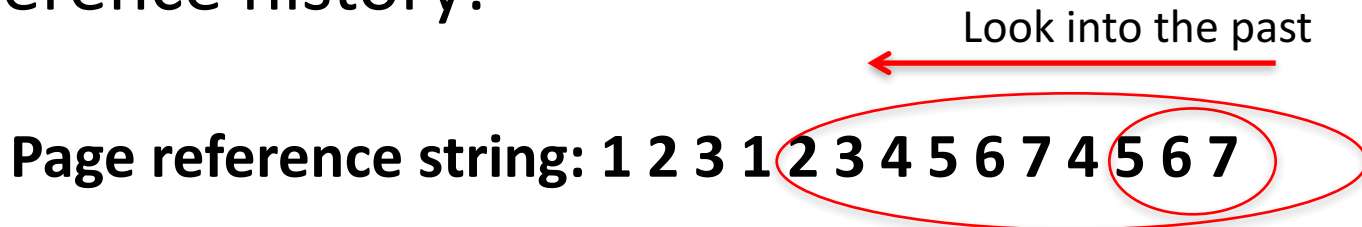
Set of pages a process is currently working with

- The working set of a process at current time t is the set of pages the process has been referencing for the past Δ time steps
- The *working set* can be larger than the *resident set* – some pages of the working set may have to be swapped in and out
- **A working set is an approximation of a program's locality**
- It is the set of pages actively used by a process and all these pages should be held in memory
- The size of this set gives us an estimate how many frames should be allocated to a process – how large the *resident set* should be

Tune the Observation Window

- Let's assume, we have 7 pages and the following reference history:

Page reference string: 1 2 3 1 2 3 4 5 6 7 4 5 6 7



The diagram shows the page reference string "1 2 3 1 2 3 4 5 6 7 4 5 6 7". A red arrow labeled "Look into the past" points from the end of the string towards the left. Two red ellipses are drawn around the string: a larger one encircling the last seven digits "2 3 4 5 6 7 4" and a smaller one encircling the last three digits "5 6 7".

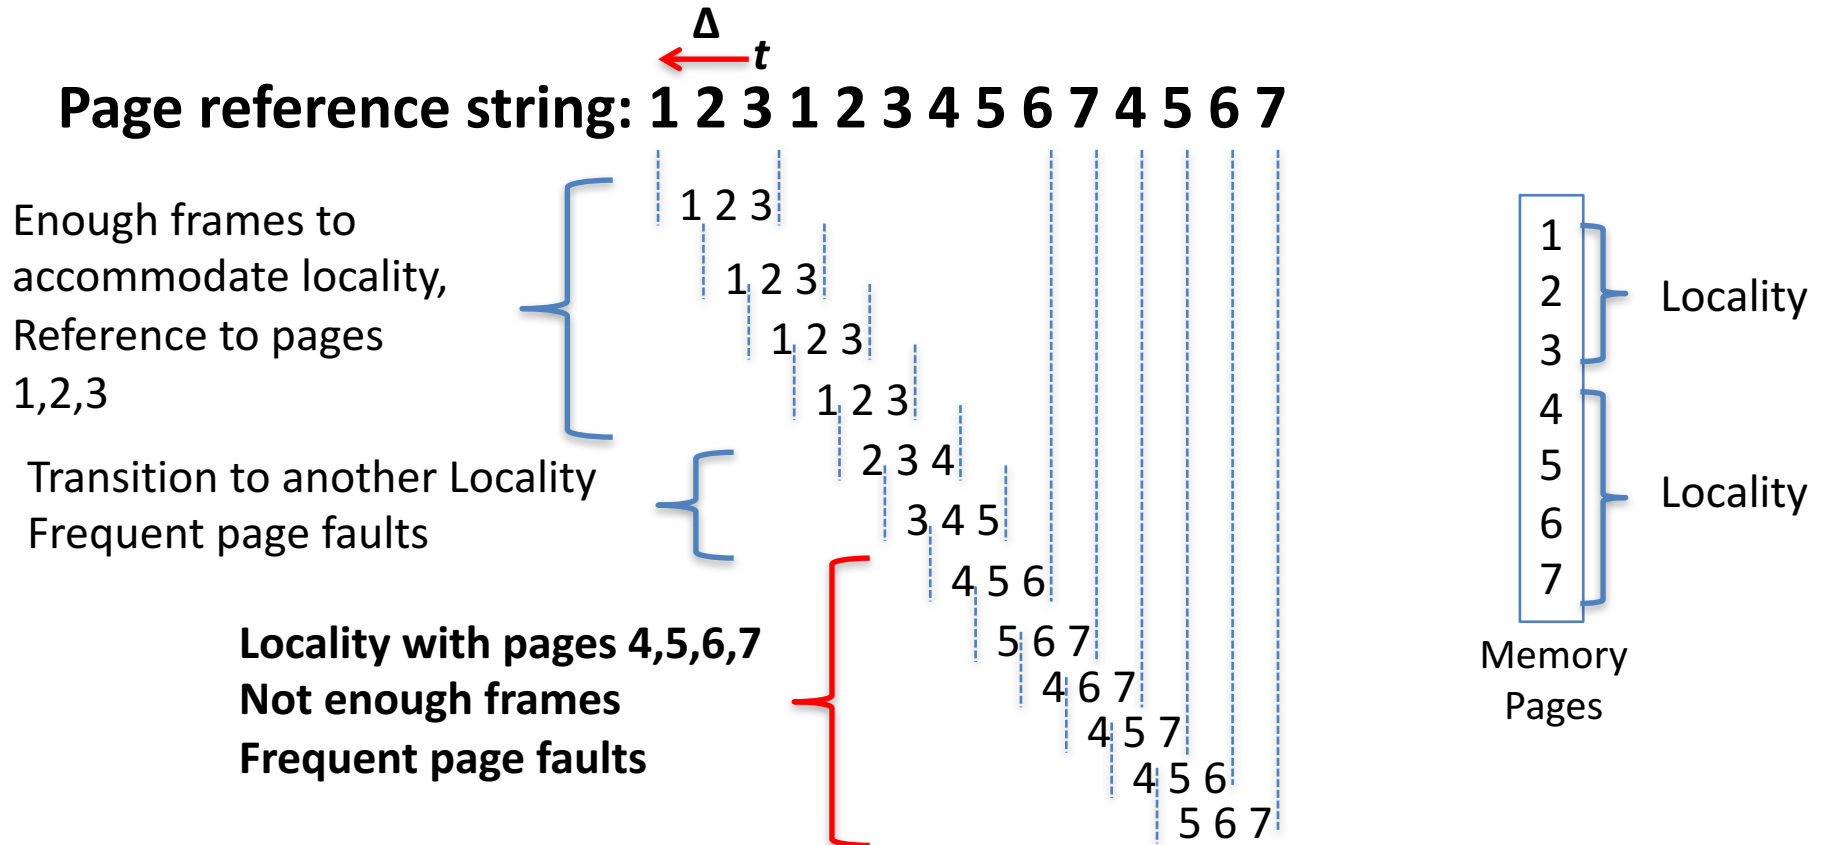
- If we choose a working set window $\Delta = 3$, then we only see part of the locality, which gives the impression that we only need 3 frames for pages 5, 6, 7
- If we choose a $\Delta = 10$, we see part of the previous locality, which give the impression that we need 6 frames for pages 2, 3, 4, 5, 6, 7

Working Set Window

- We look into the past to estimate the future size of the “working set”
 - Set of pages that have been referenced in the recent past of a process execution
- How far to look into the past:
 - Define an observation window or “working set window” Δ
- Parameter Δ : number of most recent page reference observations
 - If Δ is too small, not all actively used pages (the current “locality”) will be in the working set window
 - If Δ is too large, we may look back too far and see several different localities
 - Does not give us clear estimate how many frames we actually need for this process

Working Set

- The working set may change with each page reference
- How far shall we look into the past?
- Let's assume that we allocate **3 frames** to the process:



Working Set Strategy

- We need to ensure that we can keep the working set in memory:
 - Working set determines the minimum size of the resident set to avoid thrashing
 - If the working sets of all processes exceed total number of page frames, one or more processes have to be swapped out
- We can monitor the working set over time, provides insight in the required amount of page frames

Working Set Size WSS

- Working set size WSS_i
 - Size of the working set for process i
 - Depends on how many past observations Δ are taken into account
 - We assume that the pages accessed by process i in the observation window Δ are most likely to be accessed in the future as well
 - approximate the future working set
 - the number of pages observed determines the number of frames to be allocated
- Therefore, a process needs at least WSS_i frames to avoid thrashing
- Operating system monitors the working set of a process and allocates WSS_i frames to the process

Working Set

- What to do if the Resident set is too small
 - Danger of thrashing !
 - Consult the “working set”: Add more frames
 - If we run out of frames, suspend process and try to allocate again later
- What to do if the Resident Set is too large
 - There may be pages that are loaded into memory, but haven’t been referenced for a long time
 - Periodic discarding of pages from the resident set that are not in the working set

Working Set Strategy, Problem

- Problems
 - Page references have to be logged
 - Page references have to be ordered
 - Optimal time window size Δ is unknown at runtime and may vary
- Simplification
 - Monitor number of page faults per time unit per process, instead of monitoring the working set

Page-Fault Frequency

- Use page-fault frequency to control thrashing
 - Thrashing has a high page fault rate
- Idea
 - If page-fault rate is high, then process needs more frames
 - If page-fault rate is low, then process has too many frames
- We can establish upper and lower bounds on the desired page-fault rate, controls frame allocation
 - If actual page-fault rate exceed upper limit, process receives an extra frame
 - If actual page-fault rate falls below lower bound, process loses one frame
- If there are no more frames to allocate, a process exceeding the upper bound, will be suspended and its frames reallocated to other processes

Allocation Strategies

Frame Allocation to Processes

- Each process is allocated free memory
 - Each process has a Free-frame list
 - Demand-paging will allocate frames from this list due to page faults
 - When the free-frame list is exhausted, a page replacement algorithm will choose a page to be replaced
- With multiple processes: how many free frames should be allocated to each process?

Resident Set Management

	Local Replacement	Global Replacement
Fixed Allocation	<ul style="list-style-type: none">•Number of frames allocated to a process is fixed.•Page to be replaced is chosen from among the frames allocated to that process.	<ul style="list-style-type: none">•Not possible.
Variable Allocation	<ul style="list-style-type: none">•The number of frames allocated to a process may be changed from time to time to maintain the working set of the process.•Page to be replaced is chosen from among the frames allocated to that process.	<ul style="list-style-type: none">•Page to be replaced is chosen from all available frames in main memory; this causes the size of the resident set of processes to vary.

Resident Set Allocation

- Fixed-allocation
 - A process has a fixed pool of frames in main memory
 - If this pool is exhausted, a page fault will cause a page replacement
- Variable-allocation
 - Allow the pool of frames to vary in size over the lifetime of a process

Replacement Scope

- We can classify page replacement in terms of allocation behaviour
 - Global page replacement
 - A process selects a frame for replacement from a global list of frames (frames of all processes)
 - Competition with other processes: A page replacement will interfere with another process' memory allocation
 - Local page replacement
 - A process can only select a frame for replacement from its own local list of frames

Replacement Scope

- Disadvantages
 - Global page replacement
 - Performance of a process depends on external circumstances, as other processes may replace its pages
 - Waste of performance
 - Local page replacement
 - Process is restricted to its own allocation of physical memory
 - Even if memory is not used by other processes, it is not available globally
 - Waste of resources
- Global replacement generally results in better system throughput and is the more commonly used method

Thrashing

- Limit the effects of thrashing by using a local replacement algorithm
 - Processes can only be served from their allocated pool of free memory frames
 - Process cannot “steal” frames from other processes
- Problem
 - When process starts thrashing, it will occupy the paging device queue most of the time
 - This increases the time for other processes to be served by the paging device

Working Set

- How many processes can be started?
 - Operating system will start processes and allocate frames as needed by their working sets
 - Working sets can change in size and need for frames
 - If the sum of all needed frames exceeds the actually available number of frames, one of the processes is selected for suspension
 - Its pages are swapped out and its frames reallocated to other processes
 - The suspended process can be restarted later, when enough frames become available again

FIFO: Belady's Anomaly

- Problem of FIFO
 - Does not indicate how heavily a page is actually used
- Possible solution: increase number of physical frames
 - FIFO shows Belady's anomaly: the page-fault rate increases, when the number of frames is increased
- Example
 - Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
 - 3 frames vs 4 frames:

1	4	5
2	1	3
3	2	4

9 total page faults

3 Frames, 5 pages

1	5	4
2	1	5
3	2	
4	3	

10 total page faults

4 Frames, 5 pages