

Transaction Management

Managing Concurrent Access to Data
Objects

CS3524 Distributed Systems

Lecture 07

Transaction

A transaction is a controlled sequence of actions for the manipulation of a database

- Objective:
 - Guarantee that data always remains in a consistent state
- Avoid:
 - Loss of data due to system failure (Database Integrity, Durability of data)
 - Problems resulting from concurrent access of multiple users (consistent manipulation)

Database Integrity

- There are two categories of problems that may compromise the integrity of databases:


System failure

- Malfunction of hardware and software
- crash during data manipulation
- leaves data in a corrupted / inconsistent state

Concurrent access

- Users interfere with each others' data manipulation operations
- overwrite / delete each others data
- may leave data in a corrupted / inconsistent state

Transaction



- Transactions are a controlled sequence of read and write operations:
 - Have a clear “start”: the first operation of a transaction
 - Have a clear “end”:
 - Commit: transaction ends successfully
 - Transactions can be aborted any time
 - Abort: transaction is “undone”

Transaction Atomicity

- Transactions are regarded as ***atomic***:
 - “All-or-nothing”: the complete sequence of operations is regarded as an indivisible unit - either all of the operations occur or none of them
- A transaction transforms databases from one **consistent state** to another **consistent state**
- Transactions are either committed at the end or aborted (anytime):

Commit: transaction ends successfully, all manipulations are stored in database, database is in **new consistent state**

Abort: transaction is “undone” – manipulations performed during transaction have no effect on the database, database remains in **old consistent state**

Programming Transactions

Transaction Commit	Account A	Account B
<pre>account_a = getAccount(db,"account_1") account_b = getAccount(db,"account_2") open_transaction() balance1 = read(account_a) balance1 = balance1 - 200 write(account_a, balance1) balance2 = read(account_b) balance2 = balance2 + 200 write(account_b, balance2) commit()</pre>	1000.00 800.00	 1000.00 1200.00
Both accounts have changed	800.00	1200.00

Programming Transactions

- Transaction Abort:
 - A user can decide to abort a transaction (e.g. Pressing a button in a GUI to abort a payment transaction):

Transaction Abort	Account A	Account B
<pre>open_transaction() balance1 = read(account_a) balance1 = balance1 - 200 write(account_a, balance1) balance2 = read(account_b) balance2 = balance2 + 200 write(account_b, balance2) user_decision = getUserDecision() if(user_decision == "abort") abort()</pre>	1000.00 800.00	 1000.00 1200.00
Both accounts have to be returned to their original balance	1000.00	1000.00



How can the Database system do that?

Handling Transaction Abort

- Transaction Abort
 - Intentionally by client / user programs interacting with database server
 - Aborted by database system
 - Client session timeout
 - For resolving conflicts between concurrent transactions (e.g. deadlocks)
- Action: Transaction “Rollback”
 - Database is “rolled back” to its previous consistent state
 - All manipulative operations are undone (or not applied, depending on the implementation of the transaction management)

System Failure

Transaction Commit	Account A	Account B
<code>open_transaction()</code>		
<code>balance1 = read(account_a)</code>	1000.00	
<code>balance1 = balance1 - 200</code>		
<code>write(account_a, balance1)</code>	800.00	
<code>balance2 = read(account b)</code>		1000.00
<code>balance2 = balance2 + 200</code>		
<code>write(account_b, balance2)</code>		
<code>commit()</code>		




The table is crossed out with a large red 'X' that spans across the bottom half of the rows, from the row containing 'balance2 = balance2 + 200' down to the 'commit()' row.

- A system failure occurs, after change of account A, but before account B can be changed – the database is “inconsistent”!

Recovery from System Failure

- System failures may corrupt / destroy data
- Database management systems employ a “Recovery Manager” to re-instate a database to its last consistent state before system failure
 - Backup: record last consistent state
 - Log files / Transaction Journals:
 - A Recovery manager records each data manipulation actions since last backup
 - “Roll Forward”: recovery manager uses log / journal records to redo all committed transactions


Concurrent Transactions



- Usually many transaction may run concurrently
- They may manipulate the same data concurrently
- Isolation:
 - The Transaction management has to employ concurrency control mechanisms and protocols to properly “isolate” transactions from each other so that they do not interfere with each other

Database Integrity

Transaction



- **Consistency:**
 - Safeguard against “undesired” / “wrong” manipulation of data (user concern)
- **Durability:**
 - Safeguard against loss of data (system concern)
- Transaction management systems have to guarantee:
 - In case of system failure, data can be recovered
 - Isolate concurrently executing transactions from each other
 - Avoid race conditions (overwrite each others’ results)
 - Protect transactions from reading partial / uncommitted results of other transactions

The ACID Properties

- **Atomicity:** Either *all operations* of a transaction are performed, or *none* of them
- **Consistency:** A transaction transforms the system from one *consistent state* into another
- **Isolation:** An incomplete transaction cannot reveal its *intermediate state* to other transactions until committed
- **Durability:** Once a transaction is committed, the system must guarantee that the results of the transaction will *persist*, even if the management system subsequently fails.


Atomicity

- Principle of “All-or-nothing”:
 - A transaction either completes successfully, and the effects of all of its operations are recorded or (due to system failure or deliberate abort) it has no effect at all
- Failure atomicity (“*nothing*”):
 - the effects are atomic even when the server crashes – intermediate results are “rolled back”
- Durability (“*all*”):
 - after a transaction has completed successfully, all its effects are saved in persistent storage

Durability

- **Durability:**
 - Once a transaction is committed, the system must guarantee that the results of the transaction will persist, even if the management system subsequently fails.
- In order to support failure atomicity and durability, data objects must be *recoverable*
 - regular backups
 - transaction logs that record manipulation actions

Concurrency Problems



- The interleaving of actions of concurrently executing transactions may lead to severe problems
- Lack of Isolation:
 - Lost Update: Transactions may overwrite each others' updates
 - Inconsistent Retrieval: Transactions base their calculations on retrieved data that is not yet committed by other transactions ("dirty read")
- It depends on the sequence of actions, scheduled from different transactions, whether we create inconsistencies in our databases

The Lost Update Problem

Time	Transaction T1	Account A	Transaction T2
t1	open_transaction()	100.00	
t2	read(balance_a) ←	100.00	open_transaction()
t3	balance_a = balance_a + 100	100.00 →	read(balance_a)
t4	write(balance_a) →	200.00	balance_a = balance_a - 10
t5	commit()	90.00 ←	write(balance_a)
t6		90.00	commit()

- Schedule:

$$S = \{ T1.read(), T2.read(), T1.write(), T2.write() \}$$
- Problem:
 - Last write of transaction T2 overwrites previous write of transaction T1
- Race condition !

Inconsistent Retrieval Problem

- Also called the “inconsistent analysis” problem
- Occurs when
 - One transaction first retrieves data via a sequence of read operation
 - A second transaction performs write operations concurrently (after read operations) and overwrites some of this data in the database
- Problem:
 - The reading transaction bases its calculations on the original retrieved data (not yet committed) that have been changed in the meantime in the database itself – calculations are out-of-date!

Inconsistent Retrieval Problem

(due to “Dirty Read”)

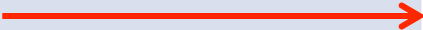
Time	Transaction T1	A	B	Transaction T2
t1	open_transaction()	100	50	
t2	sum = 0	100	50	open_transaction()
t3	read(balance_a)	100	50	read(balance_a)
t4	sum = sum + balance_a	100	50	balance_a = balance_a - 10
t5		90	50	write(balance_a)
t6	read(balance_b)	90	50	commit()
t7	sum = sum + balance_b	90	50	
t8	commit()			

- Transaction T1 calculates:
 - $\text{sum} = 100 + 50 = 150$
- Problem: transaction T2 changed the database
 - Balance of account A + balance of account B = 140 !

Problems with Aborted Transactions

- Effects of transactions are recorded in persistent memory only if they are committed
- If a transaction is aborted, then they should not have any effect on subsequent transactions
- Problems
 - Cascading Aborts
 - Premature Writes

Uncommitted Dependency Problem (due to “Dirty Read”)

Time	Transaction T1	Account A	Transaction T2
t1	<code>open_transaction()</code>	100.00	
t2	<code>read(balance_a)</code>	100.00	
t3	<code>balance_a = balance_a + 100</code>	100.00	
t4	<code>write(balance_a)</code>	200.00	<code>open_transaction()</code>
t5		200.00	<code>read(balance_a)</code>
t6	<code>rollback()</code> 	100.00	<code>balance_a = balance_a - 10</code>
t7		190.00	<code>write(balance_a)</code>
t8		190.00	<code>commit()</code>

- Problem:
 - Transaction T2 reads an intermediate, uncommitted result of T1 and bases its own calculations on that “dirty read”
 - T1 aborts transaction and performs a “rollback”

Cascading Aborts

- Cascading aborts
 - If a transaction has seen the effects of an aborted transaction, it has to abort as well
 - This can cause for other transactions to abort as well – a cascade of aborts can occur
- Solution
 - Read is only allowed as long as there are no concurrent uncommitted write operations during transaction:
 - A transaction is only allowed to read objects that are committed by other transactions
 - Any read operation must be delayed until other transactions have committed or aborted write operations on shared data objects

Concurrency Control

- How can we avoid these problems?
- We need a concurrency control mechanism


Concurrency Control

Saveguard Isolation between
Transactions

Concurrency Control

- Concurrency control enforces a particular order of operations of concurrent transactions
 - It creates a *Transaction schedule*
- The mechanisms commonly employed for concurrency control are
 - Locking:
 - Reserve a data object for exclusive access by a single transaction
 - Most practical systems use locking
 - Optimistic concurrency control
 - Timestamp ordering
- How do we know that such a concurrency control mechanism is “correct” ?
- How do we know that the enforced order of operations represents a “correct” transaction schedule?

What is a correct Transaction Schedule?



- A “schedule” is a sequence of operations from different transactions – transactions operate in an interleaved fashion
 - Such a schedule may compromise the integrity / consistency of a database


Transaction Schedule

- We specify
 - A schedule is “correct”, if it preserves the consistency of a database
 - No lost update
 - No inconsistent retrieval
- Radical solution
 - Completely serialized execution of transactions

Serial Schedule

- Serial Schedule:
 - A schedule where the operations of each transaction are scheduled for execution consecutively, without any interleaving of operations from different transactions
- There is no interference or concurrency problem
- Therefore, outcome of the execution of such a schedule preserves consistency of database
- It is a “correct” schedule!

Serial Schedule



Serialised Transaction Schedules


- Radical solution
 - Completely serialised execution of transactions
 - In order to avoid the concurrency problems described, one obvious solution would be to schedule only one transaction at a time for execution
- Such a completely “serialised” schedule can be regarded as “correct”:
 - Enforces Isolation: transactions are completely isolated and cannot interfere with each other
 - Enforces consistency: Serial execution of transactions never leaves a database in an inconsistent state
- However: we want a concurrent execution of transactions, that preserves the ACID principles

Serialised and Concurrent Execution


- Objective:
 - The aim of modern database management systems is to allow multi-user access and to maximise the degree of concurrency between transactions
- How can we guarantee that a particular transaction schedule does not violate the consistency of a database – that it is “correct”?
- We use the observation that completely serialised schedules are always correct

Transaction Schedules

Serial Schedule



Non-serial / Interleaved Schedule



Serial Equivalence, Serialisability

- We try to find a non-serial schedule that is ***equivalent*** to a corresponding serial schedule:
 - A non-serial schedule over a set of transactions is correct if it produces the same results as a completely serial execution of the same set of transactions
- If we can show or guarantee serialisability of a transaction schedule, then we can guarantee consistency of data manipulation

Serial Equivalence

- Definition: Serial Equivalence

Two or more transactions are serial equivalent, if they produce the same result operating in an interleaved fashion, as if they would operate in a completely serialized fashion.

- Serial Equivalence is used as a design criterion for concurrency control protocols!

Serialisability

- What makes a non-serial schedule equivalent to a serial schedule?
- Observation:
 - It is the sequence and order of **read / write** operations (of different transactions) in a schedule that determines whether the schedule is serialisable
- A concurrency control protocol must enforce such a correct sequencing!

Concurrency Control Protocols

- Two kinds of concurrency control behaviour
 - Transactions wait to avoid conflicts (pessimistic concurrency control)
 - Transactions are restarted after conflicts have been detected (optimistic concurrency control)
- Methods
 - Locking
 - Optimistic concurrency control
 - Timestamp ordering