

Distributed Transactions I

CS3524 Distributed Systems

Lecture 11

Distributed Transactions

- A transaction is distributed, if it invokes read() and write() operations on objects that reside on different servers

T = openTransaction

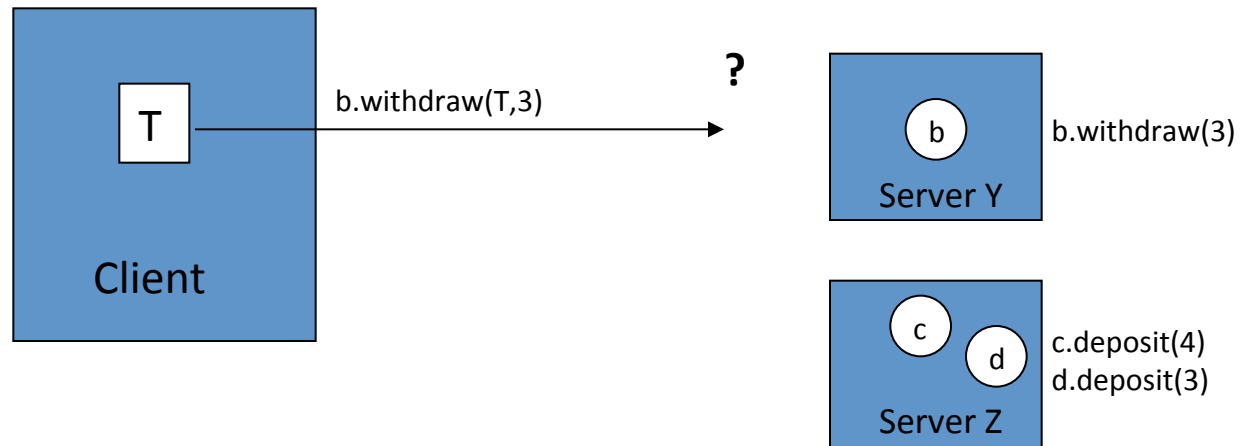
a.withdraw(4)

c.deposit(4)

b.withdraw(3)

d.deposit(3)

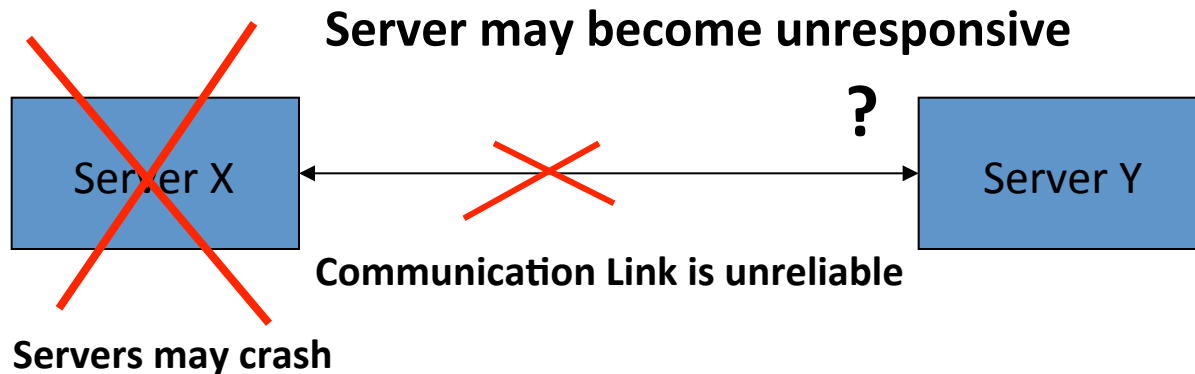
closeTransaction



Atomicity for Distributed Transactions


- **Atomicity** must be guaranteed:
 - Either **all servers** involved ***commit*** their part of distributed transaction or **all servers** involved ***abort*** their part of the distributed transaction
- This requires communication between servers
 - We need a **Commit Protocol**

Problem of Distribution



- Communication links are unreliable
 - Servers can crash and become unreachable
 - Communication links may **temporarily** fail – the client does not receive any information about success / failure of (sub-) transactions
- Sub-transactions can be affected by server crashes
 - is it possible for the server to recover and to complete the transaction?
 - Is there a need to abort the complete distributed transaction?

Distributed Commit Protocol



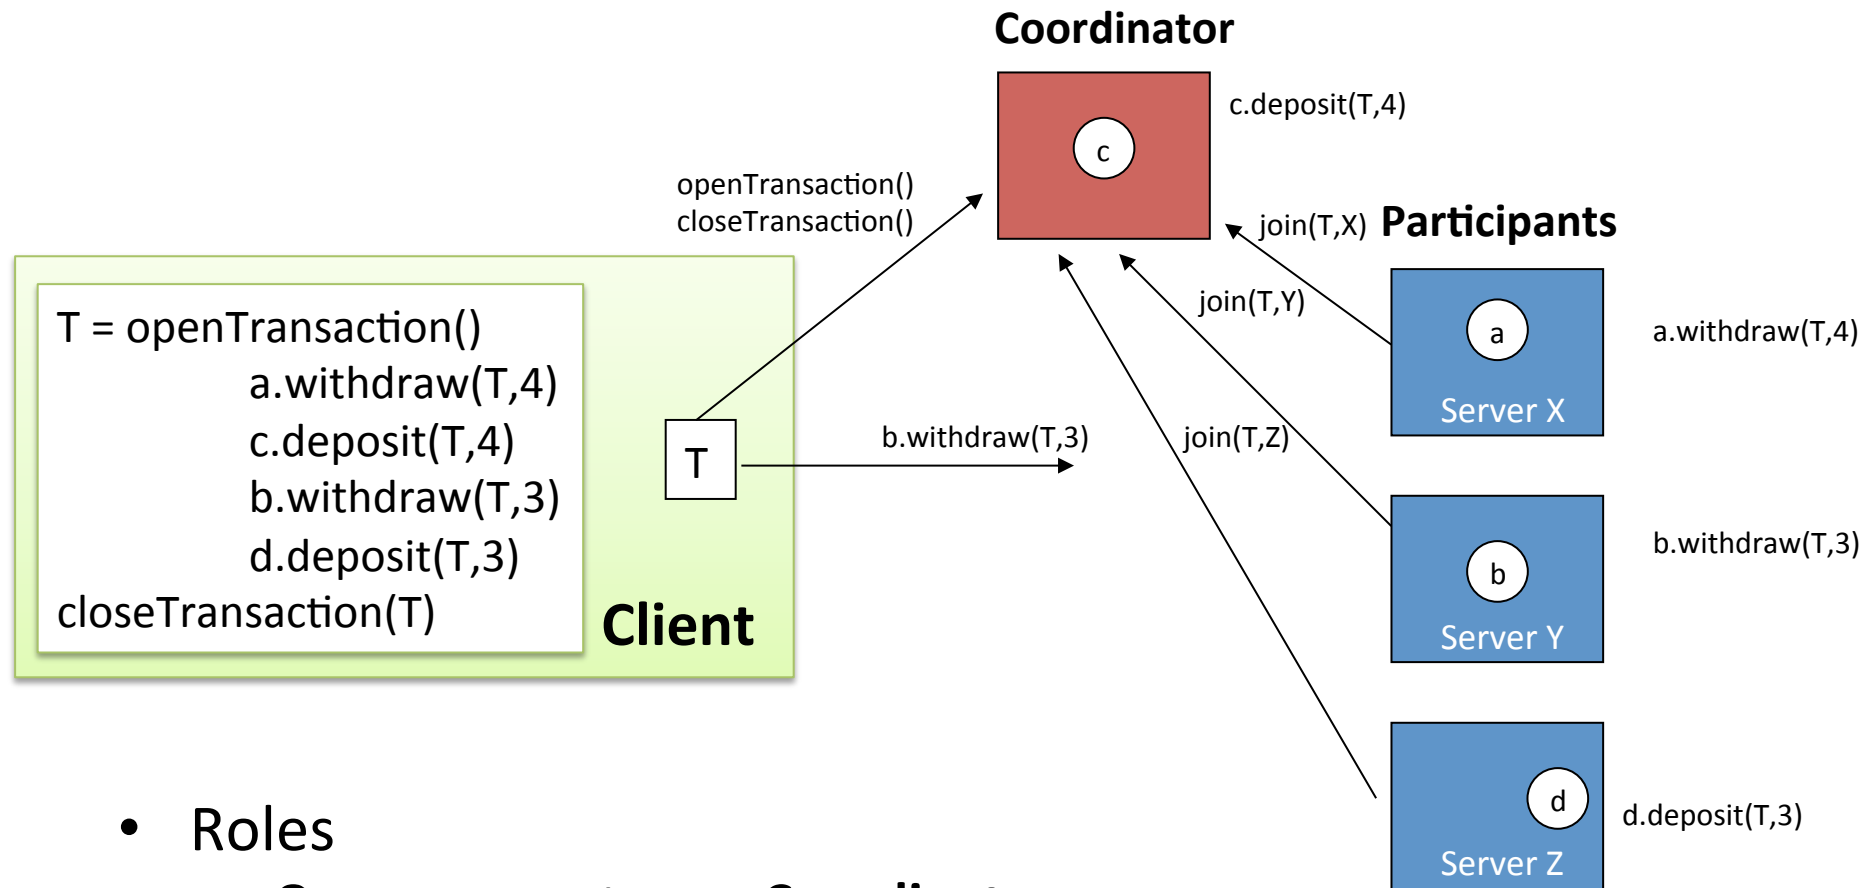
- Any lock can be released only after commit or abort
 - In distributed transactions, a set of servers may manipulate data
 - these manipulations have to be committed / aborted in an orderly fashion
- Commit protocol
 - Servers participating in such a transaction have to communicate and follow a particular protocol to complete their transaction

Distributed Commit Protocol

- Commit Protocol:
 - is a procedure that allows servers to coordinate their actions for committing a distributed transaction in an orderly fashion
- Has to guarantee Atomicity
 - Guarantee that either all of its operations (which are distributed across multiple servers) are carried out or none of them
- Has to guarantee Durability
 - has to be designed to account for possible failure situations
 - the commit phase of a transaction has to be finished (either committed or aborted) despite server crashes, lost messages etc.

Distributed Commit Protocol

Coordinator and Participant



- Roles
 - One server acts as a **Coordinator**
 - All other servers act as **Participants**

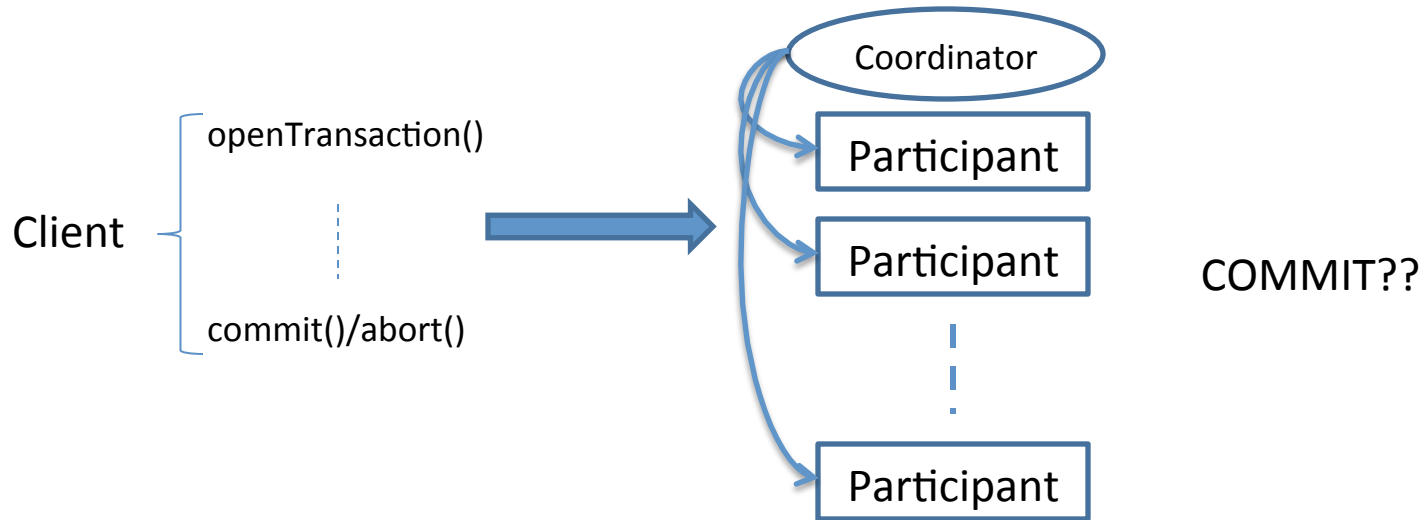
Distributed Commit Protocol

- The “Coordinator”:
 - One of the servers involved in a distributed transaction takes on the role of the “coordinator”
 - a client wishing to perform a transaction contacts a coordinator by sending an *openTransaction()* request
 - The coordinator server becomes the coordinator of the distributed transaction
 - The coordinator is responsible for committing or aborting the transaction
- The “Participant”:
 - Each server participating in a distributed transaction *registers* with the coordinator as a “participant”

Distributed Commit Protocol

- Objective
 - All participating servers have to coordinate their actions during the commit phase
 - All participating servers must be sure that the whole distributed transaction was either successfully committed or aborted

Simple Commit Protocol



- Simple (“One-Phase”) Commit Protocol:
 - The coordinator **keeps repeating a commit request**, until all of the participants have acknowledged the commit
- Problem – Is **not feasible**:
 - Coordinator may wait a long time or forever until all participants responded
 - What if one participant fails?
 - Only the coordinator can abort the transactions, servers cannot make a unilateral decision to abort a transaction
- One-phase commit protocol is inadequate !!
- Better solution: Two-phase commit protocol

Two-Phase Commit Protocol

Two-phase Commit Protocol

- Is based on a **voting scheme**:
 - All participants reach a **consensus** – participants vote *either to commit or to abort*
- Preserves **atomicity**:
 - if a part of a transaction is aborted, the complete transaction must be aborted
 - This requires informing all participants about such an abort and ***making sure*** that they get the message
- **Unilateral Abort possible**:
 - Servers (coordinator, participants) can unilaterally decide to abort their part of a transaction, **if their behaviour can be detected by coordinator**

Two-phase Commit Protocol

Phases

- Consists of two phases

Voting phase – ***prepare to commit***

Coordinator asks participants whether they are ***prepared to commit*** or ***abort***

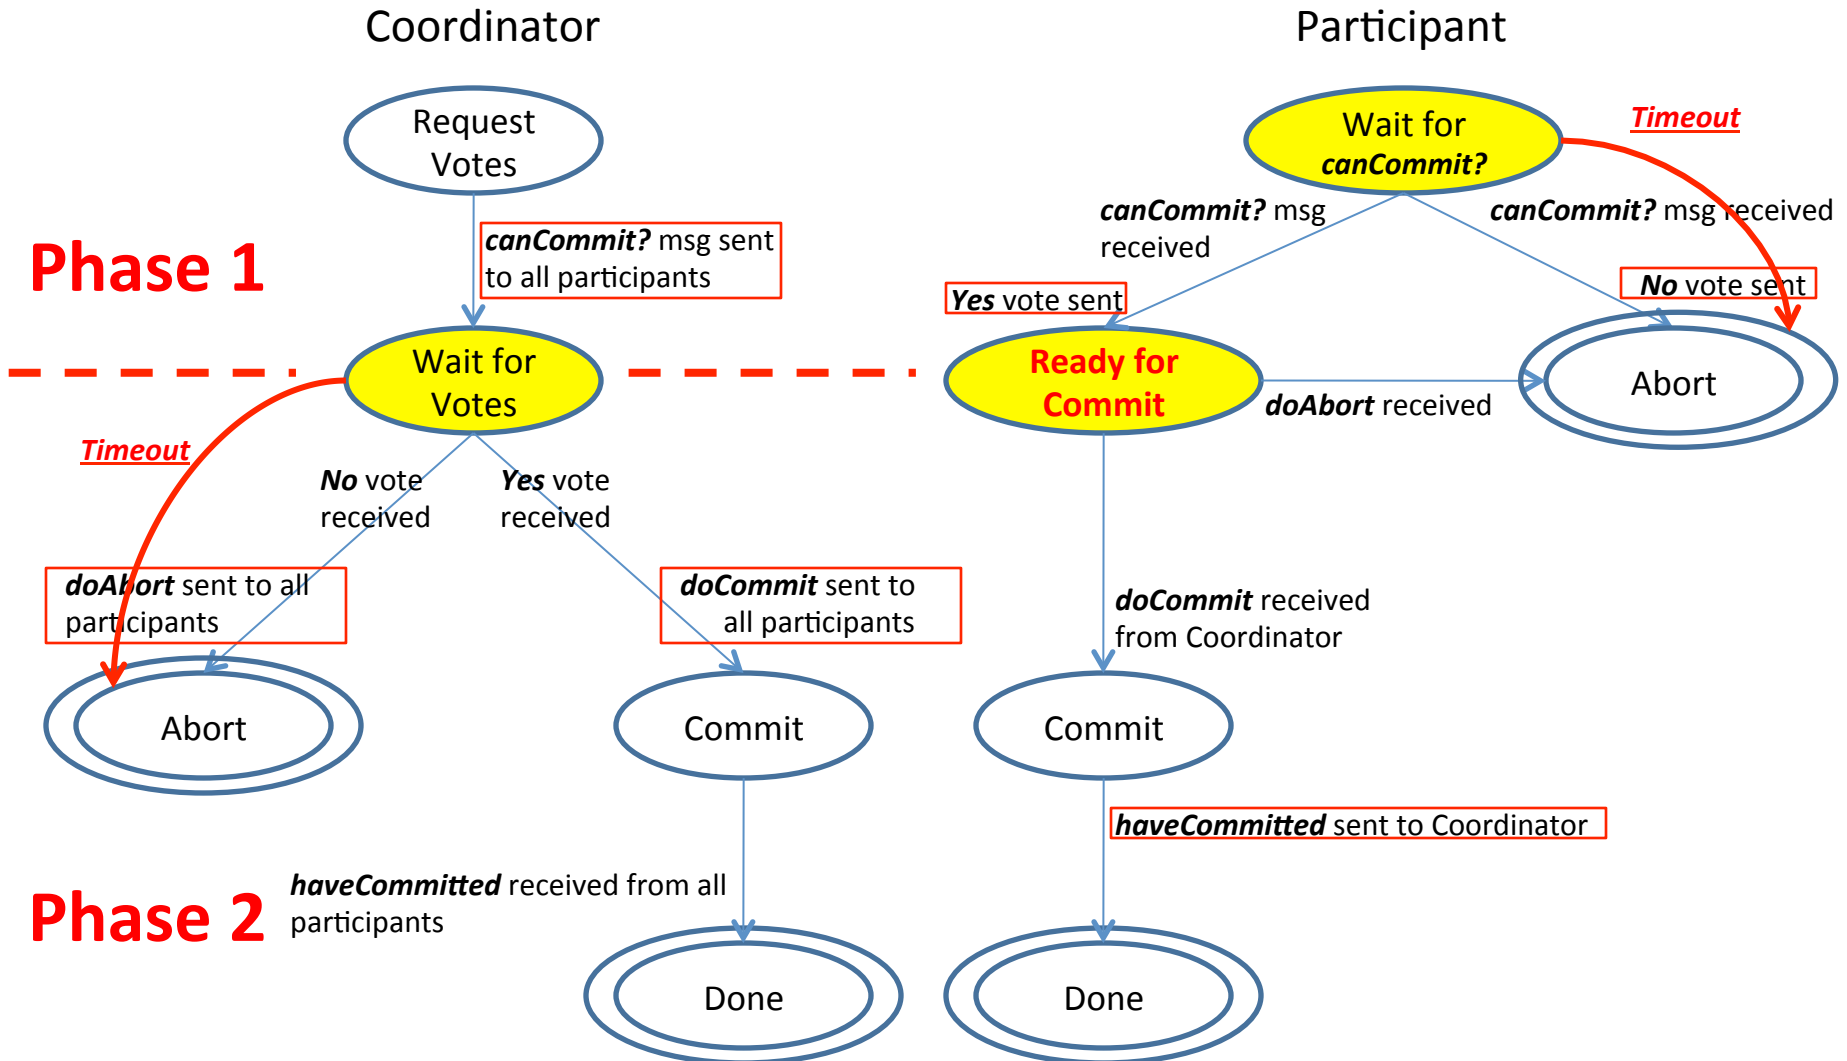
Completion phase – ***commit***

If all votes are for commit, coordinator asks participants to ***commit***

If at least one vote is for abort, coordinator asks all participants to ***abort***

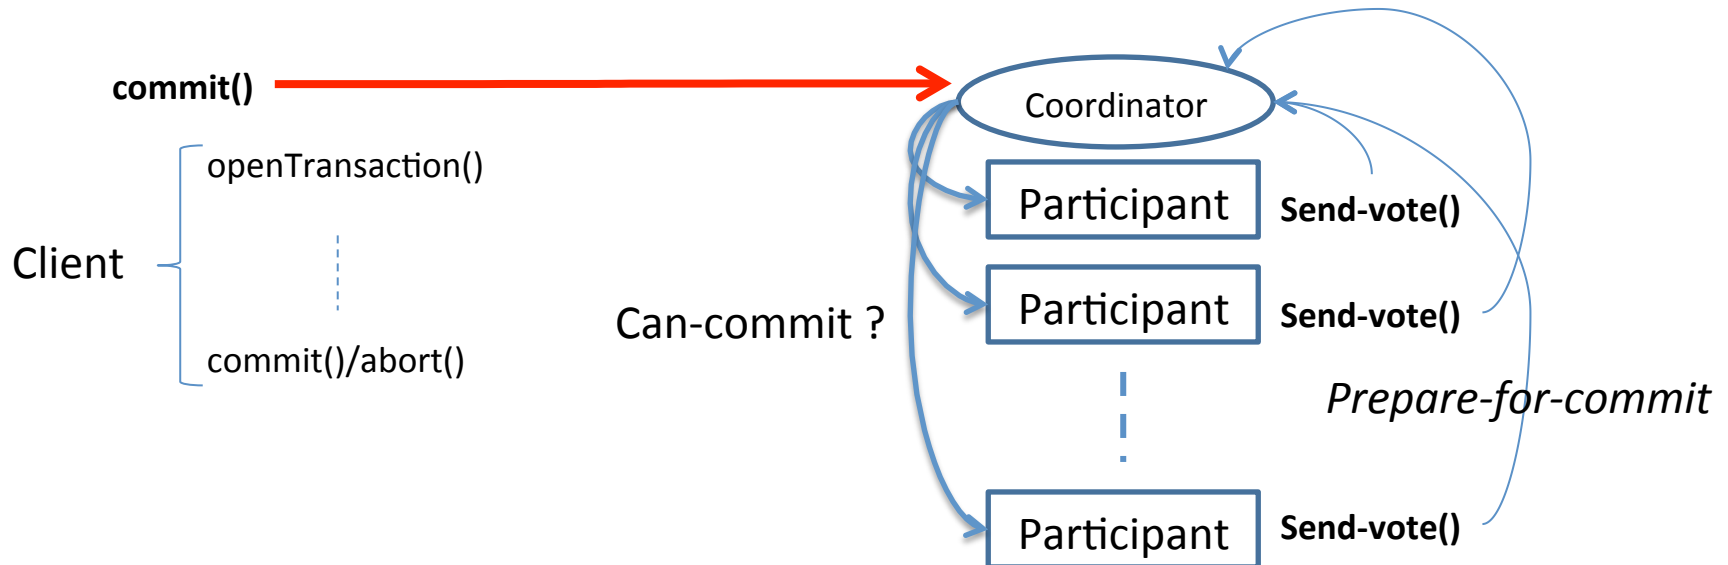
- If a participant votes to commit, it must make sure that it can commit eventually, even if it crashes – participant must be able to ***recover from system failure*** by storing intermediate state

2-Phase Commit Protocol



Two-phase Commit Protocol


Phase 1: Voting Phase



- Voting
 - Coordinator asks participants whether they can commit
 - Participants send their vote: either “Yes” or “No”
- Participants have to prepare for commit:
 - Write log files to record which write operations to commit

Phase 1: Prepare to Commit

- Each server, coordinator and participant, must be able to ***recover from system failure*** by storing intermediate results
 - Intermediate results stored in persistent storage
 - Log files record what objects are manipulated , locked etc.
 - These log files are used when a server is restarted and recovers from failure
- There may be cases where all other servers have to wait until a crashed server recovers

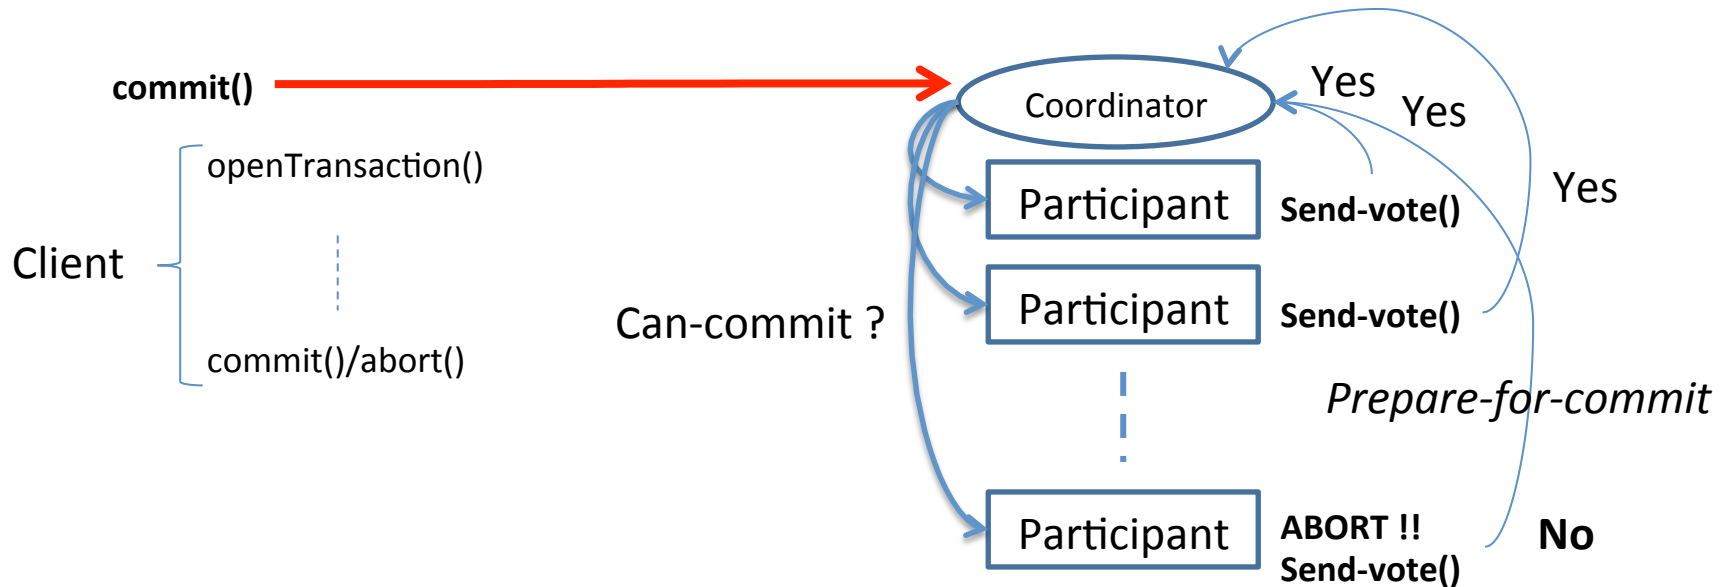


ACID Principle: Durability

- Guaranteed Commit
 - *If a participant voted for a commit of the transaction then it must guarantee that it will eventually commit its part of the transaction*
- This must be guaranteed for system failure
 - If participant recovers from a system failure, it must be able to commit its transaction
- Principle of Durability !

Two-phase Commit Protocol

Phase 1: Voting Phase



- Unilateral Abort in phase 1
 - A participant can **unilaterally** abort **any time** and send a “No” vote
 - Rollback of write operations, release of locks
- If there is at least one “No” vote than coordinator has to request abort from all other participants in phase 2

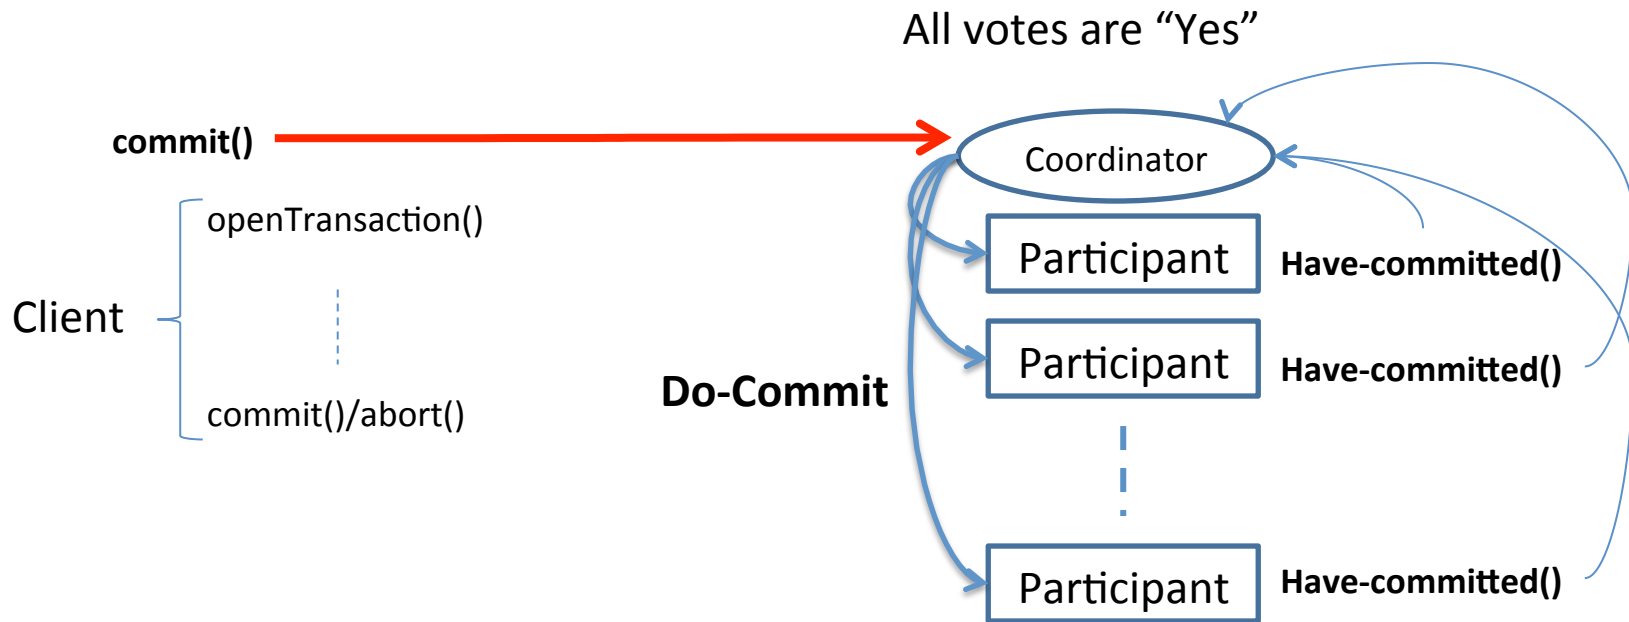
Two-phase Commit Protocol

Phase 1: Voting Phase

- ***Request for votes:***
 - Coordinator asks participants whether they are prepared to commit or abort
- ***Participant prepares for Commit***
 - If a participant sends a vote to commit its part of a transaction, it will prepare for commit by recording manipulations on data object in permanent storage and marks the transaction as *prepared*
 - A participant is, therefore, said to be in a *prepared* state for a transaction commit.
 - all participants vote whether to commit, once they voted to commit, they are not allowed to abort any more
 - When a participant is prepared for Commit, it has to wait for a Coordinator Decision – **blocked until coordinator decides**
- ***Participant aborts***
 - If a participant sends a vote to abort its part of a transaction, it will unilaterally abort (immediate abort, release locks)

Two-phase Commit Protocol

Phase 2: Completion Phase



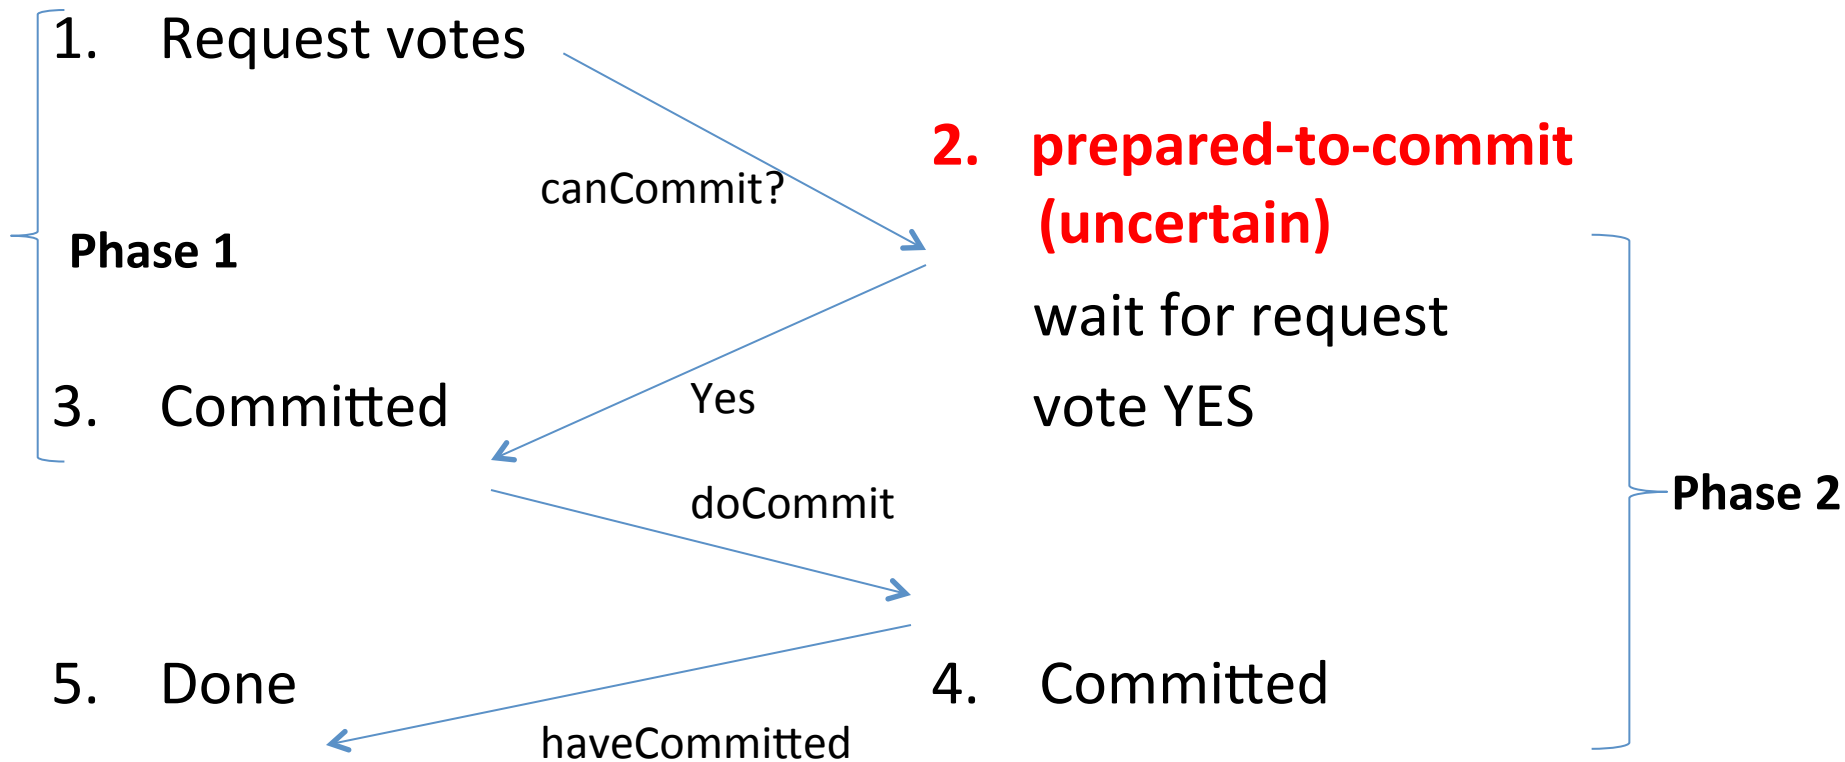
- Completion Phase – all votes “Yes”
 - Coordinator first asks participants to commit
 - Participants then send message to coordinator about success of commit

2PC Protocol

All participants vote for commit, Coordinator commits

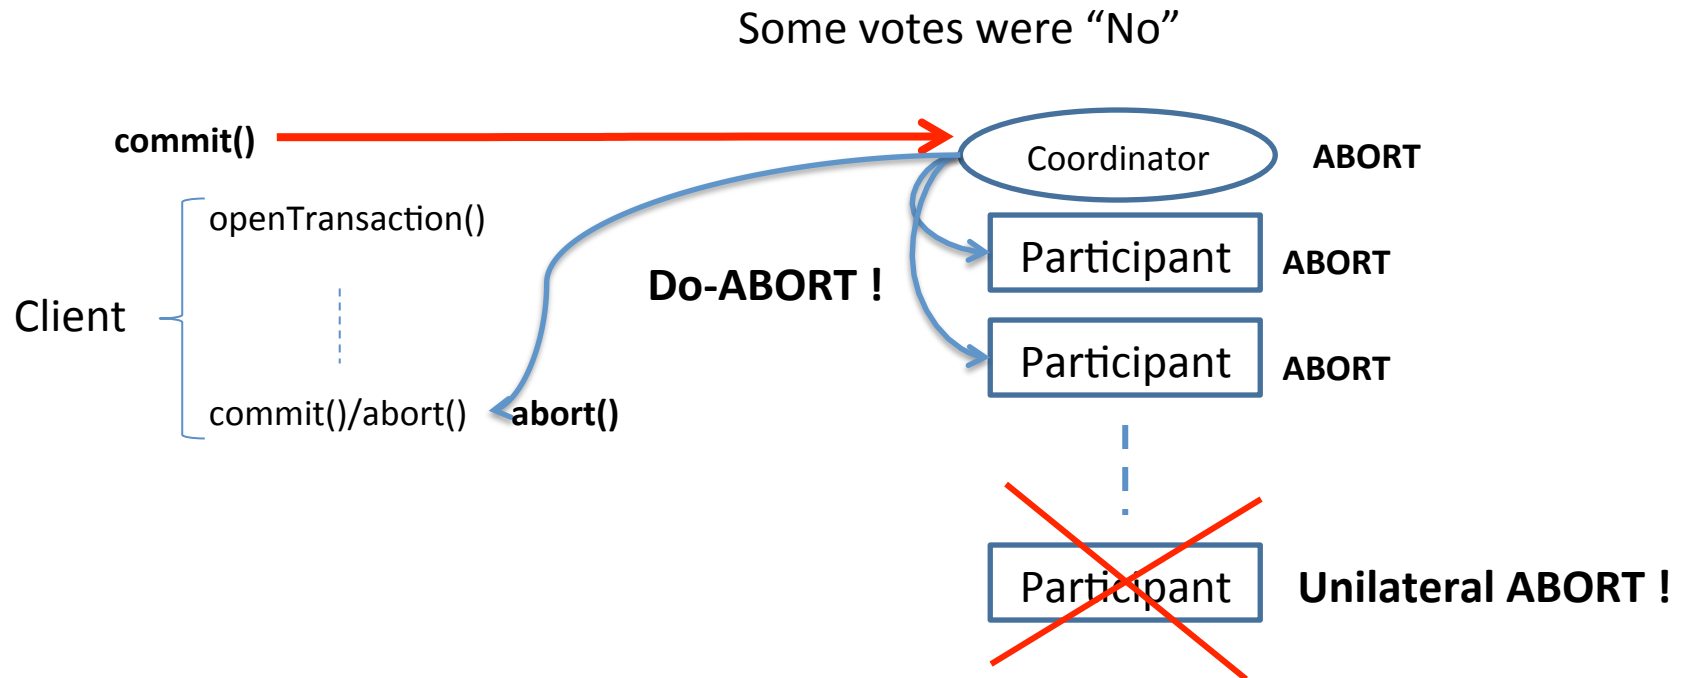
- Coordinator

- Participant



Two-phase Commit Protocol

Phase 2: Completion Phase



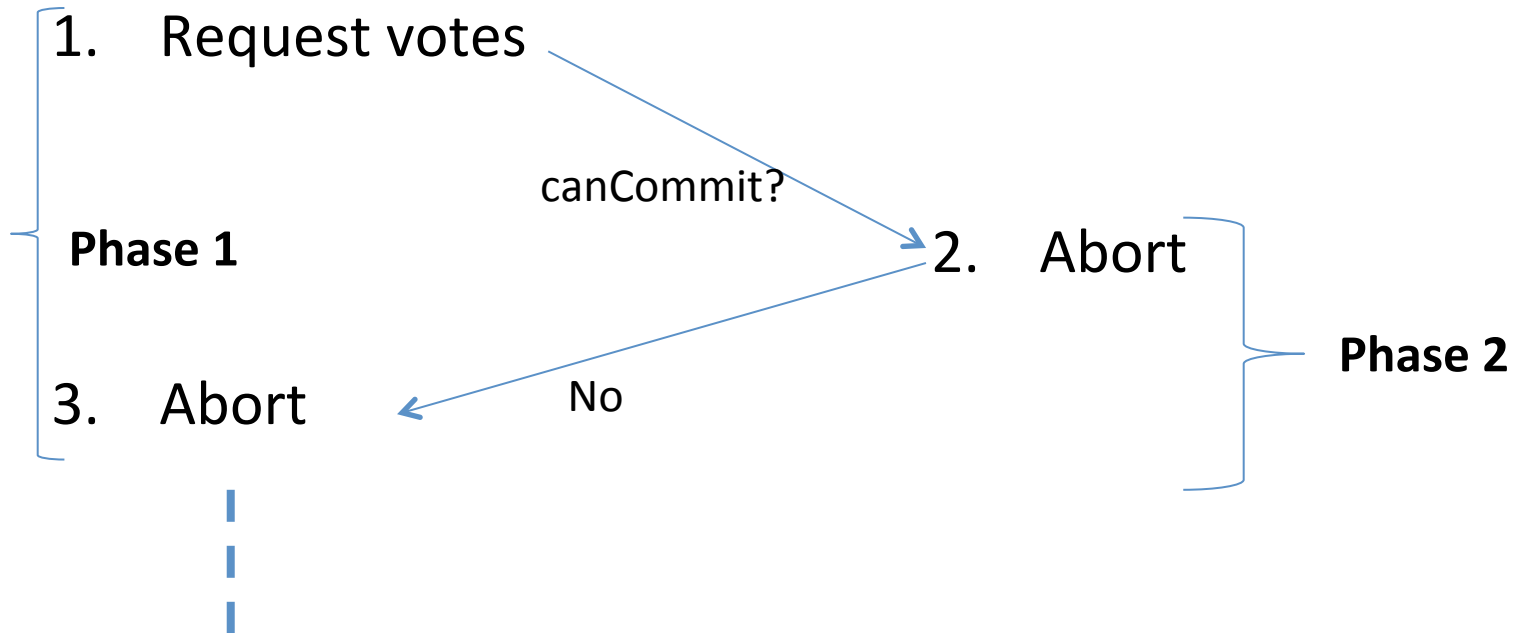
- Completion Phase – some votes are "No"
 - Coordinator asks participants to ABORT
 - No further communication between participants and coordinator
- Participants that voted "No" can unilaterally abort, without communicating with Coordinator

2PC Protocol

Participant votes Abort, unilaterally aborts

- Coordinator

- Participant




Coordinator has to abort all other participants.

Two-phase Commit Protocol

Phase 2: Completion Phase

All votes are “Yes” in phase 1



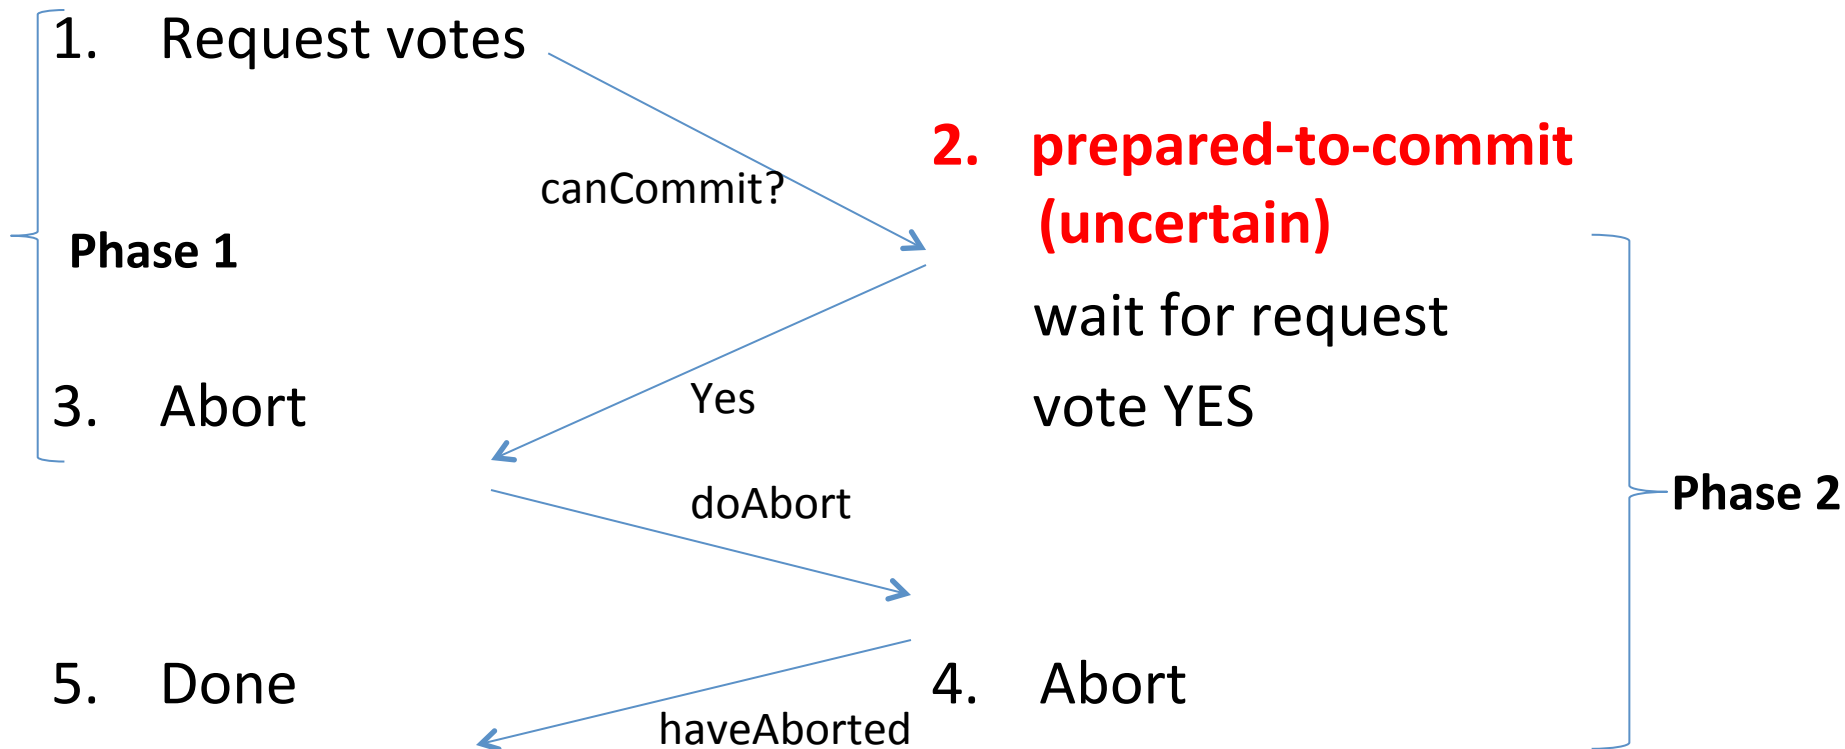
- Completion Phase – all votes “Yes”
 - Coordinator may still ask participants to abort
- Unilateral Abort:
 - In case of Abort, there is no communication between participants and coordinator

2PC Protocol

Participant votes for commit, Coordinator aborts

- Coordinator

- Participant



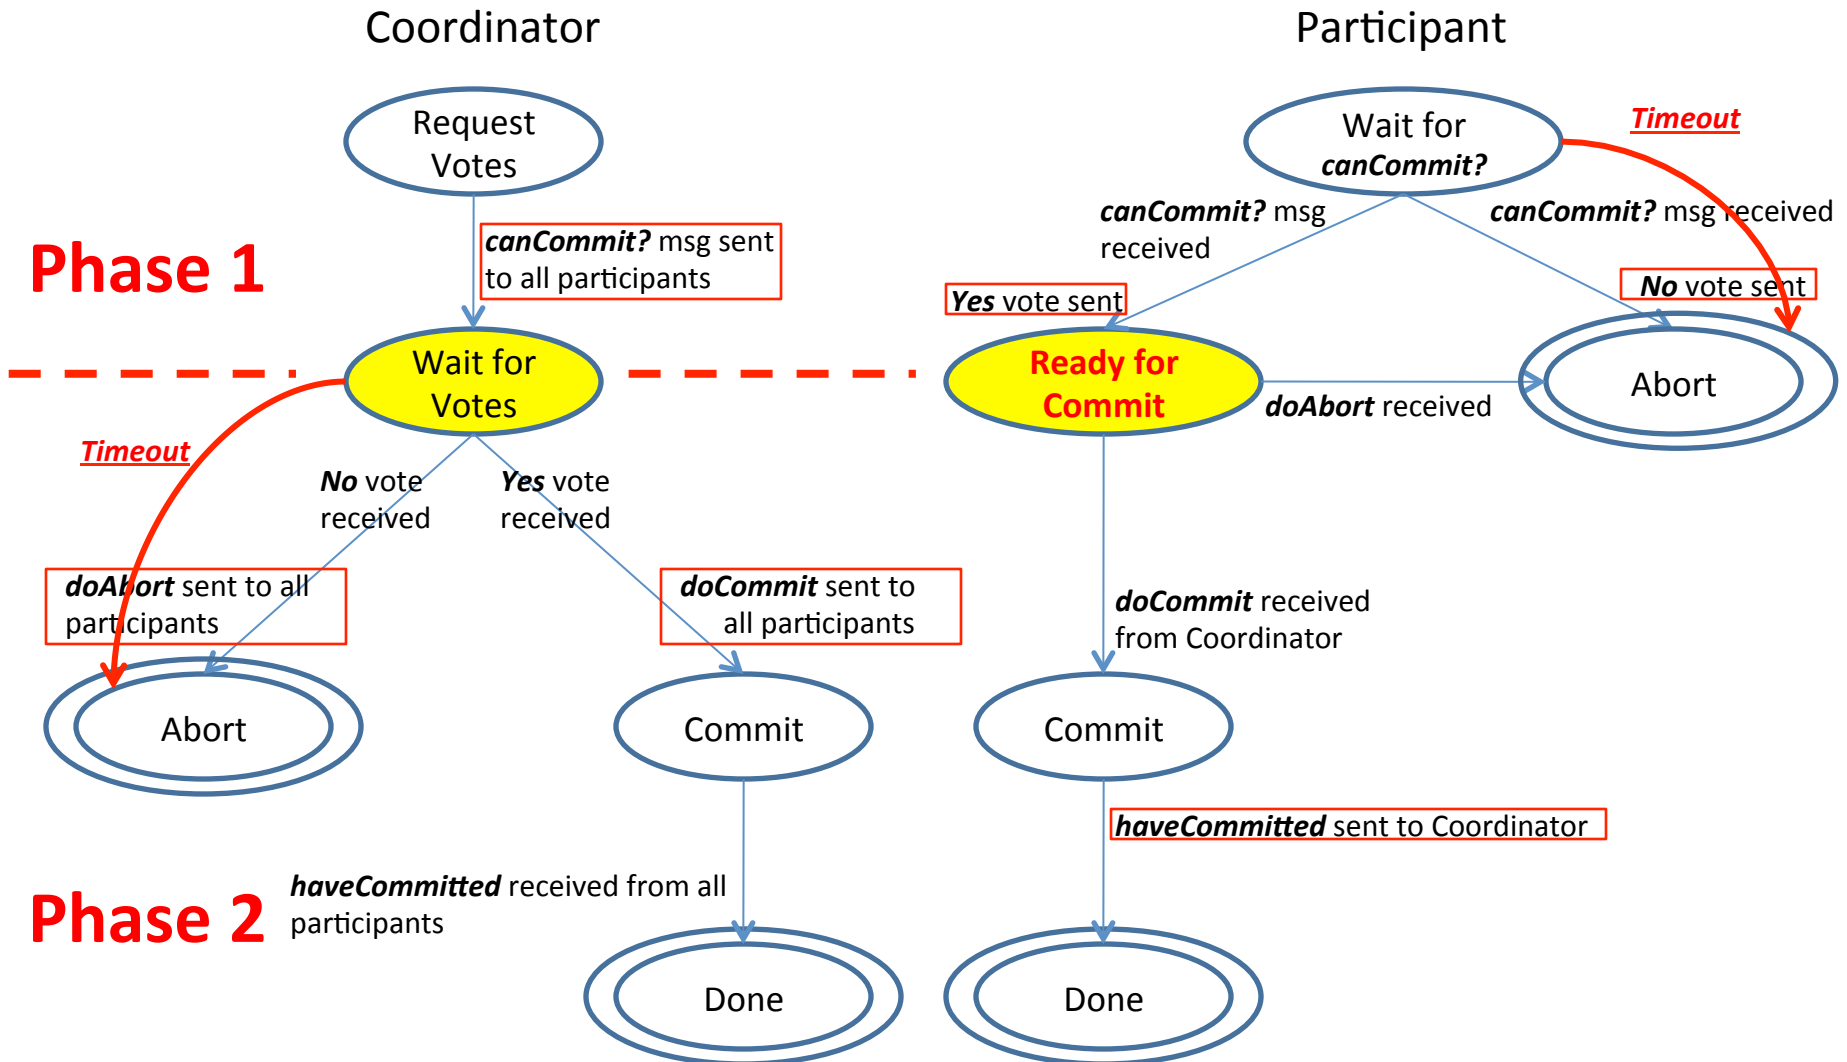
Coordinator received a **No** vote and sends a **doAbort** to all other participants.

Two-phase Commit Protocol

Phase 2: Completion Phase

- Coordinator collects votes from participants and informs participants about its decision
 - If all votes are commit-votes:
 - the coordinator informs all participants to commit their local transaction
 - If not all votes are commit-votes:
 - (This may be because at least one participant sends abort **or is not responding within a time limit**)
 - the coordinator informs all those participants that voted for commit, to abort their local transaction
- Participants that voted for abort, immediately abort their local part of the distributed transaction, without waiting for coordinator

2-Phase Commit Protocol



2PC Protocol Implementation

- Phase 1: The Voting Phase
 - The coordinator sends a **canCommit?** Request to each participant in the transaction
 - When a participant receives a **canCommit?** request, it decides whether to vote **Yes** or **No**
 - If it decides to vote **Yes**, the participant *prepares* for commit by saving data objects plus its current state in persistent storage (necessary for recovery)
 - If it decides to vote **No**, it aborts immediately
 - The participant sends a vote: **Yes** or **No**
- Phase 2: The Completion Phase
 - Coordinator collects the votes including its own
 - If there are no failures and all votes are **Yes**, the coordinator decides to commit and sends a **doCommit** instruction to all participants
 - Otherwise the coordinator sends a **doAbort**
 - Participants that voted **Yes** are waiting for a **doCommit** or **doAbort**. If the instruction is **doCommit** they act accordingly and send a **haveCommitted** confirmation to the coordinator

Operations of the 2PC Commit Protocol

- `canCommit? (Tid)` --> *Yes / No*
 - Call from coordinator to participant
 - Asks whether it can commit a transaction
 - Participant replies with its vote
- `doCommit (Tid)`
 - Call from coordinator to participant
 - Tells participant to commit its part of a transaction
- `doAbort (Tid)`
 - Call from coordinator to participant to tell participant to abort its part of a transaction
- `haveCommitted (Tid, Participant)`
 - Call from participant to coordinator
 - Confirms that it has committed the transaction
- `getDecision (Tid)` --> *Commit / Abort*
 - Is used to recover from server crash or delayed messages
 - Call from participant to coordinator, if there is no reply from the coordinator after some delay
 - Participant asks for the decision on a transaction after it has voted *Yes*

Participant Interface

- `Vote = canCommit (Tid)`
 - Returns the vote of the participant to the coordinator
 - Asks whether it can commit a transaction
- `doCommit (Tid)`
 - Coordinator tells participant to commit its part of a transaction
- `doAbort (Tid)`
 - Call from coordinator to participant to tell participant to abort its part of a transaction


Handling Failure Situations

Handling Failure Situations

Robustness of 2PC

- 2PC works in failure situations because
 - participants save their state (manipulations on data) in permanent storage as a *preparation* for commit
 - This preparation enables recovery in case of system failure
- If a participant recovers from a crash, it can continue from this saved state and complete the interaction with the coordinator
- Participant may retrieve last vote on transaction from Coordinator
 - Participant sends getDecision message to Coordinator

Participant asks Coordinator



- After recovery, a participant has to ask Coordinator for its decision
 - Participant sends a “getDecision” message to Coordinator
 - Participant will act according to Coordinator decision

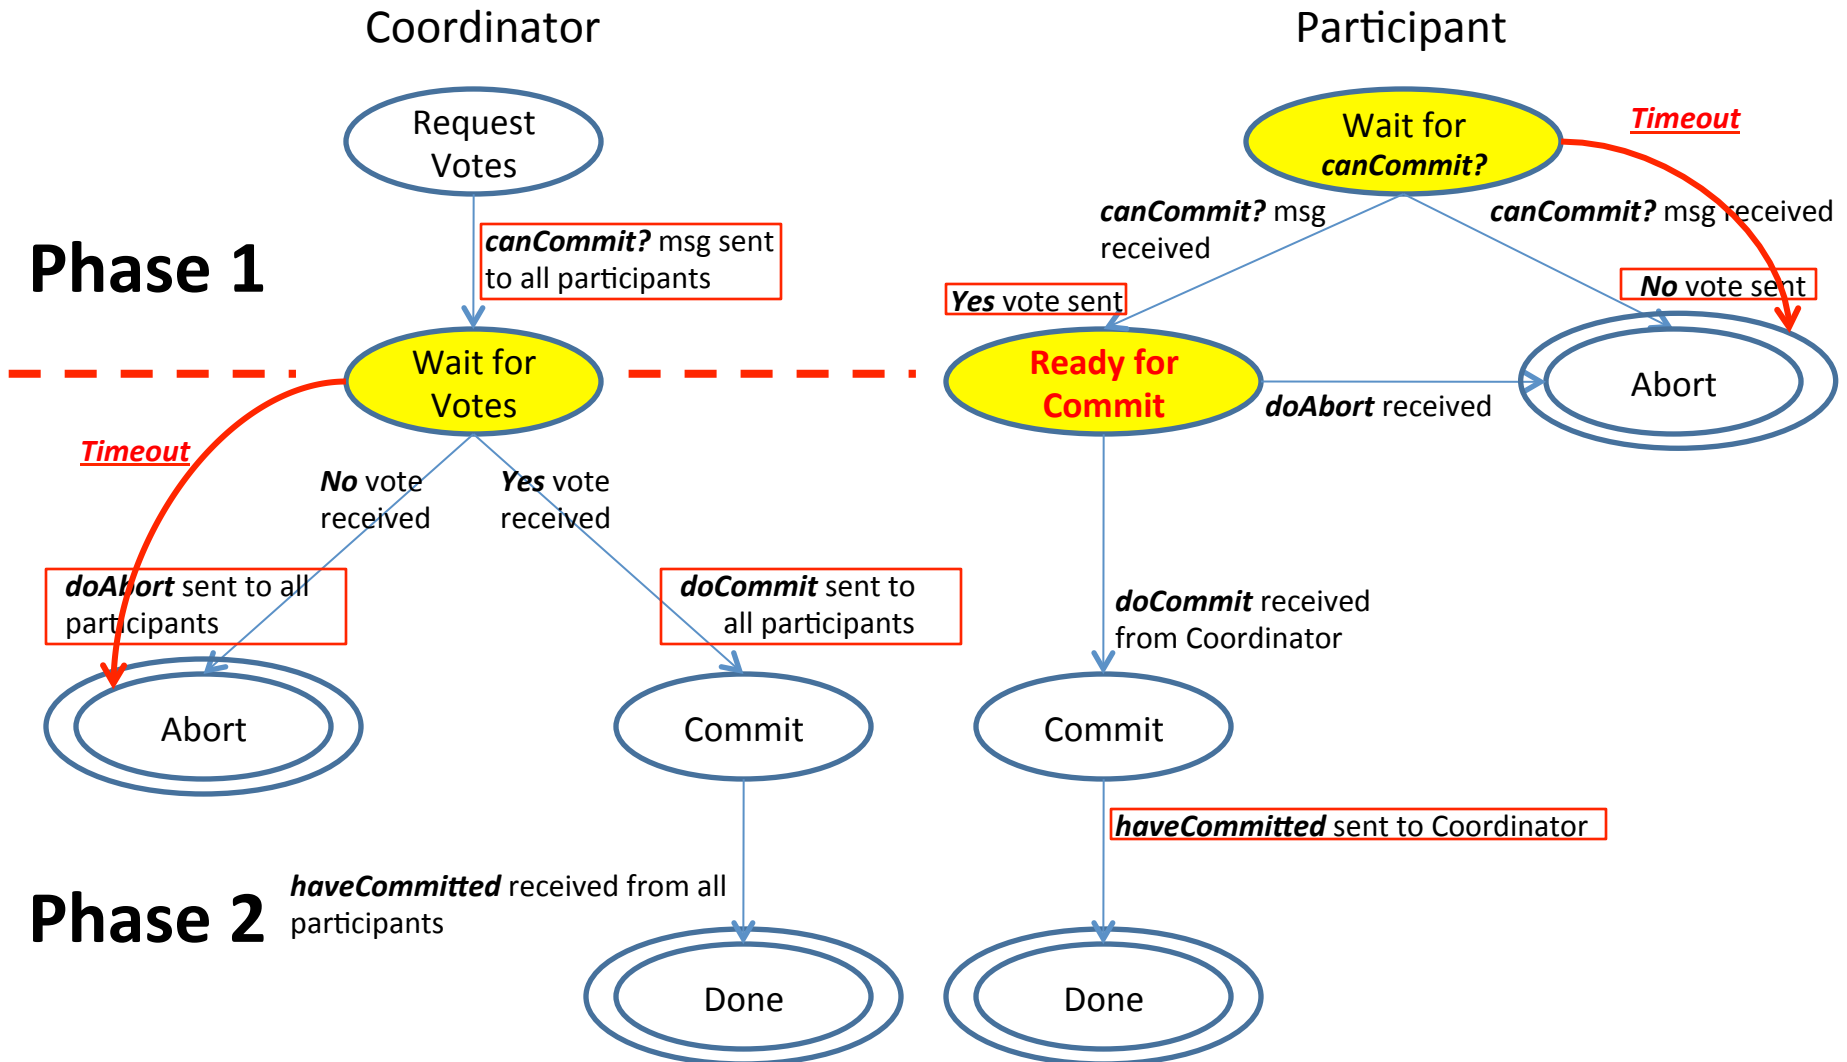
Two-phase Commit Protocol

- As long as Coordinator is alive, distributed transactions can be aborted fast and restarted
 - Coordinator will abort distributed transaction in case of failure
 - When the failed participant recovers, it may ask coordinator about decision (which was abort)
- **HOWEVER: What if the coordinator crashes ?**

Problem in System Failure Situations

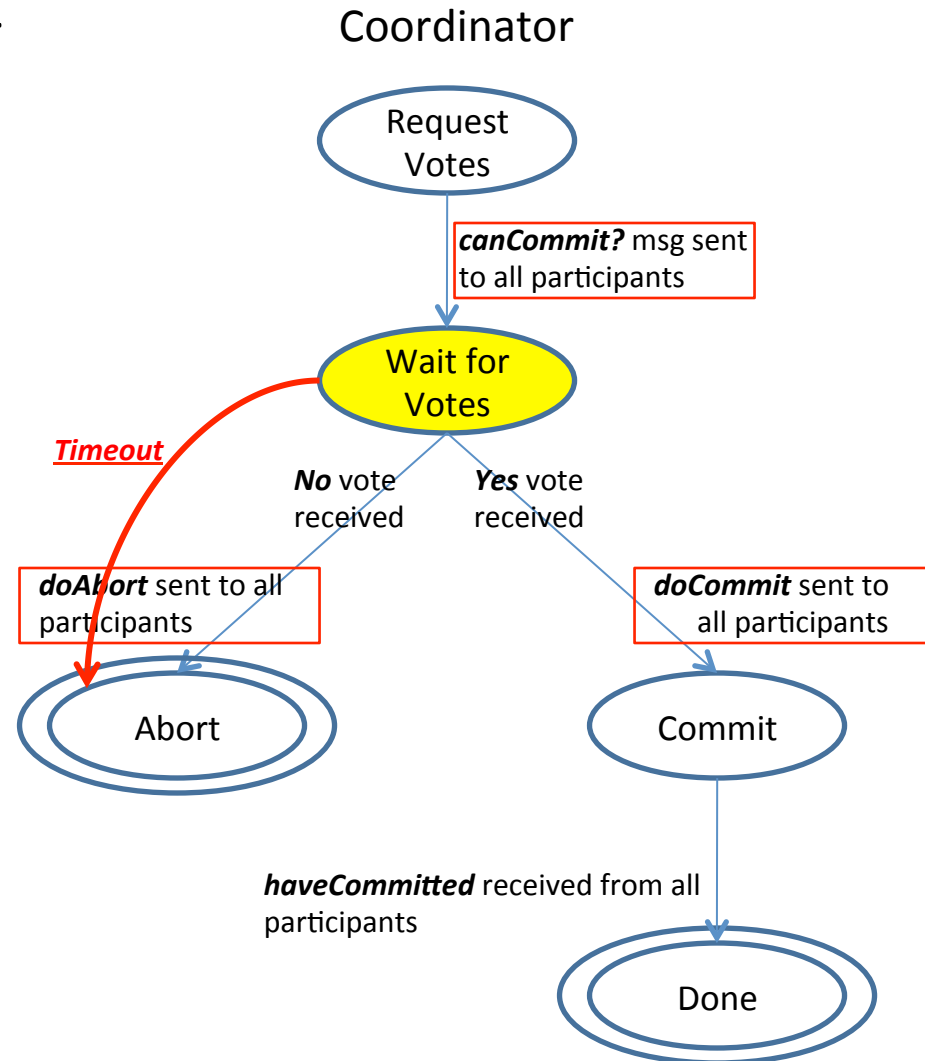
- 2PC can become a *blocking* protocol if coordinator fails in phase 2:
 - 2PC can cause considerable delays to participants in an *uncertain* state, this occurs when the coordinator fails and cannot reply to `getDecision()` requests
 - Participants have to wait, locks on data objects remain in place as the transaction cannot finish

Timeout 2-Phase Commit Protocol



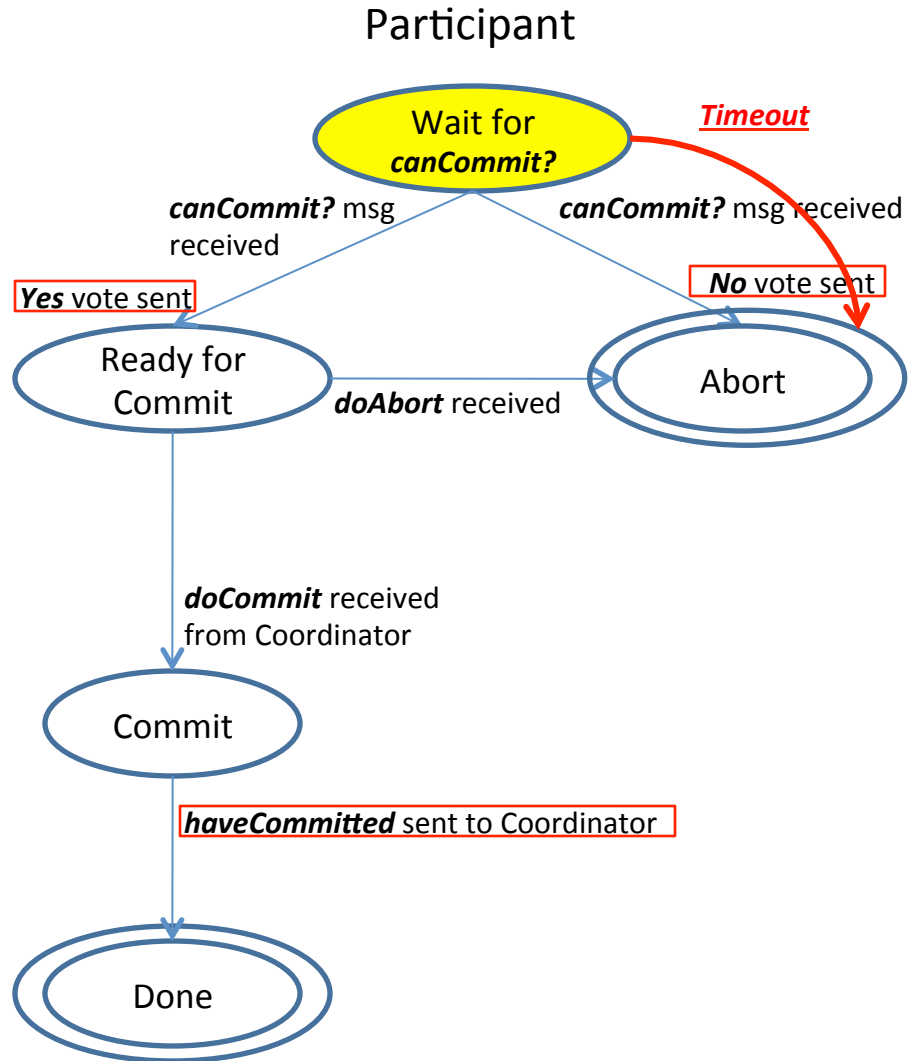
Timeout Actions in the 2PC Protocol

- Coordinator time-out in phase 1
 - No vote received from at least one participant after timeout
 - This is regarded as an indication to abort
 - Coordinator decides to abort the transaction
 - Coordinator sends “doAbort” to participants and aborts unilaterally
 - Coordinator will ignore subsequent votes
- If failed participant recovers it has to ask coordinator for its decision and will abort as well



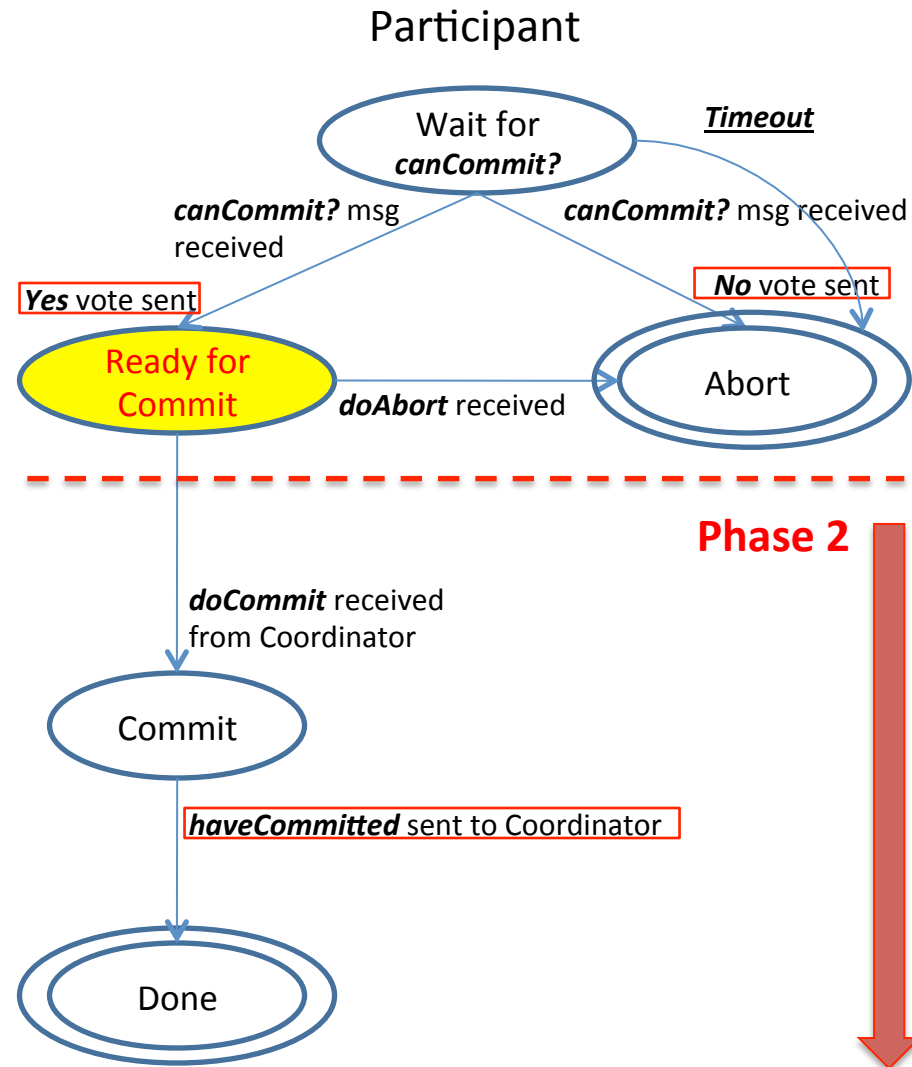
Timeout Actions in the 2PC Protocol

- Participant time-out in phase 1
 - Participant finishes its part of the distributed transaction
 - Coordinator is late in sending a request-for-vote (phase 1):
 - the participant reaches its timeout and will unilaterally abort



Blocking Situation in Phase 2

- Participant waits for coordinator
 - Participant has sent vote to commit, is ready for commit
 - As it has voted for commit, no locks on shared objects can be released, the participant has to wait for coordinator decision
 - Coordinator is late in sending a doCommit or doAbort (phase 2)
- Participant cannot abort unilaterally, it has to wait for the coordinator's decision !!**
- How to find out the coordinator's decision, if coordinator has crashed?



Strategy: Ask for Decision

- Situation: Coordinator is late in sending a doCommit or doAbort (phase 2)
 - Participant is prepared for commit (in phase 2), but *uncertain* about the outcome of the voting – it cannot proceed, therefore objects remain locked
- Strategy: Ask for coordinator's decision after a time delay:
 - call the coordinator's getDecision() method to get information about the vote
 - If there is a reply from the coordinator, it can finally commit and call the coordinator's hasCommitted() method

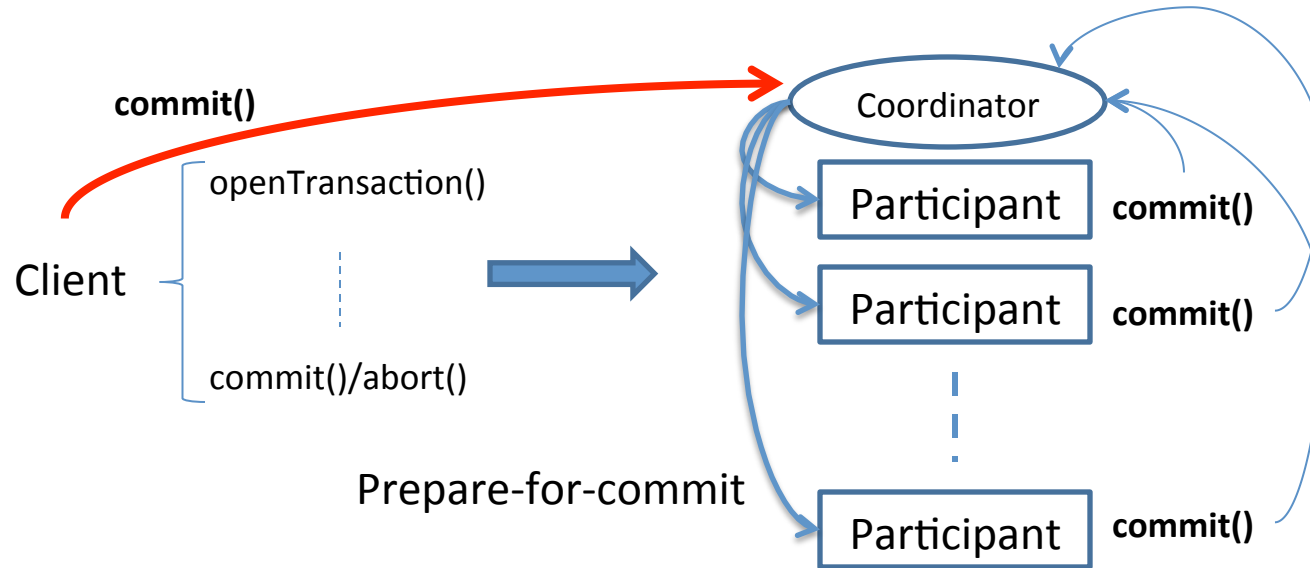
Strategy: Elect a new Coordinator

- Situation: Coordinator is late in sending a doCommit or doAbort (phase 2)
 - Participant is prepared for commit (in phase 2), but *uncertain* about the outcome of the voting – it cannot proceed, therefore objects remain locked
- What if the coordinator has failed after sending a ***request for vote*** message?
 - Participants cannot contact the coordinator, must wait until it is recovered/replaced
 - In the meantime, all locks on data objects involved in the distributed transaction cannot be released
- Participants may elect a new coordinator among them that sends abort messages to the remaining participants

Problem in System Failure Situations

- 2PC is a *blocking* protocol in failure situations:
 - 2PC can cause considerable delays to participants in an *uncertain* state, this occurs when the coordinator fails and cannot reply to getDecision() requests
 - Participants have to wait, locks on data objects remain in place as the transaction cannot finish
- Three-phase commit protocols have been designed to prevent delays due to coordinator or participant failure, but the cost is higher in the failure-free case (more messages required).

Commit Protocol, Requirements




- If the client requests a **commit**:
 - The distributed transaction can only be committed, if all local sub-transactions executed on servers finished successfully
- **All participants have to be informed** that a **commit** is imminent, as they have to prepare for commit
 - Write log information to enable recovery from possible system crashes
- The coordinator can declare the transaction committed only if it gets a **reply from each participant** that their local sub-transactions have successfully finished

Commit Protocol, Requirements

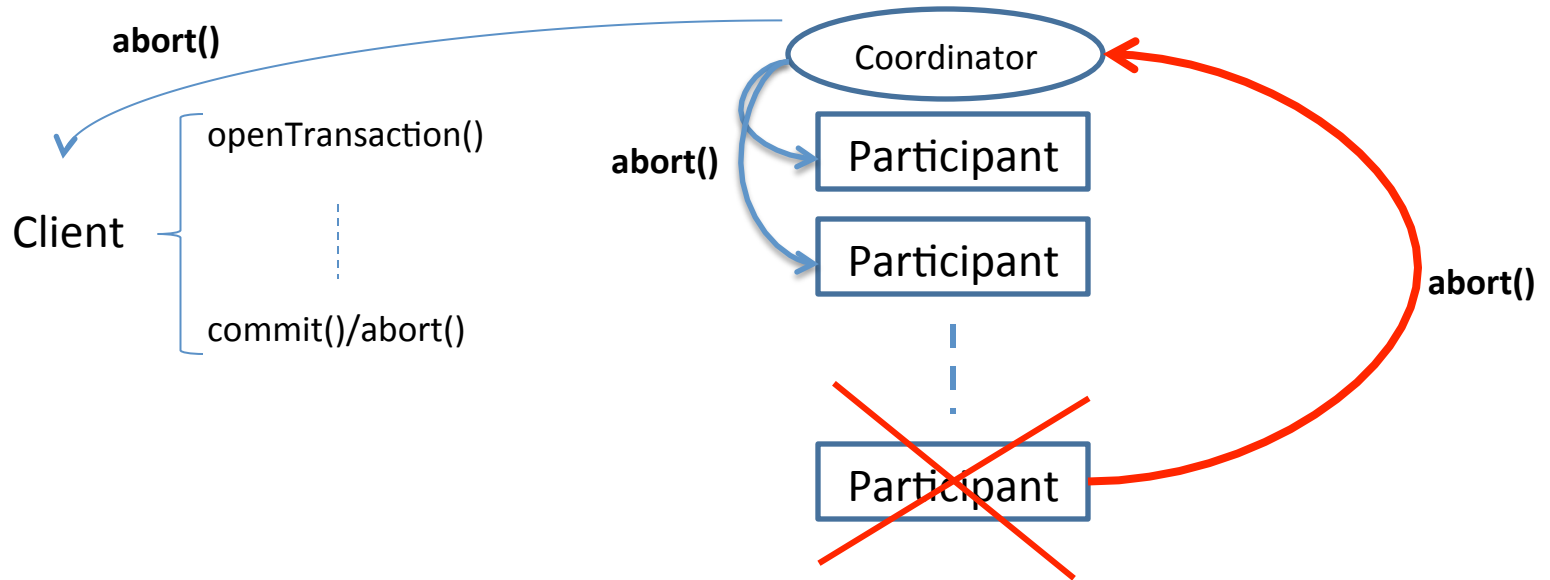
- If the client requests a **commit**:
 - The distributed transaction can only be committed, if all local sub-transactions executed on servers finished successfully
- Coordinator Action:
 - **All participants have to be informed** that a **commit** is imminent, as they have to prepare for commit
 - Write log information to enable recovery from possible system crashes
- Participant Action:
 - The coordinator can declare the transaction committed only if it gets a **reply from all participants** that their local sub-transactions have successfully finished
 - For this, they have to communicate the state of their local portion of the distributed transaction to the coordinator

Commit Protocol, Requirements



- If the client requests an **abort**:
 - All participants have to be informed that an abort is imminent
 - they have to release local locks and
 - undo/rollback any manipulation on data objects
 - Do they have to report back to Coordinator ?

Commit Protocol, Requirements



- Participant
- If a *participant* has to **abort** its portion of the distributed transaction
 - the coordinator has to abort the overall distributed transaction, all participants involved in that transaction are told to abort (rollback of complete transaction)
- How does Coordinator know that all transactions are aborted?