

Revision Distributed Systems

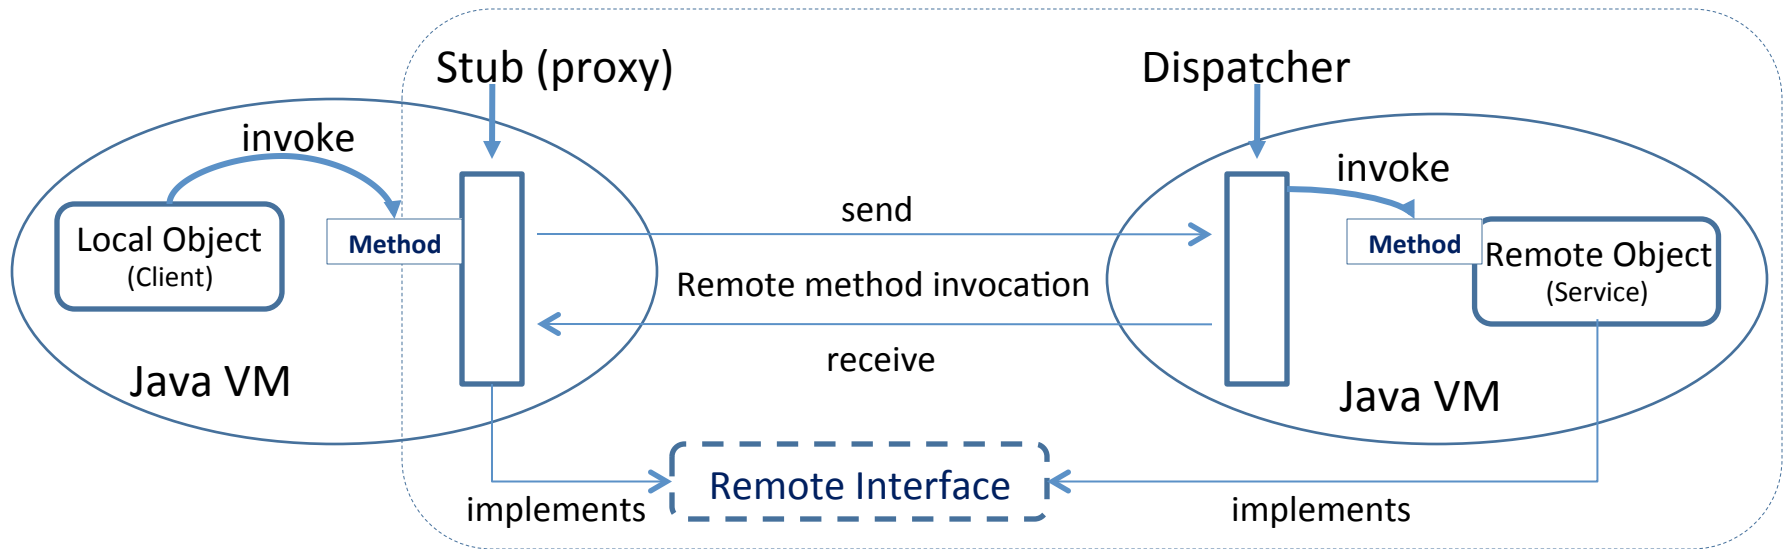
CS3524 Distributed Systems

Lecture 23

RMI

RMI Implementation

Distributed Object Model



- RMI middle-ware:
 - Stub
 - Dispatcher (uses Java reflection)

Serialization

- How to pass information between client and server?
 - So far, we have passed instances of `java.lang.String` from client to server
- How to pass more complex, programmer-defined data structures?
 - We need mechanisms to convert such data structures into and from byte streams – these byte streams can then be passed between client and server
- Java provides a simple means for object serialisation:
 - The interface **`java.io.Serializable`**
- If Java classes implement this interface, then objects instantiated from these classes can be translated into a byte stream that is sent across a network
- This byte stream is sufficient for the reconstruction of the serialized object when it is read by the receiving program

Callback

- Implementation in RMI:
 - We define a remote interface that contains a method for the server to call remotely and an implementation for it - these are the definitions for a “callback object”
 - Client instantiates, exports and registers a remote object that implements this interface – this is the “callback object”
 - Server provides a method in its remote interface that clients can use to inform the server of the remote object reference of their callback object – the server will record this reference
 - Whenever an event of interest for a client occurs, the server “calls the interested client” – it will call a remote method on the callback object

RMI Activation Service

- The ability to automatically generate remote objects, triggered by requests for references to these objects
- Robustness:
 - The distributed system is constructed in a way to handle and recover from disturbances and system failure, without loss of information, e.g. by
 - Restarting remote objects that were destroyed due to a system failure
- Scalability
 - A distributed system should be able to absorb high system load, increased/decreased service requests , e.g. by
 - Starting services/remote objects on demand

Threads

Threads

- A thread is a single sequential flow of control within a program
- Threads exist within a process (e.g. the Java Virtual Machine)
- Many threads may execute concurrently within such a process
- All threads share resources owned by this execution environment
- Multi-threaded execution is an essential feature of Java

Critical Sections

- Threads may execute concurrently and independently without interfering with other threads
- But: Threads (in most cases) interfere with each other
 - they access shared resources concurrently
 - A particular section of code specified in a Java class may be executed by multiple threads concurrently – we call this a *critical section*
 - Critical sections of code have to be save-guarded in a special way – only one thread at a time should be able to execute such a section **in its completeness**
- Goal: **Mutual Exclusion** between threads – only one thread at a time executes a critical section in its completeness

Race Condition

- A ***race condition*** is an undesirable situation that occurs when two or more threads attempt to manipulate a shared resource concurrently
 - Two or more threads are able to access shared data and they try to make changes at the same time
- The result of such a computation depends on the sequence of how these threads are scheduled over time
 - A different schedule may result in a completely different interleaving of actions of the different threads and in different results
- “the threads are ‘racing’ against each other to win access to the shared resource”
 - E.g.: multiple read and write operations by different threads may result in one thread overwriting / removing a computation of another thread.

Monitor

- We need a mechanism to *coordinate and synchronise* thread activity. Java provides simple synchronisation mechanisms based on C. A. R. Hoare's widely-used *monitor* concept
- A monitor can be thought of as a *barrier* securing a *shared resource* with a **lock** :
 - If the resource is not used by another thread, a thread can *acquire* the lock and access the resource – the thread is regarded as “entering the monitor”
 - Other threads wanting to access the resource must wait (in a queue) until the lock is *released*

Critical Section in Java

- In Java, *critical sections* of code can be marked with the **synchronized** keyword to provide mutually exclusive access to an object
- The synchronized keyword indicates where a thread must acquire the object's lock
 - When a thread is scheduled for execution and enters a section of code that is enclosed with a **synchronized** construct, it has to acquire first a lock on this object
 - If another thread already holds this lock, our currently scheduled thread has to wait
- Please note: the whole object is locked for a particular thread!

Observer Pattern

- Class `java.lang.Object` provides the following methods (from the Observer Pattern):
 - `wait()`
 - `notify()`
 - `notifyAll()`
- These methods support an efficient transfer of control from one thread to other threads
- These methods can be used to handle the competition of threads for entering a monitor / shared object with a lock

Deadlock

- Necessary conditions for a deadlock to occur (“Coffman Conditions”):
 - Mutual exclusion:
 - At least one resource must be non-shareable – only one thread / process at a time may use this resource
 - “Hold and wait”:
 - A thread / process is currently holding at least one non-shareable resource and is waiting for another non-shareable resource held by another thread / process
 - No preemption:
 - The operating system cannot simply free the resource, it must be given up by the holding thread / process voluntarily
 - Circular wait:
 - Two threads / processes: one thread waits for the release of a resource that is held by a second thread, which, in turn, waits for the release of a resource held by the first thread
 - N threads / processes: given n threads, the first thread waits for the second thread to release a resource, the second waits for the third to release a resource, etc., and the nth thread waits for the first thread to release a resource – circular wait.

Further Problems

- Livelock:
 - If thread A acts in response to the actions of thread B and thread B immediately reacts to thread A's action, they may be locked into an endless loop of mutually reacting to each others' actions, always performing the same set of actions
- Starvation:
 - can occur when threads can have different priorities
 - A thread is never scheduled for execution, because the scheduler always schedules threads with higher priority
 - Scheduling algorithm has to take this problem into account

Transactions

Transactions

- Transactions encapsulate a series of data manipulation operations (read/write)
- Transactions have a defined start where they are “opened”
- When a transaction is opened, all operations on data are regarded as occurring “preliminary”
- Transactions have a defined end with two possible outcomes
 - “**commit**”: a successful termination of a transaction is regarded as a “commit” – all operations on data that occurred since the start of the transaction are *committed*
 - “**abort**”: an unsuccessful termination of a transaction is regarded as an “abort” – all manipulation operations on data are *undone* or *rolled back*
- A transaction causes the system to move from one consistent state, at the application level, to another consistent state

The ACID Properties

- **Atomicity:** Either all of the transaction's operations are performed, or none of them
- **Consistency:** A transaction transforms the system from one consistent state into another
- **Isolation:** An incomplete transaction cannot reveal its intermediate state to other transactions until committed
- **Durability:** Once a transaction is committed, the system must guarantee that the results of the transaction will persist, even if the management system subsequently fails.

Serial Equivalence


- Serial Equivalence
 - Two or more transactions produce the same result operating in an interleaved fashion, as if they would operate in a completely serialized fashion (each transaction is completed before the next one gains access to shared data objects)
- Serial Equivalence is used as a criterion for establishing concurrency control protocols
 - A transaction manager has to schedule operations of different transactions in a way so that interference and problems such as “lost updates” are avoided

Simple Locking

- If a transaction wants to read / write objects, it establishes locks for these objects
 - The server sets a lock on each object for a particular transaction before it is accessed (read or write operations)
 - The server removes these locks when the transaction is finished and not immediately after the operation – exclusive locking
- While an object is locked
 - only the transaction holding the lock can access/manipulate this object
 - Other transaction have to wait for the lock to be released or may be able to share a lock (this depends on the locking concepts used)
- The use of locks can lead to deadlocks !

Strict Two-Phase Locking

- No alternating lock /unlock during a transaction – transactions only set/promote locks during their operation and release locks at commit / abort



- When an operation accesses an object within a transaction:
 - a) If the object is not already locked, it is locked and the operation proceeds
 - b) If the object has a conflicting lock obtained by some other process it waits until it is unlocked
 - c) If the object has a non-conflicting lock set by another transaction, it shares the lock and the operation proceeds
 - d) If the object has already been locked in the transaction, the lock will be promoted if necessary and the operation proceeds (if promotion is prevented by conflict, rule b) is used)
- When a transaction is committed or aborted, the server unlocks all objects it locked for the transaction

Simple Exclusive Lock

- This simple strategy guarantees serialised transactions
 - The order of all conflicting pairs of operations is the same, because a conflicting transaction cannot proceed while it does not hold locks on objects it needs
- Important!
 - Strict two-phase locking – a transaction cannot set and release locks arbitrarily, there has to be (1) a phase of setting locks and then (2) a phase of releasing all locks, all locks are held until the release-phase at commit/abort
- In the example
 - Transaction U cannot proceed as long as Transaction T holds a lock for object b
- Problem
 - This is not the most efficient strategy; it unnecessarily restricts access to shared resources
- For example, two transactions that simply wish to perform a read() operation on an object would not interfere – allow shared read !

Shared *read* Locks

- Concurrency of transactions can be improved by distinguishing between read and write locks
 - Before a read() is performed, a *read* lock is set on an object
 - Before a write() is performed, a *write* lock is set
- Conflict rules
 - If transaction T has set a *read* lock then transaction U can set a *read* lock as well
 - If transaction T has already set a *read* lock then transaction U is **not** allowed to set a *write* lock until T commits / aborts
 - If transaction T has already set a *write* lock then transaction U is **not** allowed to set either *read* or *write* locks (***exclusive lock***)
- *read* locks are also called *shared* locks, as many transactions can set a read lock (or “share” it) on an object at the same time
- *read* locks guarantee that the object remains readable but no other transaction can set a *write* lock and make inconsistent updates

Lock Compatibility

<i>For one object</i>	<i>Lock requested for second object</i>	
	<i>Read</i>	<i>Write</i>
<i>Lock already set</i> None	OK	OK
Read	OK	Wait
Write	Wait	Wait

- The fact that a lock has been set on an object by transaction 1 does not necessarily mean that transaction 2 cannot also obtain a lock
- The lock may be *shared* if they are both *read* locks
- Lock Promotion – make a lock more exclusive
 - A read lock can be “promoted” to a write lock if it is **not** shared by other transactions

Two-Version Locking

		Lock requested		
		read	write	Commit
Lock already set	none	OK	OK	OK
	read	OK	OK	wait
	write	OK	wait	---
	commit	wait	wait	---


- Uses three locks:
 - Read, write, commit
 - A transaction can set a read lock if the object is unlocked or read locked
 - If the object has a commit lock, the transaction has to wait
 - A transaction can set a write lock, if the object is unlocked or read locked
 - If the object has a write or commit lock, the transaction has to wait
- At commit of a transaction, the transaction manager tries to convert the write locks into commit locks – if there are any outstanding read locks on the corresponding objects, the transaction has to wait

Deadlocks

- The use of locks may lead to deadlocks
- Definiton
 - A deadlock is a state in which each member of a group of transactions is waiting for some other member to release a lock
- The scheduler is responsible for detecting and breaking the deadlock
 - Deadlocks may be broken by simply aborting one of the transactions involved
- How do you choose which transaction to abort?
 - Abort the oldest
 - Abort depending on the complexity of the transactions
- Rather than detecting deadlock (an overhead), you could just use timeouts, but how long should the timeout be?

Wait-for Graph


- Representation of a deadlock situation:
 - Nodes in this graph represent transactions
 - Edges between nodes represent wait-for relationships between current transactions
 - The dependency between transactions is indirect – via a dependency on objects
 - A cycle in the graph indicates a deadlock:



- Transaction T and U wait for each other because of a lock on one object
- None of these locks can ever be released
- One of the transactions would have to be aborted to break this deadlock

Nested Transactions

- Nested Transactions have sub-transactions
- Sub-transactions can commit / abort independently
- A nested Transaction can only commit or abort if all its sub-transactions have completed
- When a sub-transaction completes, it makes an independent decision to provisionally commit or abort (abort is final)
- When a sub-transaction aborts, the parent transaction can decide whether to abort or not
- When the parent aborts, all sub transactions abort
- When the top level transaction commits, all transactions that committed provisionally, will commit finally



Committing Nested Transactions

- A transaction may commit or abort only after all its sub-transactions have completed
- When a sub-transaction completes, it makes an independent decision either to commit *provisionally* or to abort (abort is final)
- When a parent transaction aborts, all its sub-transactions are aborted, even those that have *provisionally* committed
- When a sub-transaction aborts, the parent can decide whether to abort or not
 - E.g.: this allows a parent to repeat a failed sub-transaction
- If the top-level transaction commits, then all of the sub-transactions that have provisionally committed can commit too, provided that none of their ancestors has aborted
- Note: the effect of sub-transactions are not permanent until the top-level transaction commits

Distributed Transactions

Atomic Commit Protocols

- Remember – Atomicity of transactions:
 - When a transaction comes to an end, either all of its operations appear to be carried out or none of them
- In case of a distributed transaction, these operations take place in more than one server
- How to complete a distributed transaction in an atomic fashion
 - Guarantee that either all of its operations (which are distributed across multiple servers) are carried out or none of them

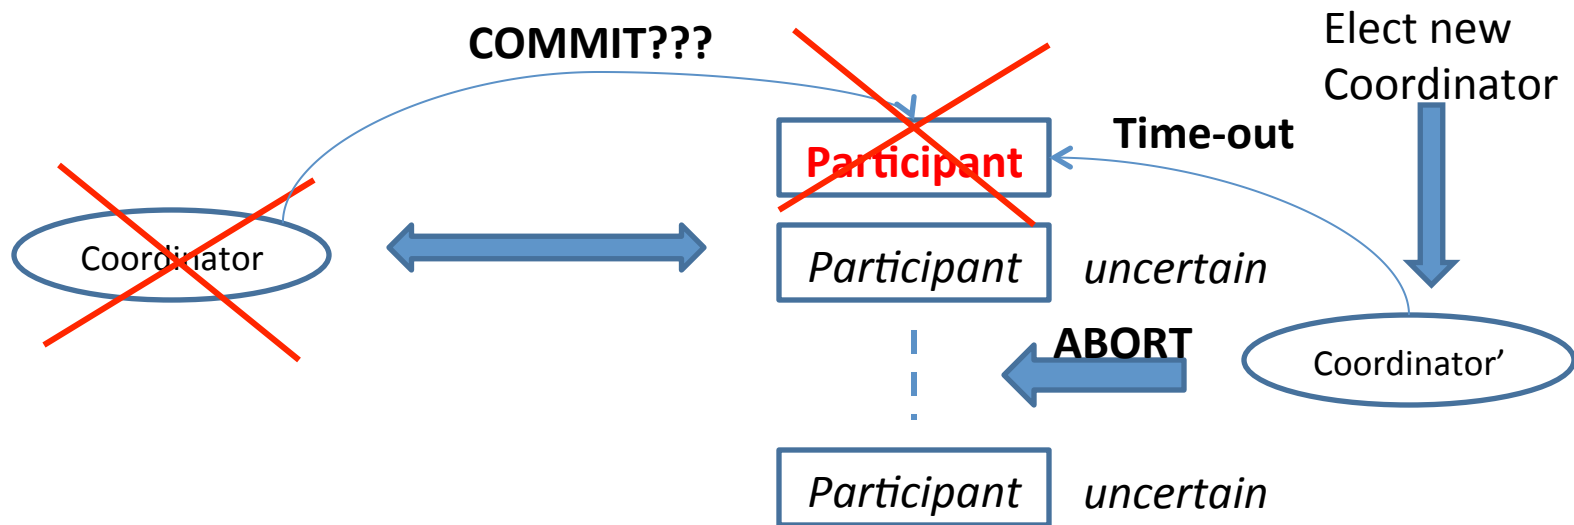
Two-phase Commit Protocol

- Is designed to allow any participant to abort the whole transaction
 - this requires informing all other participants and making sure that they get the message
 - Servers (coordinator, participants) can unilaterally decide to abort their part of a transaction
 - Due to the requirement of atomicity, if a part of a transaction is aborted, the complete transaction must be aborted
- Is based on a voting scheme: all the participants reach a consensus
 - either to commit or to abort
- Consists of two phases
 - Voting phase – ***prepare to commit***:
 - Coordinator asks participants whether they are ***prepared*** to commit
 - Completion phase - ***commit***:
 - Coordinator asks participants to ***commit***
- If a participant votes to commit, it must make sure that it can commit eventually, even if it crashes – participant must be able to recover from system failure by storing intermediate state

Coordinator Failures

- Serious problems occur in case of Coordinator failure
 - Participant received a canCommit? message in phase 1, is prepared to commit, but does not receive a doCommit message from the Coordinator
 - Participant is in an uncertain state and has to keep locks on data objects
 - Participants can wait for Coordinator to recover
 - Participants can elect a new Coordinator that restarts the 2PC protocol by sending new canCommit? Messages
- If coordinator fails in phase 2
 - It may have sent a doCommit or doAbort, before it crashed
 - One of the Participants may have received this information
 - Others can ask this Participant for the Coordinator's decision
 - If none of the Participants received a message from the failed coordinator
 - They can wait for Coordinator to recover
 - They can elect a new Coordinator that restarts the 2PC protocol by sending new canCommit? messages

Distributed Transactions



- Coordinator failed, one participant failed, no alive participant knows the coordinator's decision
- Problem
 - The coordinator's decision is unknown
 - It is unknown whether the crashed participant received a decision
 - The remaining participants may elect a new coordinator among them and vote – this new coordinator will time-out as the crashed participant does not deliver a vote and ***make an ABORT-decision***
 - **BUT:** If the crashed participant received a **COMMIT-decision**, it ***will commit data manipulations after recovery***
- Possible **CONTRADICTION !!**
- All participants have to wait for the original coordinator to recover to receive its original decision
 - All manipulated data objects remain locked


Coordinator fails in Phase 2

Scenario 3: Both Coordinator and one Participant fail

- Most critical case: both the coordinator and one participant failed
 - The failed Participant may have received a doCommit or doAbort, but which one? As it crashed, it cannot inform other participants
- If we assume that situation is resolved by electing a new coordinator / restarting the 2PC protocol:

- Participants elect new Coordinator
- New Coordinator restarts 2PC and sends canCommit? in phase 1
- as one Participant failed, it will not receive the full set of votes
- it will send a doAbort in phase 2
- the Participants alive will **abort**

- Failed Participant recovers – we assume it received a doCommit from the failed Coordinator
- Recovered Participant **commits**



- New Coordinator cannot be elected !!
- Participants have to wait until original Coordinator recovers and sends doCommit or doAbort

Reliability and Availability

Replication


- Replication is the maintenance of copies of data at multiple computers
- Replication is a technique for enhancing services, it is key to the effectiveness of distributed systems, as it provides
 - enhanced performance,
 - high availability and
 - fault tolerance
- Example:
 - Caching of resources from web servers within browsers and web proxy servers
- Replication is of particular importance for mobile computing, because it operates in a more disconnected way
- We will consider two approaches:
 - Passive (primary-backup) replication
 - Active replication
- We will also consider transactions over replicated data

Replication

- Key Advantages
 - Performance Enhancement
 - Caching (replication) of data by a client may improve performance of data access
 - E.g. web browsers cache web sites, images and other data
 - Increased Availability
 - Availability is influenced by
 - Server failures
 - Network partitioning and disconnected operations – communication disconnection are due to user mobility and often unplanned
 - The proportion of time a server is available should be as close to 100% as possible
 - Fault Tolerance
 - Fault-tolerant systems and services guarantee ***strictly correct*** behaviour despite a certain number and type of faults


Passive Replication

“Primary-backup” Replication



- At any one time there is a single primary (“master”) replica manager
- There are one or more secondary replica managers (“backups” or “slaves”)
- Front-ends communicate only with the primary replica manager
- The primary manager performs operations on data and sends copies to the backups
- If the primary fails, one of the backups becomes the primary

Active Replication




- All replica managers play the same role and are organised as group
- Client requests are multicast by the frontend to the replica manager group
- All replica managers process the request independently by identically
- All respond to the client request

Security

Symmetric-Key Encryption

(Secret Key Encryption)



- Same key is used by sender and receiver, has to be shared via some trusted channel

Secret Key Cryptography

- Relies on one key for encryption and decryption
- Symmetric model of a cryptography system
- Encryption algorithm should be hard to break
 - Attackers should be unable to decrypt ciphertext or discover the key (even if they have a set of corresponding cipher / plain texts)
- The larger the key size, the harder to attack
- Sender and receiver must obtain copies of the secret key in a secure fashion
- As long as the key is kept secret, the cryptographic procedure does not have to be secret

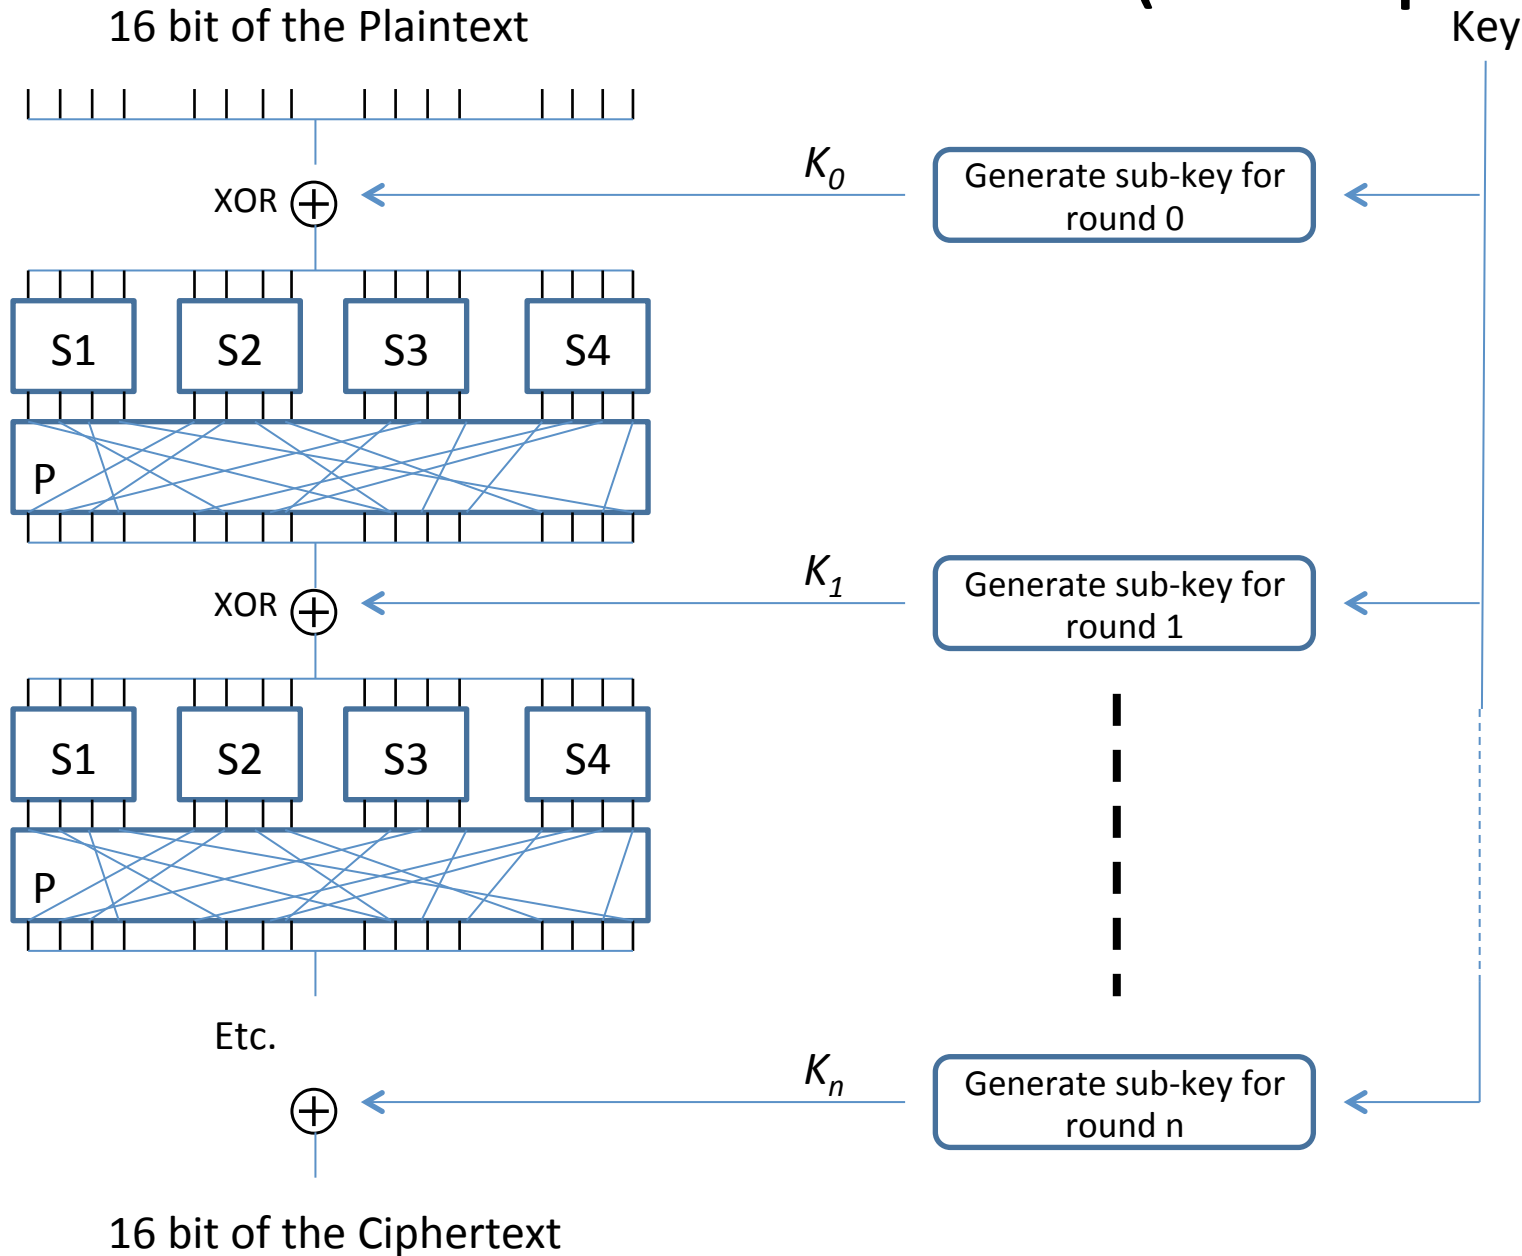
Block Ciphers

- Message is broken into blocks (usually 16 or 32bit words), each block is encrypted separately
- Operate with a fixed transformation procedure on large blocks of plaintext data
- Block Ciphers
 - Feistel Cipher
 - Is the first block cipher, which inspired subsequent cipher methods such as DES, etc.
 - Substitution-Permutation network

Substitution-Permutation Network

- SP-Networks (SPN) describe a series of substitution and permutation operations to be applied on plain text
 - The plaintext is separated into blocks (16bit words)
- The encoding operates over a sequence of rounds (“layers”), reapplying substitution and permutation operations over and over again to the output of a previous round

SP-Network (Example)



Substitution and Permutation

- Substitution in an SP-Network is performed with a so-called “substitution box” or S-box:
 - Are used to obscure or “confuse” the relationship between key and ciphertext
 - Takes as input m bits, produces a corresponding output of n bits
 - Implemented as a $m \times n$ lookup table
 - Substitution table is carefully designed to reduce vulnerability (Shannon: a change of one input bit should change at least half of all output bits)
- Permutation is performed with a so-called “permutation box” or P-box:
 - Takes the output of all substitution boxes as its input
 - Reorders (permutes) bits to produce output

Key Distribution

- With any symmetric algorithm, the key must be agreed upon by sender and receiver in a secure way
- Before 1976, key exchange was by far the biggest problem in secure communications
- Possible Strategies:
 - A key could be selected by A and physically delivered to B
 - A third party could select the key and physically deliver it to A and B
 - If A and B have previously used a key, one party could transmit the new key by encrypting it with the old key
 - If both A and B have an encrypted connection with a third party C, C could deliver a key on the encrypted links to A and B

Diffie-Hellman Key Exchange

- Developed in 1976, is a key exchange method where two parties exchange information that allows them to derive the same key, but never actually exchange the key
- Method:
 - Two parties, Alice and Bob, agree on a large prime number p and a small integer g ; these two numbers are public
 - Alice picks a secret large random integer a , and calculates a number A:

$$A = g^a \bmod p$$

- A becomes a public key, Alice transmits A to Bob
- Bob picks a secret large random integer b , and calculates a number B:


$$B = g^b \bmod p$$

- B becomes a public key, Bob transmits B to Alice
- Alice computes the secret key: $K_{Alice} = B^a \bmod p$
- Bob computes the secret key: $K_{Bob} = A^b \bmod p$

- Rules
 - p must be a prime number, $p > 2$
 - g must be a small integer, $g < p$
 - a and b are large random integers, $a < p-1$, $b < p-1$

Diffie-Hellman Key Exchange

- Example



Public Key Cryptography

- Using one secret key poses a security risk
- A solution to this problem is “Public Key Cryptography”
- Two keys: public and private (secret) key
- The keys are matched so that
 - A message encrypted with the public key can be decrypted using the private key
 - A message encrypted with the private key can be decrypted using the public key

Applications for Public Key Cryptography

- Encryption / Decryption
 - The sender encrypts the message with the recipients public key
- Digital Signature
 - The sender “signs” a message with its private key, for this, a cryptographic algorithm is applied to the whole message or to a small block of data that is a function of the message (a “fingerprint” of the message, called a message “digest”)
- Key Exchange
 - Exchange key information using the private key of one or both parties

RSA

- Encryption:
 - A ciphertext block C is the result of encryption of a plaintext block M , using the publicly known numbers e and n

$$C = M^e \bmod n$$

- Decryption
 - A plaintext block M is the result of decryption of a ciphertext block C , using the secret number d

$$M = C^d \bmod n = (M^e)^d \bmod n = M^{ed} \bmod n$$

RSA Key Generation

- To do:
 - Public Key: both sender and receiver must know the values of n and e
 - Calculate number n (maximum possible value of a plaintext / ciphertext block)
 - Calculate number e (a value needed for encryption)
 - Private Key
 - Calculate number d (a value needed for decryption)
 - only the receiver knows the value of d

RSA Key Generation

- Calculate n
 - Select two large prime numbers p and q , these are secret
 - Calculate: $n = p \times q$
- Calculate the public e
 - e is “relatively prime” to the Euler Totient $\phi(n)$
 - $e < \phi(n)$
- *What is*
 - “relatively prime” ?
 - Euler Totient $\phi(n)$?

RSA

- Relatively prime numbers:
 - Two integers n and m are relatively prime, if their greatest common divisor is 1: $\gcd(n,m) = 1$
 - N and m do not share any common positive prime factors (divisors) except 1
- Euler Totient $\phi(n)$
 - Is the **number** of positive integers that are $< n$ and that are relatively prime to n
 - E.g.: $n = 10$, $\{1,3,7,9\}$ is the set of positive integers relative prime to 10, therefore: $\phi(n) = 4$
 - If n is the product of two prime numbers, p and q , then $\phi(n) = (p-1)(q-1)$
 - E.g.:
 - $p = 3$, $q = 5$, $n = p \times q = 3 \times 5 = 15$, therefore: $\phi(n) = (p-1)(q-1) = 2 \times 4 = 8$
 - $n = 15$, $\{1,2,4,6,7,8,11,13\}$

RSA Key Generation

- Calculate n
 - Calculate: $n = p \times q$, p and q are two large prime numbers
- Calculate the public e
 - We know:
 - If $n = p \times q$, p and q are prime numbers, then $\phi(n) = (p-1)(q-1)$
 - Choose e :
 - e is relatively prime to $(p-1)(q-1)$ and $1 < e < (p-1)(q-1)$
- Calculate the private key d
 - $e \times d = 1 \bmod (p-1)(q-1)$
 - $d = e^{-1} \bmod (p-1)(q-1)$
- Result
 - Public key $K_{\text{PUB}} = \{e, n\}$
 - Private key $K_{\text{PRIV}} = \{d, n\}$
- *Encryption: encrypt a plaintext M to generate a ciphertext C via $C = M^e \bmod n$*
- *Decryption: decrypt a ciphertext C to generate a plaintext M via $M = C^d \bmod n$*

Digital Signatures

Digital Signatures

- Public key cryptography can also be used for creating digital signatures
 - Identifies reliably the originator of a send digital object (file / document, message etc.)
- Encryption of message with private key
 - Instead of a sender using the public key of the receiver, the sender's private key is used to create a unique signature

Digital Signatures


Creation from a Message

- A digital signature for a particular message is created by encoding a message with the private key
 - The message itself is not secret as anybody can decode it with the public key
- Only the person with the private key can produce the signature




Digital Signature Verification

- A transmitted digital signature can be verified with the public key, identifies uniquely the sender.
- As public key is public, anyone can verify that the signature is valid



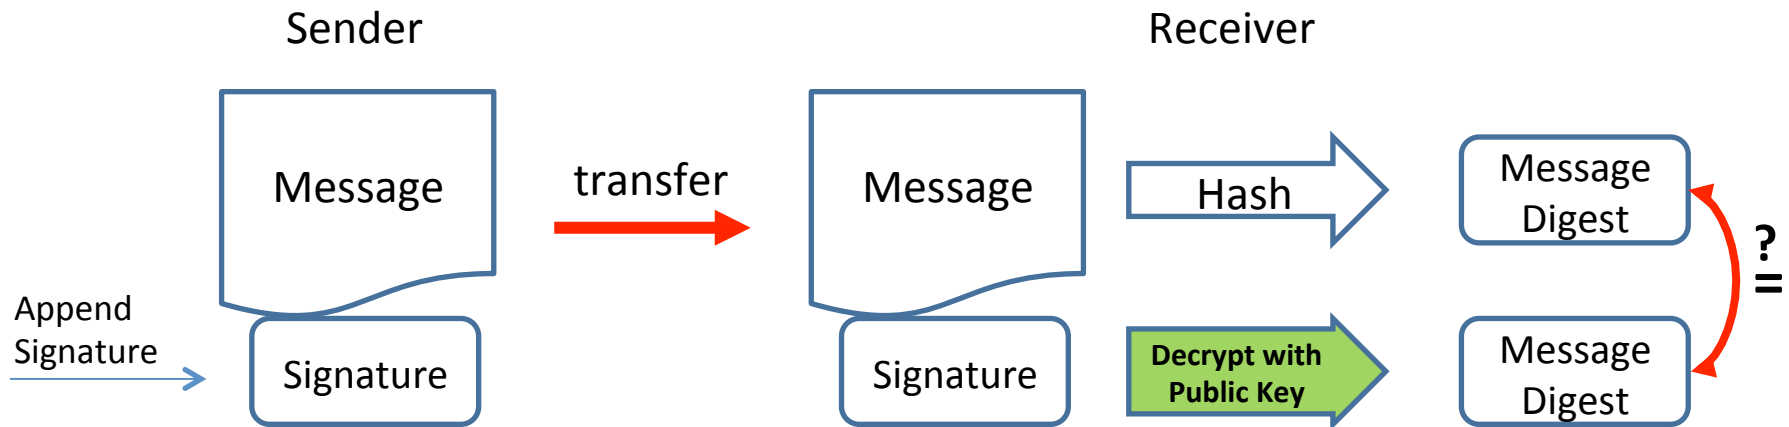
Digital Signature – Message Digest

- Encoding the whole message to produce a digital signature is not feasible
- Solution: create a “Message Digest”
 - Is a kind of “fingerprint” of a message, which is much shorter
 - the message digest is encrypted with the Private Key to create a unique signature for the message
- Calculation of the message digest
 - Apply a hash function to the message



Digital Signature – Message Digest

- Sender:
 - Creates signature from message digest
 - Appends signature to message and sends it
- Receiver
 - Recreates the message digest from the received message
 - Decrypts the signature with the public key of the sender
 - Compares the two results to check origin of message



Digital Certificate

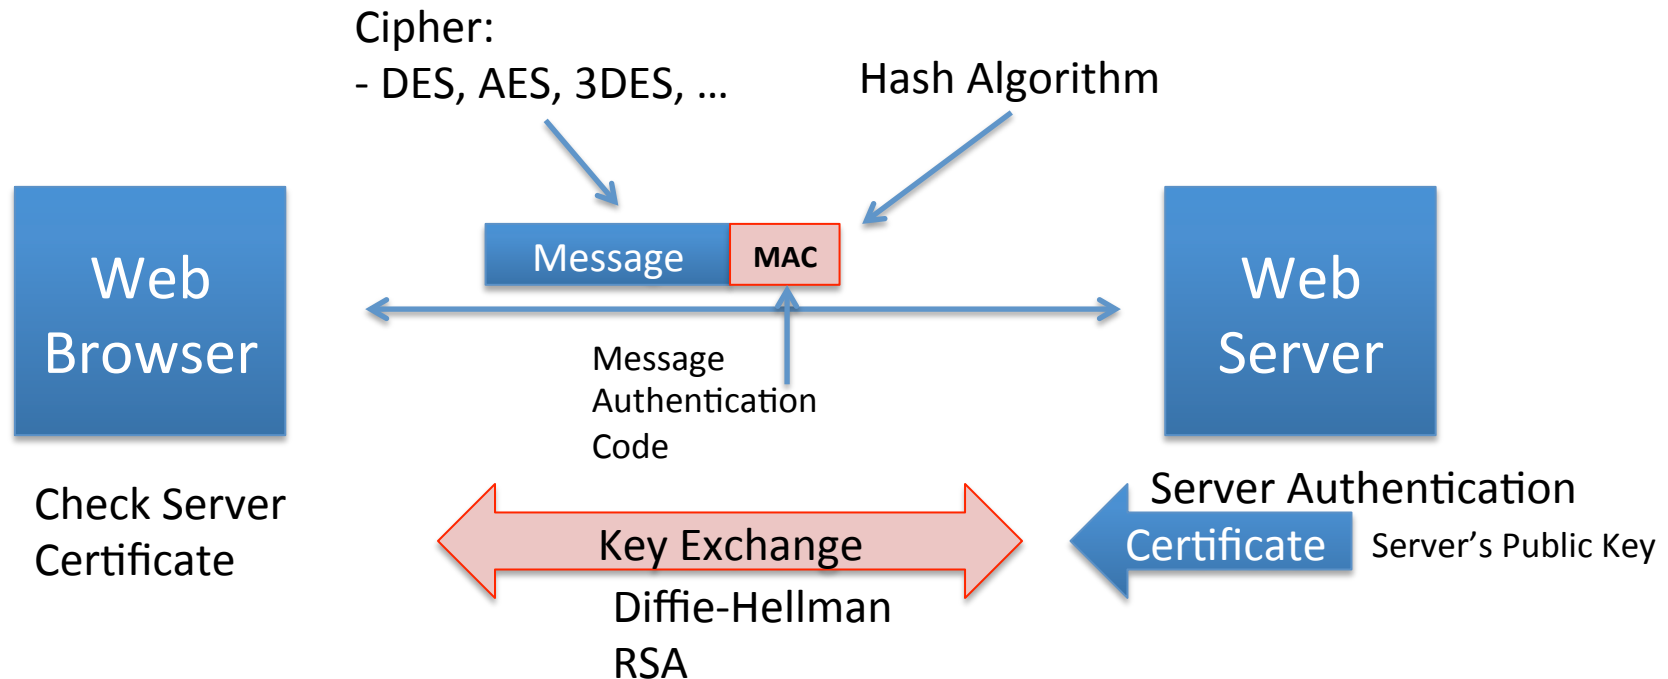
- Helps to address the problem of public-key distribution and authenticity
 - Certification:
 - digital certificates are introduced to proof authenticity of public keys
 - Validation:
 - ability to check that the binding of a public key to a certificate is authentic
- Is a “statement of originality” by a third party, the “Certificate Authority” (CA)
 - Binds a public key to a particular owner
 - A digital certificate itself is authenticated with a digital signature, signed by the CA with its private key
 - One’s trust that a certified public key is original relies on one’s trust in the validity of the CA’s key

Public Key Infrastructure

- A Public Key Infrastructure provides the environment for the management of digital certificates
- Elements of a PKI
 - “Certificate Authority” (CA)
 - Is a trusted third party that issues and verifies digital certificates
 - Uses its own private key to sign digital certificates
 - “Registration Authority” (RA)
 - verifies the identity of users requesting information from the CA
 - A central repository for storing public keys

SSL Secure Socket Layer

SSL Secure Socket Layer



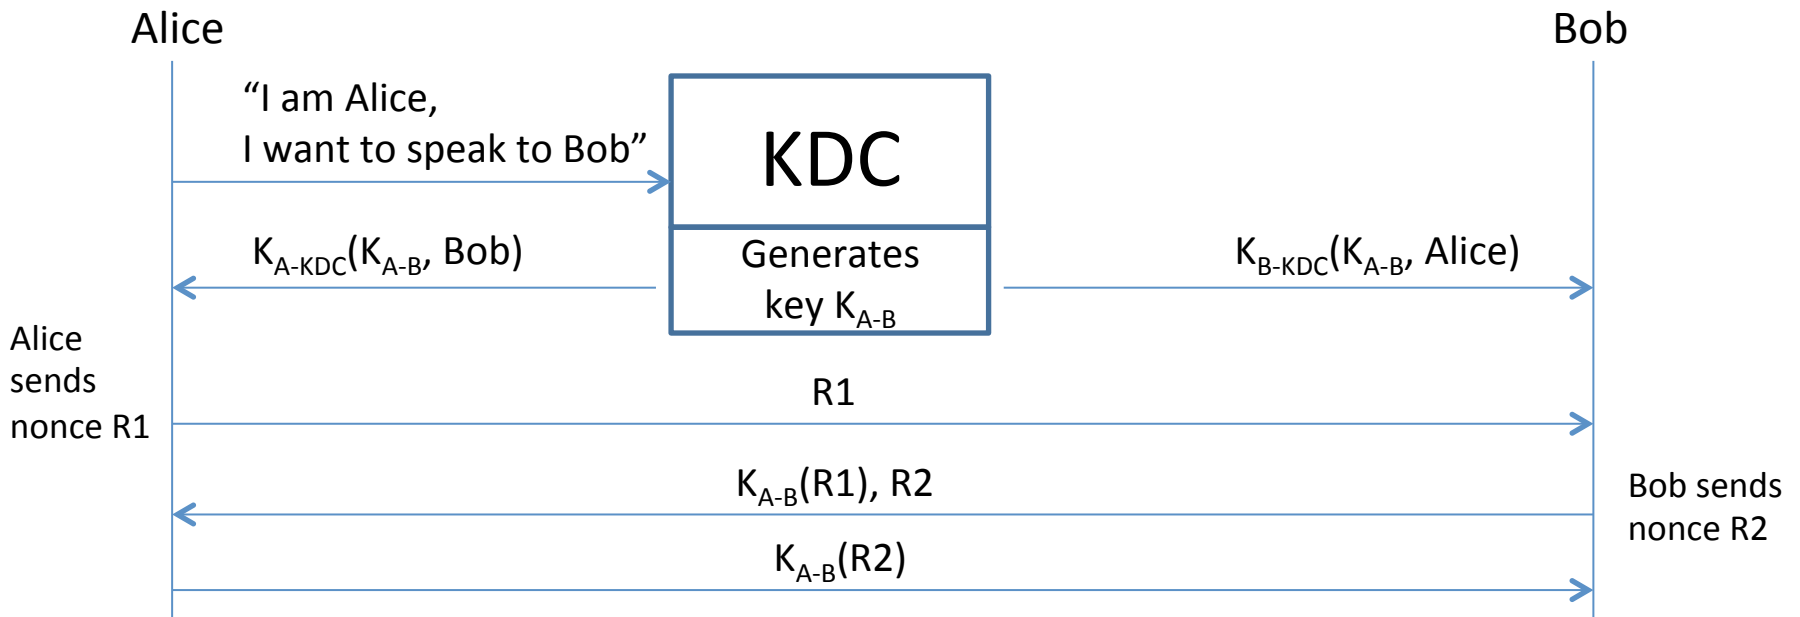
- Before encrypted communication
 - Negotiate cipher suite: crypto algorithm, hash algorithms for MAC
 - Authenticate Server (one-way authentication)
 - (maybe server also requests client authentication with certificate: two-way authentication)
 - Exchange information for secret key with Diffie-Hellman, RSA etc. key exchange

SSL Secure Socket Layer

- SSL can be viewed as a security layer that sits between the application layer and the transport layer
 - The client hands over data to SSL (e.g. An HTTP message for a web server), SSL then encrypts this data and writes it to a TCP socket
 - The server receives this encrypted data via its TCP socket, SSL takes this data, decrypts it and directs this data to the server for processing
- SSL uses symmetric key cryptography for encryption and decryption of data that is transferred
- To Do:
 - Verify that server is trustworthy (certificates)
 - Exchange a symmetric key between server and client

Mediated Authentication

- Goal: two parties eventually share a secret symmetric key (K_{A-B}) for communication
 - Communication partners have keys to contact the KDC (K_{A-KDC} , K_{B-KDC})
 - Two nonces, “R1” and “R2” are used to proof authenticity



"Alice" has a key K_{A-KDC} to communicate with KDC

"Bob" has a key K_{B-KDC} to communicate with KDC

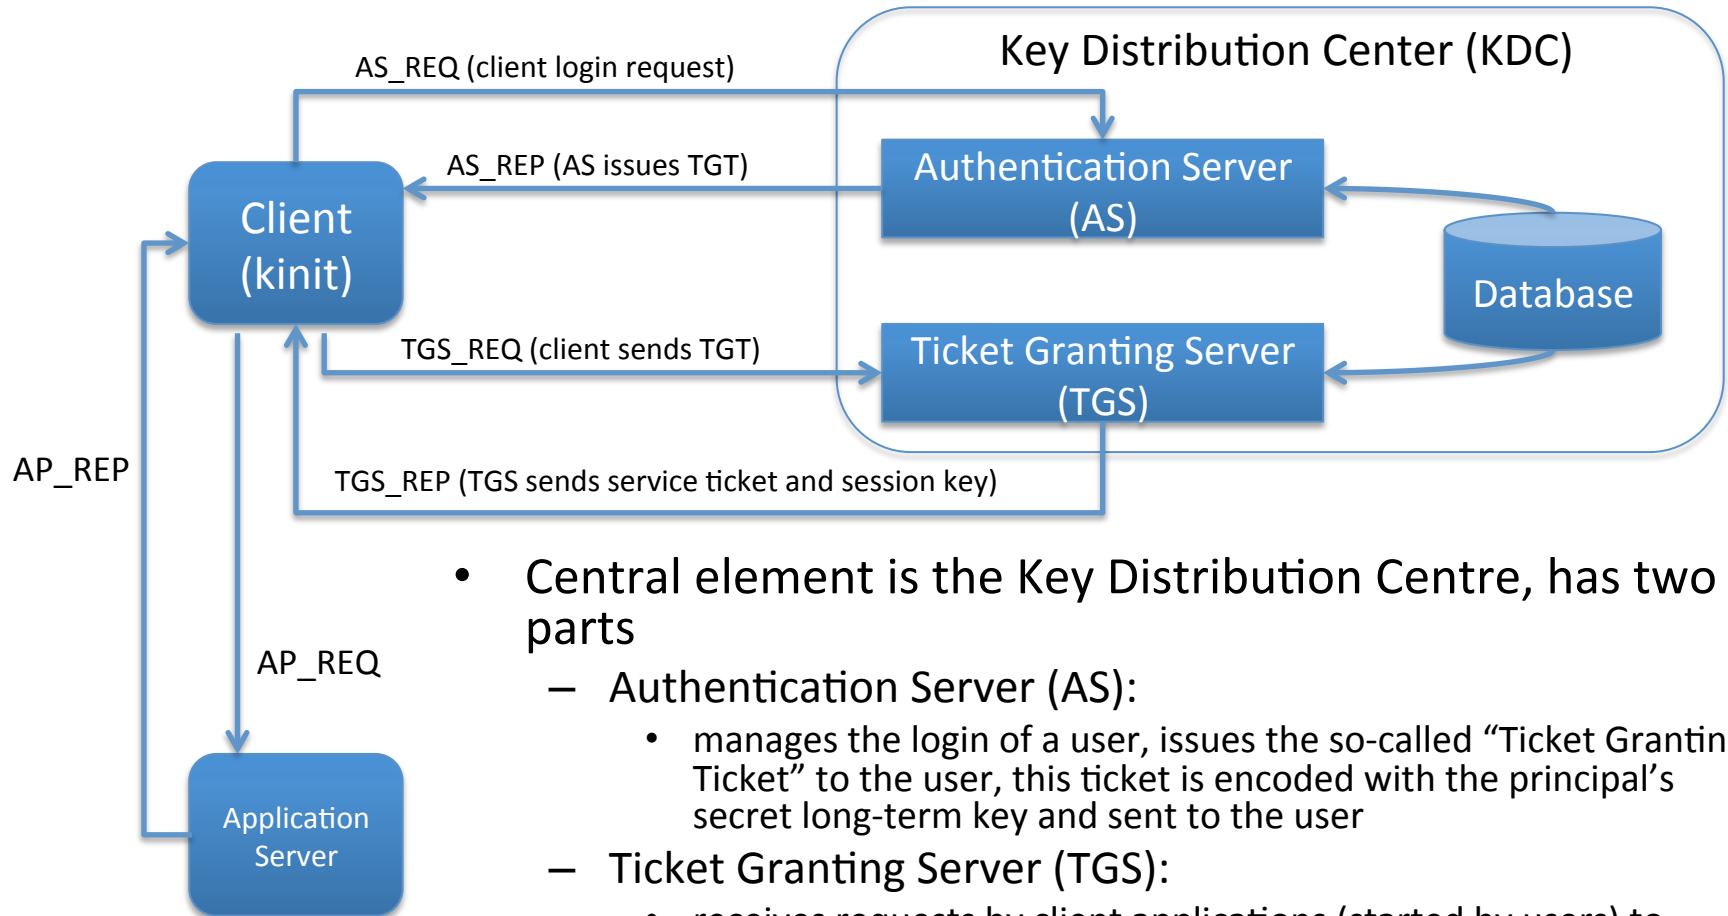
Kerberos

- Widely adopted and implemented in popular operating systems
 - See <http://www.kerberos.org>
- Kerberos uses time stamps as “nonces” in the mutual authentication phase of the protocol
- It uses mediated authentication with tickets

Kerberos Infrastructure

- Users and services are called “principals” – they have to be registered with a Kerberos domain, called a “realm” (managed by the Key Distribution Centre,
 - each principal is stored in a central database with an ID and a secret key
 - This secret key is regarded the Principal’s “long-term” secret key
- Central element is the Key Distribution Centre, has two parts
 - Authentication Server (AS):
 - manages the login of a user, issues the so-called “Ticket Granting Ticket” to the user, this ticket is encoded with the principal’s secret long-term key and sent to the user
 - Ticket Granting Server (TGS):
 - receives requests by client applications (started by users) to access servers, client sends the TGT to the KDC for such a request
 - TGS calculates a secret session key for communication between client and server, is sent to client with a new service ticket (service ticket is encoded with the server’s long-term secret key)

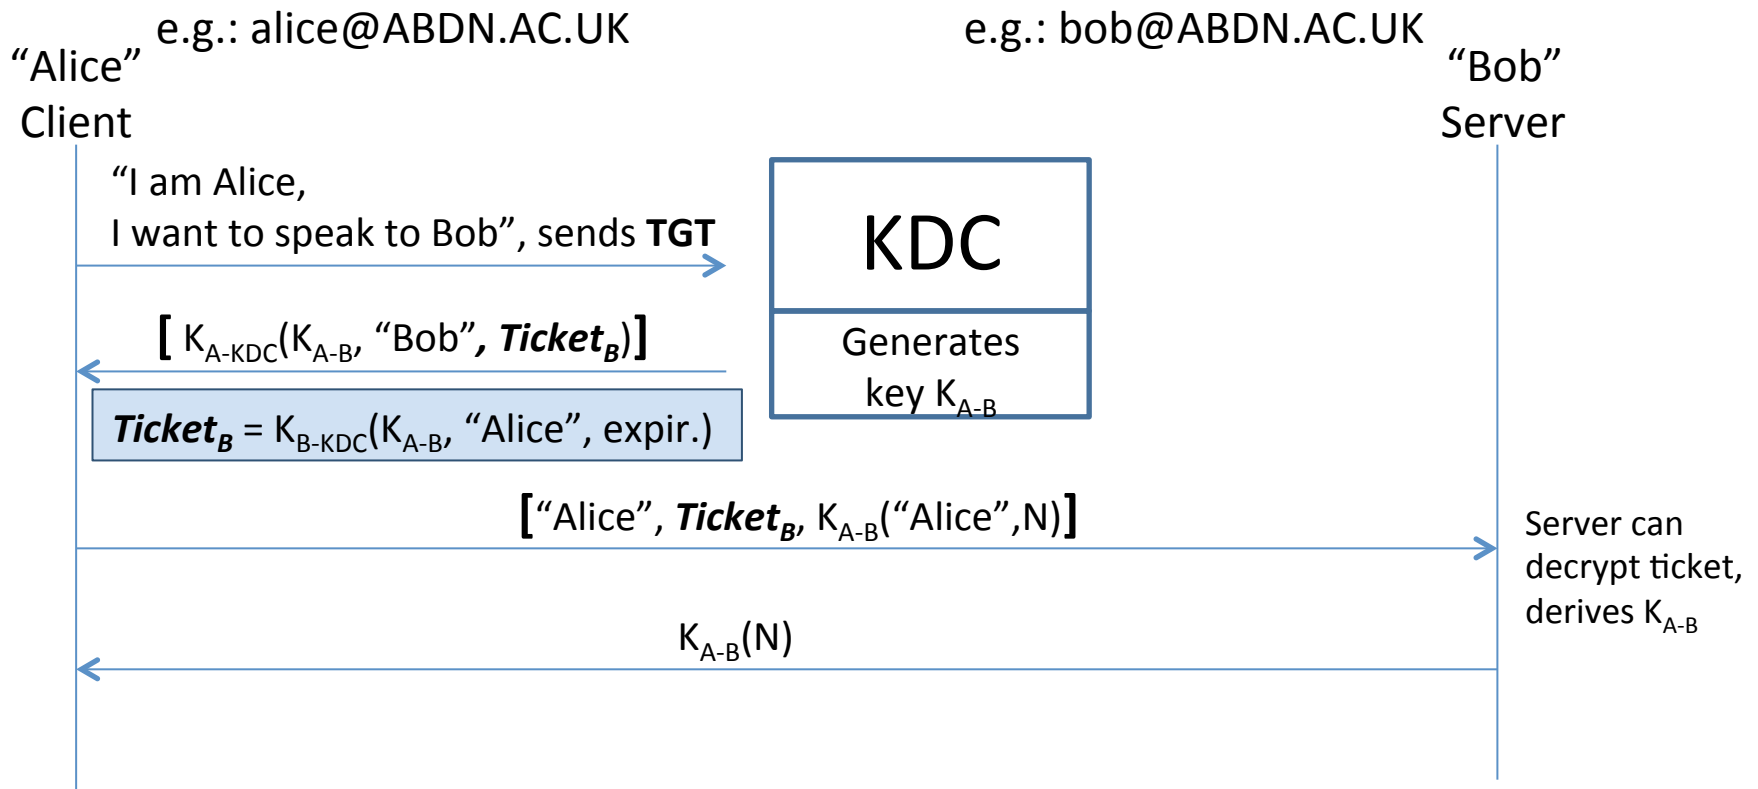
Kerberos Infrastructure



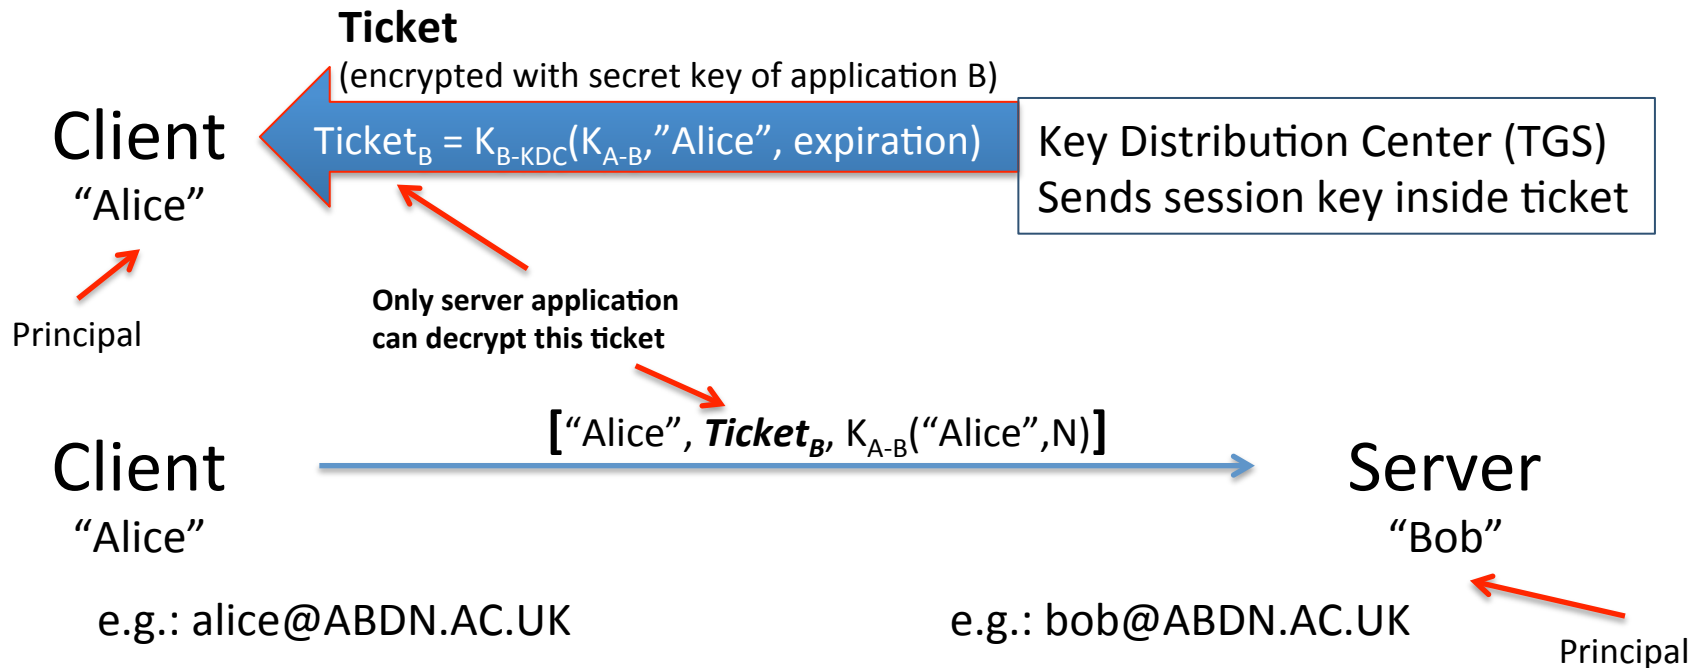
- Central element is the Key Distribution Centre, has two parts
 - Authentication Server (AS):
 - manages the login of a user, issues the so-called “Ticket Granting Ticket” to the user, this ticket is encoded with the principal’s secret long-term key and sent to the user
 - Ticket Granting Server (TGS):
 - receives requests by client applications (started by users) to access servers, client sends the TGT to the KDC for such a request
 - TGS calculates a secret session key for communication between client and server, is sent to client with a new service ticket (service ticket is encoded with the server’s long-term secret key)

Kerberos Authentication and Access to Applications

- Client sends TGT for authentication to KDC (TGS)
- KDC (TGS) responds with a session ticket encoded with the server's long-term key (K_{B-KDC})
- A timestamp N is used as a nonce



Ticket Transfer



- Key Distribution Center transmits a ticket that is encrypted with the server's long-term secret key (K_{B-KDC})
- Only server can decrypt this ticket, therefore authenticating the client

Kerberos Protocol

- Accessing Services
 - A client application (used by an authenticated user) sends user's TGT to the KDC, indicating that it wants to use a particular service
 - The KDC authenticates the client, checks access privileges to service, generates a random symmetric (short-term) session key K_{A-B} for communication between client and server
 - The KDC sends a message back to the client, encoded with the shared key K_{A-KDC} : the value of K_{A-B} , and a **ticket** for accessing the service
 - $K_{A-KDC}(K_{A-B}, \text{Ticket}_B)$, where $\text{Ticket}_B = K_{B-KDC}(\text{"client"}, K_{A-B}, \text{expir.})$
 - The client sends the ticket to the service, client also sends an **authenticator** for message to the service; the authenticator consists of the client name and a timestamp (nonce) N encrypted with K_{A-B} , that is $K_{A-B}(\text{"client"}, N)$
 - The service decrypts the ticket, using the secret key K_{B-KDC} , with that it will learn about the session key K_{A-B}
 - The service sends back the nonce to the client, encoded with K_{A-B} to show that it received the secret session key and is "alive". This is the **mutual authenticator**.