

CS2521: Sorting

N. Oren

`n.oren@abdn.ac.uk`

University of Aberdeen

Why?

- Sorting is one of the most fundamental problems in CS
- Used as a building block for many other algorithms
- Many interesting generic ideas emerge
- Computers spend most of their time sorting (1/4 of all mainframe time was spent sorting)
- It's the most studied problem in CS
- Many problems, once sorted, become a lot easier. E.g. binary search (see Skiena for many others)

- Good sorting algorithms run in $O(n \log n)$.
- Naïve ones in $O(n^2)$
- For 1000000 elements, it's 1000 seconds vs 0.6 seconds
- Applications: bioinformatics, astronomy, particle physics ($n > 10^9$ typically)

- How can we determine whether two sets are disjoint (assume sets of size m and n , $m \ll n$)?

```
1: for all  $i \in S_m$  do  
2:    $j=0$ ,  $\text{found}=\text{false}$   
3:   while  $j < |S_n|$  and  $\text{!found}$  do  
4:     if  $i == S_n[j]$  then  $\text{found}=\text{true}$   
5:   end while  
6:   if  $\text{found}==\text{false}$  then return false  
7: end for  
8: return true
```

- Complexity?

- How can we determine whether two sets are disjoint (assume sets of size m and n , $m \ll n$)?
- Sort the big set, do a binary search for each element in the smaller set. Complexity: $O(n \log n)$ for sorting, $O(m \log n)$ for search. Total: $O((n + m) \log n)$
- Sort the small set and search from big set. Complexity: $O((n + m) \log m)$

- How can we determine whether two sets are disjoint (assume sets of size m and n , $m \ll n$)?
- Sort both sets.
- Now no need to do a binary search: compare smallest elements, discard smaller one if not identical. Keep repeating on the now smaller sets. Complexity?

- How can we determine whether two sets are disjoint (assume sets of size m and n , $m \ll n$)?
- Sort both sets.
- Now no need to do a binary search: compare smallest elements, discard smaller one if not identical. Keep repeating on the now smaller sets. Complexity?
- We visit each element of each set once at most, i.e. $O(n + m)$
- Total complexity: $O(n \log n + m \log m + n + m)$

Sample Application

- How can we determine whether two sets are disjoint (assume sets of size m and n , $m \ll n$)?
- $O(nm)$, $O((n+m) \log n)$, $O((n+m) \log m)$, $O(n \log n + m \log m + n + m)$
Which is fastest?
- Clearly not $O(nm)$.
- Now $\log m < \log n$ so we eliminate $O((n+m) \log n)$
- Expanding, we see this is better than $O(n \log n + m \log m + n + m)$
- So best to sort the small set only. If m is constant, complexity is linear!
- Note that we could build a hash table with elements of set n , and verify that collisions occur for all m elements in linear time.

Choices...

- Increasing or decreasing order?
- Sort just key or entire record?
- What should be done with equal keys?
- What do we do with non-numerical data?

Choices...

- Increasing or decreasing order?
- Sort just key or entire record?
- What should be done with equal keys?
- What do we do with non-numerical data?
- We assume the existence of a comparison function


Selection Sort

```
1: function SelectionSort(s)
2:   for all i=0 to n-1 do
3:     min=i
4:     for all j=i+1 to n-1 do
5:       if s[j]<s[min] then
6:         min=j
7:       end if
8:     end for
9:     swap s[i],s[min]
10:  end for
11: end function
```

6
2
4
8
1
3


Selection Sort

```
1: function SelectionSort(s)
2:   for all i=0 to n-1 do
3:     min=i
4:     for all j=i+1 to n-1 do
5:       if s[j]<s[min] then
6:         min=j
7:       end if
8:     end for
9:     swap s[i],s[min]
10:  end for
11: end function
```




Selection Sort

```
1: function SelectionSort(s)
2:   for all i=0 to n-1 do
3:     min=i
4:     for all j=i+1 to n-1 do
5:       if s[j]<s[min] then
6:         min=j
7:       end if
8:     end for
9:     swap s[i],s[min]
10:  end for
11: end function
```



Selection Sort

```
1: function SelectionSort(s)
2:   for all i=0 to n-1 do
3:     min=i
4:     for all j=i+1 to n-1 do
5:       if s[j]<s[min] then
6:         min=j
7:       end if
8:     end for
9:     swap s[i],s[min]
10:  end for
11: end function
```



Selection Sort

```
1: function SelectionSort(s)
2:   for all i=0 to n-1 do
3:     min=i
4:     for all j=i+1 to n-1 do
5:       if s[j]<s[min] then
6:         min=j
7:       end if
8:     end for
9:     swap s[i],s[min]
10:  end for
11: end function
```

```
1: function SelectionSort(S)
2:   N=new array of size |S|
3:   for all i=0 to n-1 do
4:     sml=find smallest(S)
5:     N[i] = sml
6:     delete sml from S
7:   end for
8:   return N
9: end function
```

Selection Sort

- find smallest = $O(n)$ (sweep through average array size $n/2$)
- For loop $O(n)$
- Total time $O(n^2)$
- But we can speed up “find smallest”?

```
1: function SelectionSort( $S$ )
2:    $N$  = new array of size  $|S|$ 
3:   for all  $i=0$  to  $n-1$  do
4:      $sml$  = find smallest( $S$ )
5:      $N[i] = sml$ 
6:     delete  $sml$  from  $S$ 
7:   end for
8:   return  $N$ 
9: end function
```


Selection Sort


- find smallest= $O(n)$ (sweep through average array size $n/2$)
- For loop $O(n)$
- Total time $O(n^2)$
- But we can speed up “find smallest”?
- Yes, through a better data structure: balanced binary tree or heap. $O(\log n)$
- New time: $O(n \log n)$

```
1: function SelectionSort( $S$ )
2:    $N$ =new array of size  $|S|$ 
3:   for all  $i=0$  to  $n-1$  do
4:      $smI$ =find smallest( $S$ )
5:      $N[i] = smI$ 
6:     delete  $smI$  from  $S$ 
7:   end for
8:   return  $N$ 
9: end function
```

- Given that we don't (yet) know how to construct a balanced binary tree, we'll look at heaps.
- A heap is a data structure that efficiently supports priority queue operations insert and find/delete-min.
- A heap keeps some sort of loose ordering on the set of elements.
- A heap encodes a tree structure, wherein each element in the tree dominates its children.
- A organisational hierarchy is a good example of a heap; the boss sits on the top, then managers below, then “normal” employees (whose boss is a specific manager) etc. The employer dominates the employees.
- In a min-heap a node dominates its children by containing a smaller key than they do.
- In a max-heap a node dominates its children by containing a larger key than they do.

- Keys are typically small; if we use a binary tree representation, the pointers can use up more space than the keys themselves. We seek a different representation.
- Heaps allow us to represent binary trees without pointers.

Heaps



- We store the keys in an array, with the position of the keys implicitly acting as pointers.
- The first index stores the root, index 2 stores the left hand child, 3 stores the right, etc.
- Indexes 2^{l-1} to $2^l - 1$ store the 2^{l-1} keys at depth l of the tree.

- The left child of an element at index k sits at position $2k$ in the tree, the right child at $2k + 1$.
- The parent of k at position $\lfloor k/2 \rfloor$
- What's the catch?

- The left child of an element at index k sits at position $2k$ in the tree, the right child at $2k + 1$.
- The parent of k at position $\lfloor k/2 \rfloor$
- What's the catch?
- All missing internal nodes take up space as we need to represent the full tree (except the right part of the last level) to maintain positional mapping.
- Without this requirement, worst case is array of size 2^n to store n elements.
- Tradeoff between flexibility and space.

Using heaps

- How can we search for a key in a heap?

Using heaps




- How can we search for a key in a heap?
- We can't. A heap is not a binary search tree.
- We have to do a linear search.

Where are we?

- We've described the structure of a heap, a pointer free representation of trees.
- In a heap, each node dominates its children.
- We want to show that we can find the minimum element of a heap quickly, and feed it into our sort operation to speed it up.
- We need to be able to
 - Build the heap
 - Find its minimum
 - Delete its minimum

Building a Heap







- We can build a heap incrementally inserting each new element into the left-most open spot in the array (i.e. the $n+1$ 'st position of a n element heap).
- This maintains balance of the heap.
- BUT, does not maintain dominance.



Building a Heap

- To solve: we swap the new element with its parent as necessary.
- Old parent is clearly dominated.
- Other child is clearly dominated (even more).
- But new parent may still dominate its parent.
- So repeat, bubbling up as far as needed.
- Heap order is preserved.

Bubbling









```

1: function Insert(x,heap)
2:   heap[|heap|]=x
3:   bubble(heap,|heap|)
4: end function
5: function bubble(heap,index)
6:   if heap[index]>heap[index/2] then
7:     swap(heap[index],heap[index/2])
8:     bubble(heap,index/2)
9:   end if
10: end function
    
```

• Complexity?

Bubbling









```

1: function Insert(x,heap)
2:   heap[|heap|]=x
3:   bubble(heap,|heap|)
4: end function
5: function bubble(heap,index)
6:   if heap[index]>heap[index/2] then
7:     swap(heap[index],heap[index/2])
8:     bubble(heap,index/2)
9:   end if
10: end function
    
```

- n inserts
- $\log n$ bubbles per insert

Bubbling



```

1: function Insert(x,heap)
2:   heap[|heap|]=x
3:   bubble(heap,|heap|)
4: end function
5: function bubble(heap,index)
6:   if heap[index]>heap[index/2] then
7:     swap(heap[index],heap[index/2])
8:     bubble(heap,index/2)
9:   end if
10: end function
  
```

- n inserts
- $\log n$ bubbles per insert
- $O(n \log n)$

Making the heap

```
1: function makeHeap(array)
2:   for all i=0 to |array| do
3:     heap=array[0..i]
4:     insert(array[i],heap)
5:   end for
6: end function
```

- A heap is made in-place, using no additional memory.
- Since the original array was unsorted, we end up with a new unsorted array, which happens to be a heap.


Finding the Minimum

- Easy: look at the first element of the heap.
- $O(1)$

Deleting the Minimum

- Once we delete the minimum, we have a hole in our heap.
- Replace by moving right most leaf into first position.
- Tree shape is restored. BUT heap property might be violated.
- The root could be dominated by both of its children.
- Swap root with its smallest child.
- Repeat for next level until dominance is restored for heap.
- Non-dominated element bubbles down to its place. Also called heapify as the two subtrees (heaps) below original root are merged into a new heap.


Heapifying



1
4
12
42
93
21
34
54


```
1: function ExtractMin(heap)
2:   ans=heap[0]
3:   move heap[|heap|-1] to heap[0] and resize
4:   heapify(heap,0)
5:   return ans
6: end function
7: function heapify(heap,index)
8:   sIndex=index
9:   if 2*index>|heap| then return
10:  if heap[2*index]<heap[sIndex] then
      sIndex=2*index
11:  if heap[1+2*index]<heap[sIndex] then
      sIndex=1+2*index
12:  swap heap[index] and heap[sIndex]
13:  heapify(heap,sIndex)
14: end function
```

Heapifying




```
1: function ExtractMin(heap)
2:   ans=heap[0]
3:   move heap[|heap|-1] to heap[0] and resize
4:   heapify(heap,0)
5:   return ans
6: end function
7: function heapify(heap,index)
8:   sIndex=index
9:   if 2*index>|heap| then return
10:  if heap[2*index]<heap[sIndex] then
    sIndex=2*index
11:  if heap[1+2*index]<heap[sIndex] then
    sIndex=1+2*index
12:  swap heap[index] and heap[sIndex]
13:  heapify(heap,sIndex)
14: end function
```

Heapifying




```
1: function ExtractMin(heap)
2:   ans=heap[0]
3:   move heap[|heap|-1] to heap[0] and resize
4:   heapify(heap,0)
5:   return ans
6: end function
7: function heapify(heap,index)
8:   sIndex=index
9:   if 2*index>|heap| then return
10:  if heap[2*index]<heap[sIndex] then
      sIndex=2*index
11:  if heap[1+2*index]<heap[sIndex] then
      sIndex=1+2*index
12:  swap heap[index] and heap[sIndex]
13:  heapify(heap,sIndex)
14: end function
```

Heapifying



```
1: function ExtractMin(heap)
2:   ans=heap[0]
3:   move heap[|heap|-1] to heap[0] and resize
4:   heapify(heap,0)
5:   return ans
6: end function
7: function heapify(heap,index)
8:   sIndex=index
9:   if 2*index>|heap| then return
10:  if heap[2*index]<heap[sIndex] then
11:    sIndex=2*index
12:  if heap[1+2*index]<heap[sIndex] then
13:    sIndex=1+2*index
14:  swap heap[index] and heap[sIndex]
15:  heapify(heap,sIndex)
16: end function
```

Heapifying



- Get root and swap last/first: $O(1)$
- Maximum of $\lfloor \log n \rfloor$ heapify operations.
- Total complexity $O(\log n)$

Where are we?

- Making the heap: $O(n \log n)$
- Extracting minimal element n times $O(n \log n)$
- We can shrink the heap each time (incrementing its start position), leaving extracted element in place.
- Result: sorted list, in $O(n \log n)$ time which uses no additional memory.
- This is heap sort
- Heap construction can actually be done in linear time, is this useful?

Faster heap construction

- An array is almost a heap - it doesn't have the heap property, but can be represented as a tree.
- But the leaves in this pseudo-heap do have the heap property.
- Consider parents of leaves - they may not dominate children, but we can swap things around so they do (how?)

Faster heap construction

- An array is almost a heap - it doesn't have the heap property, but can be represented as a tree.
- But the leaves in this pseudo-heap do have the heap property.
- Consider parents of leaves - they may not dominate children, but we can swap things around so they do (how?)
- By calling `heapify` at that index.

Faster heap construction

- Total number of heapify calls is n , and running time is $O(\log n)$ so no saving.
- But only the call at the root would take all $\log n$ steps.
- Heapify takes time proportional to the height of tree it is merging, and most heaps are small. Given a tree of n elements
 - $n/2$ are leaves
 - $n/4$ are parents of leaves
 - ...
- So there are at most $\lceil n/2^{h+1} \rceil$ nodes of height h .

Faster heap construction

- Total cost of building heap is therefore

$$\sum_{h=0}^{\lfloor \log n \rfloor} \lceil n/2^{h+1} \rceil h \leq n \sum_{h=0}^{\lfloor \log n \rfloor} h/2^h$$

- $\sum h/2^h$ converges (below 2)
- Therefore, complexity is $O(n)$.

A different approach to sorting

- Basic idea: select a random element from the unsorted set, and put it in the proper position in the sorted set.

```
1: function InsertionSort(A)
2:   for i = 2 to n do
3:     j = i
4:     while (A[j] < A[j - 1]) & j > 0 do
5:       swap(A[j], A[j - 1])
6:       j = j - 1
7:     end while
8:   end for
9: end function
```

- $O(n^2)$ in the worst case.
- Works well if array is almost sorted already
- An example of incremental insertion; building up a structure on n by first building on $n - 1$.

Divide and Conquer

- We've seen before (e.g. in binary search) that dividing and conquer is a useful approach in many situations.
 - divide the problem into smaller problems.
 - solve each of the smaller problems (typically by further division and solving).
 - combine the solutions of the subproblems into a solution of the bigger problem.
- Works easily for binary search as there are no interdependencies between subproblems.
- Can we use the approach for sorting?

Mergesort

```
1: function mergeSort( $A[1..n]$ )  
2:   merge(mergeSort( $A[1..\lfloor n/2 \rfloor]$ ), mergeSort( $A[\lfloor n/2 \rfloor + 1..n]$ ))  
3: end function
```

- Base case?

```
1: function mergeSort( $A[1..n]$ )  
2:   merge(mergeSort( $A[1..\lfloor n/2 \rfloor]$ ), mergeSort( $A[\lfloor n/2 \rfloor + 1..n]$ ))  
3: end function
```

- Base case? Occurs when there's only a single element as there's nothing to sort.
- Efficiency of mergesort critically depends on how efficiently we can merge the sorted sublists.
- It would be silly to call heapsort or the like to do the combining.

Merging

- Assume we have 2 sorted lists.
- The smallest item must be at the top of one of them.
- We pick that item, put it in the new list, and repeat until no items exist in either list.

1	4
3	5
9	11

Merging

- Assume we have 2 sorted lists.
- The smallest item must be at the top of one of them.
- We pick that item, put it in the new list, and repeat until no items exist in either list.

3
9

4
5
11

1

Merging

- Assume we have 2 sorted lists.
- The smallest item must be at the top of one of them.
- We pick that item, put it in the new list, and repeat until no items exist in either list.



9

4
5
11

1
3

Merging

- Assume we have 2 sorted lists.
- The smallest item must be at the top of one of them.
- We pick that item, put it in the new list, and repeat until no items exist in either list.



Merging

- Assume we have 2 sorted lists.
- The smallest item must be at the top of one of them.
- We pick that item, put it in the new list, and repeat until no items exist in either list.

9

11

1
3
4
5

Merging

- Assume we have 2 sorted lists.
- The smallest item must be at the top of one of them.
- We pick that item, put it in the new list, and repeat until no items exist in either list.

11

1
3
4
5
9

Merging

- Assume we have 2 sorted lists.
- The smallest item must be at the top of one of them.
- We pick that item, put it in the new list, and repeat until no items exist in either list.

1
3
4
5
9
11

Merging

- Assume we have 2 sorted lists.
- The smallest item must be at the top of one of them.
- We pick that item, put it in the new list, and repeat until no items exist in either list.
- Merging takes $n - 1$ comparisons, i.e. $O(n)$

1
3
4
5
9
11

Merging

```
1: function merge( $A_1, A_2$ )
2:    $A_3 = []$ 
3:    $i = 0$ 
4:   while  $|A_1 \cup A_2| > 1$  do
5:     if  $A_1[0] \leq A_2[0]$  then
6:        $A_3[i] = A_1[0]$ 
7:       Remove  $A_1[0]$  from  $A_1$ 
8:     else
9:        $A_3[i] = A_2[0]$ 
10:      Remove  $A_2[0]$  from  $A_2$ 
11:    end if
12:     $i = i + 1$ 
13:  end while
14:  add last element of  $A_1 \cup A_2$  to
     $A_3[i]$ 
15: end function
```

1
3
4
5
9
11

- To compute the running time of mergesort, we must consider how much work is done at each step of the divide & conquer approach.
- Let us assume we have n elements, such that $\log n \in \mathbb{Z}$
- The top level consists of 1 call to merge, processing all n elements.
- Level 2 consists of 2 calls to mergesort, and each processing $n/2$ elements.
- Level 3 has 4 calls to mergesort, each processing $n/4$ elements.
- Level k calls mergesort 2^k , each processing $n/2^k$ elements.

- Level k calls mergesort 2^k , each processing $n/2^k$ elements.
- At level 0, we merge 2 sorted lists, each of size $n/2$, so maximum of $n - 1$ comparisons, $O(n)$.
- At level 1, two pairs of sorted lists are merged, each of size $n/4$, for a total of $n - 2$ comparisons maximum.
- At level k we merge 2^k sorted lists, each of size $n/2^{k+1}$, for a total of at most $n - 2^k$ operations.
- linear work is done merging all the elements on each level
- The number of elements in the subproblem is halved at each level.
- So recursion goes $\log n$ levels deep, and linear work done at each level.
- Worst case: $O(n \log n)$.

- Mergesort does not rely on random access to the elements in order to sort; it moves from start to end.
- Great for linked lists.
- By rearranging pointers, no extra space is used.
- But for arrays, a third array is required to store merge results, otherwise overwriting will occur.
- http://youtu.be/XaqR3G_NVoo

Sorting and Randomness

- Pick a spot in the array. This is the pivot
- Move all elements around pivot. If smaller than pivot, to the left of it, if bigger, to the right.
- Repeat for sub-arrays to the left and right.

5	1	7	3	4	6	8	2
---	---	---	---	---	---	---	---

Sorting and Randomness

- Pick a spot in the array. This is the pivot
- Move all elements around pivot. If smaller than pivot, to the left of it, if bigger, to the right.
- Repeat for sub-arrays to the left and right.

5	1	7	3	4	6	8	2
---	---	---	---	---	---	---	---



Sorting and Randomness

- Pick a spot in the array. This is the pivot
- Move all elements around pivot. If smaller than pivot, to the left of it, if bigger, to the right.
- Repeat for sub-arrays to the left and right.

5	1	2	3	4	6	8	7
---	---	---	---	---	---	---	---


Sorting and Randomness

- Pick a spot in the array. This is the pivot
- Move all elements around pivot. If smaller than pivot, to the left of it, if bigger, to the right.
- Repeat for sub-arrays to the left and right.




Sorting and Randomness

- Pick a spot in the array. This is the pivot
- Move all elements around pivot. If smaller than pivot, to the left of it, if bigger, to the right.
- Repeat for sub-arrays to the left and right.




Sorting and Randomness

- Pick a spot in the array. This is the pivot
- Move all elements around pivot. If smaller than pivot, to the left of it, if bigger, to the right.
- Repeat for sub-arrays to the left and right.




Sorting and Randomness

- Pick a spot in the array. This is the pivot
- Move all elements around pivot. If smaller than pivot, to the left of it, if bigger, to the right.
- Repeat for sub-arrays to the left and right.




Sorting and Randomness

- Pick a spot in the array. This is the pivot
- Move all elements around pivot. If smaller than pivot, to the left of it, if bigger, to the right.
- Repeat for sub-arrays to the left and right.



Sorting and Randomness

- Pick a spot in the array. This is the pivot
- Move all elements around pivot. If smaller than pivot, to the left of it, if bigger, to the right.
- Repeat for sub-arrays to the left and right.



Sorting and Randomness

- Pick a spot in the array. This is the pivot
- Move all elements around pivot. If smaller than pivot, to the left of it, if bigger, to the right.
- Repeat for sub-arrays to the left and right.

1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---

Formalising the idea


- This algorithm (known as quick sort) is another divide and conquer algorithm.
- The partition function does an in-place partitioning in a single scan.

```
1: function quicksort( $A, l, h$ )
2:   if  $l < h$  then
3:      $p = \text{partition}(A, l, h)$ 
4:     quicksort( $a, l, p - 1$ )
5:     quicksort( $a, p + 1, h$ )
6:   end if
7: end function
```

```
1: function partition( $A, l, h$ )
2:    $p = h$            ▷ pivot element index
3:    $div = l$           ▷ divider element
4:   for all  $i = l$  to  $h - 1$  do
5:     if  $A[i] \leq A[p]$  then
6:       swap  $A[i], A[div]$ 
7:        $div++$ 
8:     end if
9:   end for
10:  swap  $A[p], A[div]$ 
11:  return  $div$ 
12: end function
```

Formalising the idea


- This algorithm (known as quick sort) is another divide and conquer algorithm.
- The partition function does an in-place partitioning in a single scan.



```
1: function partition( $A, l, h$ )
2:    $p = h$            ▷ pivot element index
3:    $div = l$           ▷ divider element
4:   for all  $i = l$  to  $h - 1$  do
5:     if  $A[i] \leq A[p]$  then
6:       swap  $A[i], A[div]$ 
7:        $div++$ 
8:     end if
9:   end for
10:  swap  $A[p], A[div]$ 
11:  return  $div$ 
12: end function
```

Formalising the idea


- This algorithm (known as quick sort) is another divide and conquer algorithm.
- The partition function does an in-place partitioning in a single scan.



```
1: function partition( $A, l, h$ )
2:    $p = h$            ▷ pivot element index
3:    $div = l$           ▷ divider element
4:   for all  $i = l$  to  $h - 1$  do
5:     if  $A[i] \leq A[p]$  then
6:       swap  $A[i], A[div]$ 
7:        $div++$ 
8:     end if
9:   end for
10:  swap  $A[p], A[div]$ 
11:  return  $div$ 
12: end function
```


Formalising the idea


- This algorithm (known as quick sort) is another divide and conquer algorithm.
- The partition function does an in-place partitioning in a single scan.



```
1: function partition( $A, l, h$ )
2:    $p = h$            ▷ pivot element index
3:    $div = l$           ▷ divider element
4:   for all  $i = l$  to  $h - 1$  do
5:     if  $A[i] \leq A[p]$  then
6:       swap  $A[i], A[div]$ 
7:        $div++$ 
8:     end if
9:   end for
10:  swap  $A[p], A[div]$ 
11:  return  $div$ 
12: end function
```

Formalising the idea


- This algorithm (known as quick sort) is another divide and conquer algorithm.
- The partition function does an in-place partitioning in a single scan.



```
1: function partition( $A, l, h$ )
2:    $p = h$            ▷ pivot element index
3:    $div = l$           ▷ divider element
4:   for all  $i = l$  to  $h - 1$  do
5:     if  $A[i] \leq A[p]$  then
6:       swap  $A[i], A[div]$ 
7:        $div++$ 
8:     end if
9:   end for
10:  swap  $A[p], A[div]$ 
11:  return  $div$ 
12: end function
```

Formalising the idea


- This algorithm (known as quick sort) is another divide and conquer algorithm.
- The partition function does an in-place partitioning in a single scan.



```
1: function partition( $A, l, h$ )
2:    $p = h$            ▷ pivot element index
3:    $div = l$           ▷ divider element
4:   for all  $i = l$  to  $h - 1$  do
5:     if  $A[i] \leq A[p]$  then
6:       swap  $A[i], A[div]$ 
7:        $div++$ 
8:     end if
9:   end for
10:  swap  $A[p], A[div]$ 
11:  return  $div$ 
12: end function
```

Formalising the idea


- This algorithm (known as quick sort) is another divide and conquer algorithm.
- The partition function does an in-place partitioning in a single scan.



```
1: function partition( $A, l, h$ )
2:    $p = h$            ▷ pivot element index
3:    $div = l$           ▷ divider element
4:   for all  $i = l$  to  $h - 1$  do
5:     if  $A[i] \leq A[p]$  then
6:       swap  $A[i], A[div]$ 
7:        $div++$ 
8:     end if
9:   end for
10:  swap  $A[p], A[div]$ 
11:  return  $div$ 
12: end function
```

Formalising the idea


- This algorithm (known as quick sort) is another divide and conquer algorithm.
- The partition function does an in-place partitioning in a single scan.



```
1: function partition( $A, l, h$ )
2:    $p = h$            ▷ pivot element index
3:    $div = l$           ▷ divider element
4:   for all  $i = l$  to  $h - 1$  do
5:     if  $A[i] \leq A[p]$  then
6:       swap  $A[i], A[div]$ 
7:        $div++$ 
8:     end if
9:   end for
10:  swap  $A[p], A[div]$ 
11:  return  $div$ 
12: end function
```

Formalising the idea


- This algorithm (known as quick sort) is another divide and conquer algorithm.
- The partition function does an in-place partitioning in a single scan.



```
1: function partition( $A, l, h$ )
2:    $p = h$            ▷ pivot element index
3:    $div = l$           ▷ divider element
4:   for all  $i = l$  to  $h - 1$  do
5:     if  $A[i] \leq A[p]$  then
6:       swap  $A[i], A[div]$ 
7:        $div++$ 
8:     end if
9:   end for
10:  swap  $A[p], A[div]$ 
11:  return  $div$ 
12: end function
```

Formalising the idea

- This algorithm (known as quick sort) is another divide and conquer algorithm.
- The partition function does an in-place partitioning in a single scan.



```
1: function partition( $A, l, h$ )
2:    $p = h$            ▷ pivot element index
3:    $div = l$           ▷ divider element
4:   for all  $i = l$  to  $h - 1$  do
5:     if  $A[i] \leq A[p]$  then
6:       swap  $A[i], A[div]$ 
7:        $div++$ 
8:     end if
9:   end for
10:  swap  $A[p], A[div]$ 
11:  return  $div$ 
12: end function
```

Running Time

- Partition: $O(n)$
- But how many partitions are made?

Running Time

- Partition: $O(n)$
- But how many partitions are made?
- Depends on the pivot element picked.
- If in the middle, then each subarray created will be half the size of its parent, and time will be $O(n \log n)$ in total.
- But what happens if pivot element is the biggest or smallest element in the subarray?
- Then pivot is placed into “correct” position, and the new subarray is 1 element smaller than the original array.
- Total tree depth will be $n - 1$, with n operations in each, worst case time is $\Theta(n^2)$
- So why is quicksort called quicksort?

Average Case Running Time

- Assume that the typical partition produces a terrible split of x/y (e.g. $x=9, y=10$ means 9 elements are always to the left or right of the partition, and only 1 on the other side). How deep is our tree?
- If our tree is n elements in size, at level 1, we have 2 subarrays of size $n(1 - x/y)$ and nx/y . The latter is bigger, and will lead to the deeper tree, so let's consider it.
- At level 2, the subarray will have $n(x/y)^2$ elements, at level 3 $n(x/y)^3$ etc, until at level h we will have 1 element.

$$n(x/y)^h = 1 \Rightarrow n = (y/x)^h$$

- $h = \log_{y/x} n$
- Remember we can change bases of logs by constant multiplication.
- As long as x and y are constants, the height of the tree is $\Theta(\log n)$
- Given n partitioning operations per level, total work is $\Theta(n \log n)$ in the average case!

More about running time

- Our claim is, that given randomly ordered data, quicksort will run in $\Theta(n \log n)$ time with a very high probability.
- For any pivot selection method, there is some worst case input which guarantees $\Theta(n^2)$ running time.
- If we use the same pivot selection method all the time, then will always fall foul of the same worst case input.
- One way to overcoming this involves randomising the pivot element.

Pivot Element Randomization

- Randomly permute the array before starting.
- At every step of the tree, select a random element to be the pivot (e.g. swap a random element from the subarray with the pivot element).
- Randomised quick sort runs in $\Theta(n \log n)$ time on any input with high probability!

An aside, Randomization

- Given a collection of n nuts and n bolts, you need to match each bolt to each nut.
- They are all a similar size, so you can't compare two nuts or two bolts.
- $O(n^2)$ brute force approach: for each nut, compare to a bolt until a match is found.
- Randomly selecting a nut means we'll work our way through half the bolts before a match is found, still $O(n^2)$.

An aside, Randomization

- Given a collection of n nuts and n bolts, you need to match each bolt to each nut.
- They are all a similar size, so you can't compare two nuts or two bolts.
- $O(n^2)$ brute force approach: for each nut, compare to a bolt until a match is found.
- Randomly selecting a nut means we'll work our way through half the bolts before a match is found, still $O(n^2)$.
- A quicksort like approach:
 - Select a nut at random, and partition bolts to those less than the nut size, and those bigger.
 - Once we find the matching bolt, partition the nuts in the same way. Total cost: $2n - 2$.
 - We can repeat this for the partitions, and the cost (as per quicksort) is $\Theta(n \log n)$ in the average case.
- No simple efficient algorithm exists for this problem!

- Mergesort, heapsort and quicksort all outperform selection sort (and insertion sort) on sufficiently large instances, as they're all $\Theta(n \log n)$
- But which is the fastest?

- Mergesort, heapsort and quicksort all outperform selection sort (and insertion sort) on sufficiently large instances, as they're all $\Theta(n \log n)$
- But which is the fastest?
- O notation doesn't tell us, implementation details are typically the most important in this case.
- Experiments have shown that a good quick sort is 2-3 times faster than the others due to the simplicity of the inner loop operations.
- <http://youtu.be/ywWBy6J5gz8>
- **Homework (to be discussed in practical):** Which type of sort should Google use?

Can we do better?


- So can we do better than $\Theta(n \log n)$?

Can we do better?

- So can we do better than $\Theta(n \log n)$?
- yes and no...


- All our sorts so far have, in one way or another, compared 2 elements and made a decision based on the difference between those two elements.
- Such comparison sorts did not use any other information about the elements (e.g. the actual value of the element).
- We can view any comparison sort abstractly in terms of decision trees.
- A decision tree is a full binary tree that represents the comparisons between elements that are performed by a sorting algorithm operating on an input of a given size.

Decision Trees




- We ignore control, datamovement etc.
- Each internal node is labelled $i : j$ for $1 \leq i, j \leq n$ for n input elements.
- Leaves are labelled with permutations of $1 \dots n$.
- Executing a sorting algorithm traces a path from the root of the tree down to a leaf.

Decision Trees



- Each internal node represents a comparison, with edges representing decisions based on this comparison.
- When a leaf is reached, the permutation shows the ordering required for the list to be sorted.


Decision Trees



• Input: 6 9 1

```
1: function InsertionSort(A)
2:   for  $i = 2$  to  $n$  do
3:      $j = i$ 
4:     while ( $A[j] \leq A[j - 1]$ ) &
        $j > 0$  do
5:       swap( $A[j], A[j - 1]$ )
6:        $j = j - 1$ 
7:     end while
8:   end for
9: end function
```

Decision Trees



- How many leaf nodes must exist?
- The leaves must represent any possible ordering of the list, so $n!$ leaves.
- Each leaf must be reachable from the root, as every possible ordering is possible.

A Theorem!

- The length of the longest simple path from the root of a tree to any of its reachable leaves represents the worst case number of comparisons that the corresponding sorting algorithm would perform.
- This is equivalent to the height of the tree.
- A lower bound on the height of any decision tree in which each permutation appears as a reachable leaf is therefore a lower bound on the running time of any comparison sorting algorithm.
- Any comparison sort algorithm requires $\Omega(n \log n)$ comparisons in the worst case.
- Why?

- Assume a tree of height h with l reachable leaves for a sort on n elements.
- Since $n!$ permutations must appear in the leaves, $n! \leq l$. Why \leq and not $=$?
- But a binary tree of height h has no more than 2^h leaves. So,

$$n! \leq l \leq 2^h$$

- Taking logarithms we obtain $h \geq \log(n!)$
- Now it can be shown that $\log(n!) = \Theta(n \log n) = \Omega(n \log n)$ □
- Since heapsort and mergesort are $O(n \log n)$ they are asymptotically optimal comparison sorts (as they match the worst case lower bound from the above).

So how can we do better?

- In order to do better, we need to find a sort that doesn't utilise comparisons.
- The counting sort assumes that all input lies in the range $0 \dots k$. If $k = O(n)$ it runs in $\Theta(n)$ time.
- Core idea is to determine for any input x , the number of elements less than x .
- This element is then placed in the correct position in the output array. E.g. if 17 elements are less than x , x will be in position 18.
- We must be able to handle elements of the same value.

Counting Sort

```
1: function CountingSort(A,k)
2:   for all i=0 to k do
3:     C[i]=0
4:   end for
5:   for all j=0 to |A|-1 do
6:     C[A[j]]=C[A[j]]+1
7:   end for
8:   for all i=1 to k do
9:     C[i]=C[i]+C[i-1]
10:  end for
11:  for all j=|A|-1 down to 0 do
12:    B[C[A[j]]-1]=A[j]
13:    C[A[j]]=C[A[j]]-1
14:  end for
15: end function
```

- Array C is used to initially store the number of times an element (used as the index) is encountered (line 6).
- Line 9 modifies C to count how many elements are less than or equal to the index.
- Line 11 populates the output array by obtaining the appropriate number from C and inserting it into the output, before decrementing the number (line 13).

Counting Sort

```
1: function CountingSort(A,k)
2:   for all i=0 to k do
3:     C[i]=0
4:   end for
5:   for all j=0 to |A|-1 do
6:     C[A[j]]=C[A[j]]+1
7:   end for
8:   for all i=1 to k do
9:     C[i]=C[i]+C[i-1]
10:  end for
11:  for all j=|A|-1 down to 0 do
12:    B[C[A[j]]-1]=A[j]
13:    C[A[j]]=C[A[j]]-1
14:  end for
15: end function
```

A

2	5	3	0	2	3	0	3
---	---	---	---	---	---	---	---

C

2	0	2	3	0	1
---	---	---	---	---	---

Counting Sort

```
1: function CountingSort(A,k)
2:   for all i=0 to k do
3:     C[i]=0
4:   end for
5:   for all j=0 to |A|-1 do
6:     C[A[j]]=C[A[j]]+1
7:   end for
8:   for all i=1 to k do
9:     C[i]=C[i]+C[i-1]
10:  end for
11:  for all j=|A|-1 down to 0 do
12:    B[C[A[j]]-1]=A[j]
13:    C[A[j]]=C[A[j]]-1
14:  end for
15: end function
```

C

2	2	4	7	7	8
---	---	---	---	---	---

Counting Sort

```
1: function CountingSort(A,k)
2:   for all i=0 to k do
3:     C[i]=0
4:   end for
5:   for all j=0 to |A|-1 do
6:     C[A[j]]=C[A[j]]+1
7:   end for
8:   for all i=1 to k do
9:     C[i]=C[i]+C[i-1]
10:  end for
11:  for all j=|A|-1 down to 0 do
12:    B[C[A[j]]-1]=A[j]
13:    C[A[j]]=C[A[j]]-1
14:  end for
15: end function
```

A

2	5	3	0	2	3	0	3
---	---	---	---	---	---	---	---

B

						3	
--	--	--	--	--	--	---	--

C

2	2	4	6	7	8
---	---	---	---	---	---

Counting Sort

```
1: function CountingSort(A,k)
2:   for all i=0 to k do
3:     C[i]=0
4:   end for
5:   for all j=0 to |A|-1 do
6:     C[A[j]]=C[A[j]]+1
7:   end for
8:   for all i=1 to k do
9:     C[i]=C[i]+C[i-1]
10:  end for
11:  for all j=|A|-1 down to 0 do
12:    B[C[A[j]]-1]=A[j]
13:    C[A[j]]=C[A[j]]-1
14:  end for
15: end function
```

A

2	5	3	0	2	3	0	3
---	---	---	---	---	---	---	---

B

	0					3	
--	---	--	--	--	--	---	--

C

1	2	4	6	7	8
---	---	---	---	---	---

Counting Sort

```
1: function CountingSort(A,k)
2:   for all i=0 to k do
3:     C[i]=0
4:   end for
5:   for all j=0 to |A|-1 do
6:     C[A[j]]=C[A[j]]+1
7:   end for
8:   for all i=1 to k do
9:     C[i]=C[i]+C[i-1]
10:  end for
11:  for all j=|A|-1 down to 0 do
12:    B[C[A[j]]-1]=A[j]
13:    C[A[j]]=C[A[j]]-1
14:  end for
15: end function
```

A

2	5	3	0	2	3	0	3
---	---	---	---	---	---	---	---

B

	0				3	3	
--	---	--	--	--	---	---	--

C

1	2	4	5	7	8
---	---	---	---	---	---

Counting Sort

```
1: function CountingSort(A,k)
2:   for all i=0 to k do
3:     C[i]=0
4:   end for
5:   for all j=0 to |A|-1 do
6:     C[A[j]]=C[A[j]]+1
7:   end for
8:   for all i=1 to k do
9:     C[i]=C[i]+C[i-1]
10:  end for
11:  for all j=|A|-1 down to 0 do
12:    B[C[A[j]]-1]=A[j]
13:    C[A[j]]=C[A[j]]-1
14:  end for
15: end function
```

B


0	0	2	2	3	3	3	5
---	---	---	---	---	---	---	---

Counting Sort


```
1: function CountingSort(A,k)
2:   for all i=0 to k do
3:     C[i]=0
4:   end for
5:   for all j=0 to |A|-1 do
6:     C[A[j]]=C[A[j]]+1
7:   end for
8:   for all i=1 to k do
9:     C[i]=C[i]+C[i-1]
10:  end for
11:  for all j=|A|-1 down to 0 do
12:    B[C[A[j]]-1]=A[j]
13:    C[A[j]]=C[A[j]]-1
14:  end for
15: end function
```

- Loop in line 5 is $O(n)$
- Loop in line 8 is $O(k)$
- Loop in line 11 is $O(n)$
- $O(2n + k) = O(n + k) = O(n)$ when $k = O(n)$
- This is a stable sort: numbers with the same value appear in the output array in the same order as they do in the input array.

A challenge!




A challenge!



IBM Model 82 card sorter, capable of sorting 650 cards/minute!

A challenge!



- Scanning only one column at a time, how can we sort these punch cards quickly?
- We only have 13 bins, and someone must handle intermediate piles.

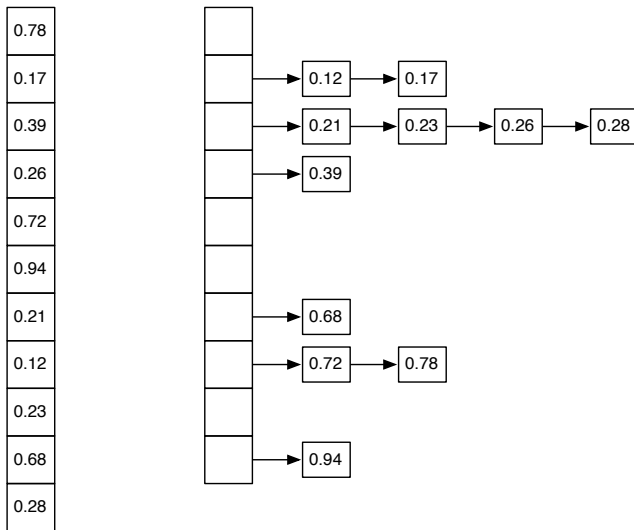
- Most intuitive approach is to sort most significant digit first, then 2nd most significant digit, etc.
- But this requires many intermediate piles.
- The radix sort sorts on the least significant digit.
- Cards are then gathered from various bins in order and sort proceeds on 2nd least significant digit, etc.
- Given d digits n cards and k possible values per digit, and using Counting sort ($O(n + k)$ to do the sorting), radix sort is $\Theta(d(n + k))$. If d is a constant, and $k = O(n)$, then sort proceeds in $\Theta(n)$

```
1: function RadixSort( $A, d$ )  
2:   for all  $i=1$  to  $d$  do  
3:     use a stable sort to sort  $A$  on digit  $d$   
4:   end for  
5: end function
```

- **Homework (for practical):** Prove that it works!
- There are several tweaks that can be performed to lower the constants hidden in the running time.
- If Radix sort is $\Theta(n)$ and comparison sorts are (at best) $\Theta(n \log n)$, why not always use Radix sort?

- The counting sort works when we have a small range of numbers.
- What if we have an arbitrary range instead $0 \dots \text{max}$?
- We can create m buckets:
 $[0..\text{max}/m], [\text{max}/m..2\text{max}/m], \dots [(m-1)\text{max}/m..\text{max}]$
- We then distribute the n inputs into the buckets.
- We can then take the numbers out of the buckets in order, sorting the numbers in the buckets with another sort (e.g. heapsort) if needed.
- If the items are uniformly distributed and $m = O(n)$ then there will be very few buckets containing multiple elements, sort will be $\Theta(n)$.
- But non-uniform distributions can be relatively common (see Skiena for an example).

Bucket Sort



Where are we?

- We've now covered different sorting algorithms in detail.
- Comparison sort ($\Theta(n \log n)$)
 - Selection sort, Insertion sort, mergesort, heapsort, quick sort, ...
- Non-comparison sorts ($\Theta(n)$)
 - Counting sort, radix sort, bucket sort
- Different applications require different sorting algorithms.
- We've been able to analyse the complexity of many of these algorithms. Next, we will look at how to perform the analysis in a more principled manner.