

CS2521: Graphs


N. Oren

`n.oren@abdn.ac.uk`

University of Aberdeen

- What is a graph?
- How do we encode graphs, and why choose one encoding over another?
- Graph related concepts
- How do we use graphs?

A bit of history - The bridges of Königsberg



Can you find a way to walk through the city, going over each bridge only once?


A bit of history - The bridges of Königsberg

- Simplify the problem:
 - All we care about is being on the same landmass as the start of the next bridge we're crossing.
 - So we can remove roads.



A bit of history - The bridges of Königsberg

- The abstract structure on which the analysis is based is called a graph.



- A graph consists of a set of nodes or vertices, and edges.
- Rather than exhaustively searching for a solution, we can formalise the problem and answer it by reasoning about the graph's properties.


A bit of history - The bridges of Königsberg

- This analysis was performed by Leonhard Euler in 1736.
 - Consider a point not at the start or end of the tour. We must have reached the landmass, and also leave it.
 - So in the middle, we need an even number of bridges (so as not to get stuck).
 - At the start, we can leave but not return, and similarly at the finish.
 - So we can also have zero or two landmasses with an odd number of bridges.
 - Why can't we end up with one landmass with an odd number of bridges?
 - But problem has odd number of bridges, so no solution exists.
- Generally, the problem is to find a sequence of adjacent edges, starting and finishing at the same vertex, using each edge exactly once. This is called an Eulerian circuit.


- A directed graph is a pair (V, E) where V is a set of vertices, and E is a set of edges such that for any edge $(s, t) \in E$, $s, t \in V$.
- Note
 - The vertices within the edge represent the edge's ends.
 - For any edge, it could be the case that $s = t$, in which case a loop exists.

- An undirected graph is a pair (V, E) where V is a set of vertices, and E is a set of edges such that for any edge $\{s, t\} \in E$, $s, t \in V$.
- Note
 - Now the edge is a set of two vertices, meaning we cannot distinguish a start vertex from an end vertex.


Examples



Examples




Examples




Examples




Other types of graphs




undirected




directed




unweighted




weighted




simple




non-simple




sparse




dense




cyclic




acyclic




embedded




topological




explicit



implicit



unlabeled



labeled

Other types of graphs

- Simple: no multi-edges or loops (makes things simpler to implement)
- Sparse: only a small fraction of vertex pairs (edges) exist.
- Acyclic: no cycles
- Embedded: geometric position of vertices and edges have meaning. Otherwise, topological
- Implicit/explicit: Is the graph fully represented, or is it conceptual?
- Labelled/Unlabelled: Each vertex is associated with a unique identifier.

- Imagine you were building a social network application.
 - How are graphs useful in such a setting?
- What are the properties of these graphs?

- Imagine you were building a social network application.
 - How are graphs useful in such a setting?
- What are the properties of these graphs?
 - Is friendship two way? (directed/undirected)
 - Are you your own friend? (simple)
 - Can you have different types of friends (multi-edge)
 - How good a friend are you? (weighted)
 - Who has the most friends? (degree of a vertex)


Graph Representation

- We can store a graph in different ways.
- Most commonly, we use either an adjacency matrix or adjacency list.

Adjacency Matrices

- Given a graph (V, E) , create a $|V| \times |V|$ matrix of 0s, and place a 1 at position i, j if $(i, j) \in E$

	1	2	3	4	5
1	0	1	0	0	1
2	1	0	1	1	1
3	0	1	0	1	0
4	0	1	1	0	1
5	1	1	0	1	0



- How easy is it to check if (i,j) is an edge?

Adjacency Matrices

- How easy is it to check if (i,j) is an edge? $O(1)$
- How easy is it to update or delete an edge?

Adjacency Matrices

- How easy is it to check if (i,j) is an edge? $O(1)$
- How easy is it to update or delete an edge? $O(1)$
- So what's the downside?



- How easy is it to check if (i, j) is an edge? $O(1)$
- How easy is it to update or delete an edge? $O(1)$
- So what's the downside?
- Unless your graph is dense, most entries of the matrix will be 0. Nevertheless, you need $O(|V|^2)$ memory to store the graph.
- Facebook: $\sim 10^9$ users. Memory requirements: 10^{18} (a petabyte is 10^{15}).

Dense vs Sparse

- A complete graph has every vertex connected to every other.
- If most vertices are not connected, the graph is sparse.
- A graph is dense if it is not sparse.
- An array is sparse if most entries have identical (default) values (normally 0).
- Many CS problems involve dealing with sparse data.

Adjacency Lists

- For each vertex, store a list of vertices to which it connects.



- Most common implementations use a (singly) linked list.
- Linked lists
 - A set of nodes containing a value and a pointer to the next node.
 - A special null pointer signifies the end of the list.

Adjacency Lists

- How do we check if (i, j) is an edge?

- How do we check if (i, j) is an edge?

L = linked list for vertex i

repeat

l = Get next element of L

if $l == j$ **then**

return true

end if

until $l = \text{null}$

return false

- If v is the vertex with highest degree (say d) for which w is not present, then complexity is $O(d)$.

Adjacency List vs Adjacency Matrix

- Print all edges of a graph

Require: L elements in the adjacency list

```
function PrintGraph(L)
  for all  $l \in L$  do
    while  $l \neq \text{null}$  do
       $m = \text{succ}(l)$ 
      print( $l, m$ )
       $l = m$ 
    end while
  end for
end function
```

Require: M adjacency matrix


```
function PrintGraph(M)
  for all  $c \in |M|$  do
    for all  $r \in |M|$  do
      if  $M[c, r] == 1$  then
        print( $c, r$ )
      end if
    end for
  end for
end function
```

Comparison

Comparison	Winner
Faster to test if (x, y) is in graph?	adjacency matrices
Faster to find the degree of a vertex?	adjacency lists
Less memory on small graphs?	adjacency lists $(m + n)$ vs. (n^2)
Less memory on big graphs?	adjacency matrices (a small win)
Edge insertion or deletion?	adjacency matrices $O(1)$ vs. $O(d)$
Faster to traverse the graph?	adjacency lists $\Theta(m + n)$ vs. $\Theta(n^2)$
Better for most problems?	adjacency lists


Given two vertices u, v .

- If there is an edge from u to v , then it is incident from, or leaves u , and is incident to, or enters v . Here, v is adjacent to u .
- The number of edges entering a vertex is the vertex's in-degree, the number of edges leaving, its out-degree.
- The degree of the vertex is the sum of its in-degree and out-degree.
- For undirected graphs, the degree is the number of edges attached to the vertex, and adjacency is symmetric.



In-degree 3
Out-degree 2
Degree 5


- A path of length k from a vertex i to a vertex j — in a given graph — is a sequence of vertices $\langle v_0, \dots, v_k \rangle$ such that $v_0 = i$ and $v_k = j$ and (v_{i-1}, v_i) is an edge, for all $i = 1, \dots, k$.
- The path is said to contain vertices v_0, \dots, v_k , and edges $(v_0, v_1), \dots, (v_{k-1}, v_k)$.
- If a path exists between i and j , then j is reachable from i .
- A path is simple if all of its vertices are distinct.
- $\langle v_i, \dots, v_j \rangle$ is a subpath of $\langle v_0, \dots, v_k \rangle$ iff $0 \leq i \leq j \leq k$.



- A path is a cycle if $v_0 = v_k$ and it has at least one edge.
- A self-loop is a cycle of length 1.
- It is a simple cycle if all vertices v_1, \dots, v_k are distinct. Note that a simple cycle is the smallest possible cycle.
- A graph with no cycles is acyclic.

Connectedness

- An undirected graph is connected there is a path from every vertex to every other vertex.
- Connected components are (intuitively) those vertices that are reachable from one another.



- A binary relation on a set X is a set $R \subseteq X \times X$.
 - The "less than" relation is a binary relation on the real numbers, i.e., $< \subset \mathbb{R} \times \mathbb{R}$. Thus, $(4, 5) \in <$ but $(5, 4)$ and $(5, 5) \notin <$.
- A relation R is an equivalence relation if, for all $x, y, z \in X$, the relation satisfies

reflexivity: $(x, x) \in R$

symmetry: if $(x, y) \in R$, then $(y, x) \in R$

transitivity: if (x, y) and $(y, z) \in R$, then $(x, z) \in R$.

- Example: "has the same birthday as". Non-example: \leq .

- For an $x \in X$, its equivalence class is those elements $y \in X$ which satisfies the equivalence relation.

$$[x] = \{y \in X \mid (x, y) \in R\}$$

- Example: All people whose birthday falls on the same day, for every day of the year.
- Connected components are the equivalence classes of vertices under the reachable from relation (i.e., where there is a path between them).

Connectedness in directed graphs



- A directed graph is strongly connected if every two vertices are reachable from each other.
- The strongly connected components are the equivalence classes of vertices that are mutually reachable.

Subgraphs

- Given a graph (V, E) , we can identify a subgraph $S = (V', E')$ where


$$V' \subseteq V \quad \text{and} \quad E' \subseteq E \cap V' \times V'$$

- S is the subgraph induced by V' iff $x, y \in V'$ and $(x, y) \in E$, then $(x, y) \in E'$





Weighted Graph

- A weighted graph has a value (typically a real or integer number) associated with it, this is the edge's weight.



Trees

- A free tree is an acyclic, connected undirected graph.
- A forest is an acyclic undirected graph.



- A rooted tree is a free tree with a single distinguished vertex called the root.
- A node v on the path between the root and node w is an ancestor of w , and w is a descendant of v . Note, if $v \neq w$, then we have a proper ancestor or descendant.
- If (v, w) is an edge in the tree, with v the ancestor of w , then v is the parent of w , and w is v 's child.
- Nodes sharing a parent are called siblings.
- Child-free nodes are leaves; non-leaf nodes are internal.
- The number of children a node has are referred to as the node's degree.

- The path length from root to node is that node's depth.
- The tree's height is the maximum depth of nodes in the tree.
- A tree is ordered if the children of each node are ordered (i.e., one can refer to the first, second, etc child of a node).

- A binary tree is a finite set of nodes and consists of either
 - ① An empty set of nodes; or
 - ② Three disjoint sets of nodes: a singleton set containing a root node; a binary tree called the left subtree and a binary tree called the right subtree.
- Since the sets are disjoint, each child must be in either the left, or right tree (and is called a left child or right child).
- A binary tree is a full binary tree if no nodes have a single child.
- Definition can be extended to k -ary trees, where each node has at most k children, and each child is labelled.

Free Tree Properties

If $G = (V, E)$ is an undirected graph, then the following properties are equivalent.

- ① G is a free tree
 - ② Any two vertices in G are connected by a unique simple path
 - ③ G is connected, but if any edge is removed from E , the resulting graph is disconnected
 - ④ G is connected, and $|E| = |V| - 1$
 - ⑤ G is acyclic, and $|E| = |V| - 1$
 - ⑥ G is acyclic, but if any edge is added to E , the resulting graph contains a cycle.
- For a set X , $|X|$ is the cardinality of the set. For finite sets, this is the number of elements in the set.

Free tree properties

- We won't prove all these properties (See Cormen for details), but consider $6 \rightarrow 1$
- Suppose that G is acyclic, but that adding any edge to E creates a cycle. We must show that G is connected.
- Let u and v be arbitrary vertices in G .
- Since adding the edge creates a cycle, if we assume that u and v are not adjacent, adding such an edge results a cycle whose edges, except for (u, v) belong to G .
- Thus, there must already be a path from u to v .
- Since u, v were chosen arbitrarily, G is connected.

Where are we?

- We've introduced graphs - structures made up of nodes and edges.
- Considered how graphs can be represented.
- Examined different types of graphs (weighted, directed, undirected, trees).
- Defined a variety of concepts over graphs.
- We are now able to "speak in the language of graphs", allowing us to examine algorithms which manipulate and use graphs for various purposes.