

# Operating System Architecture

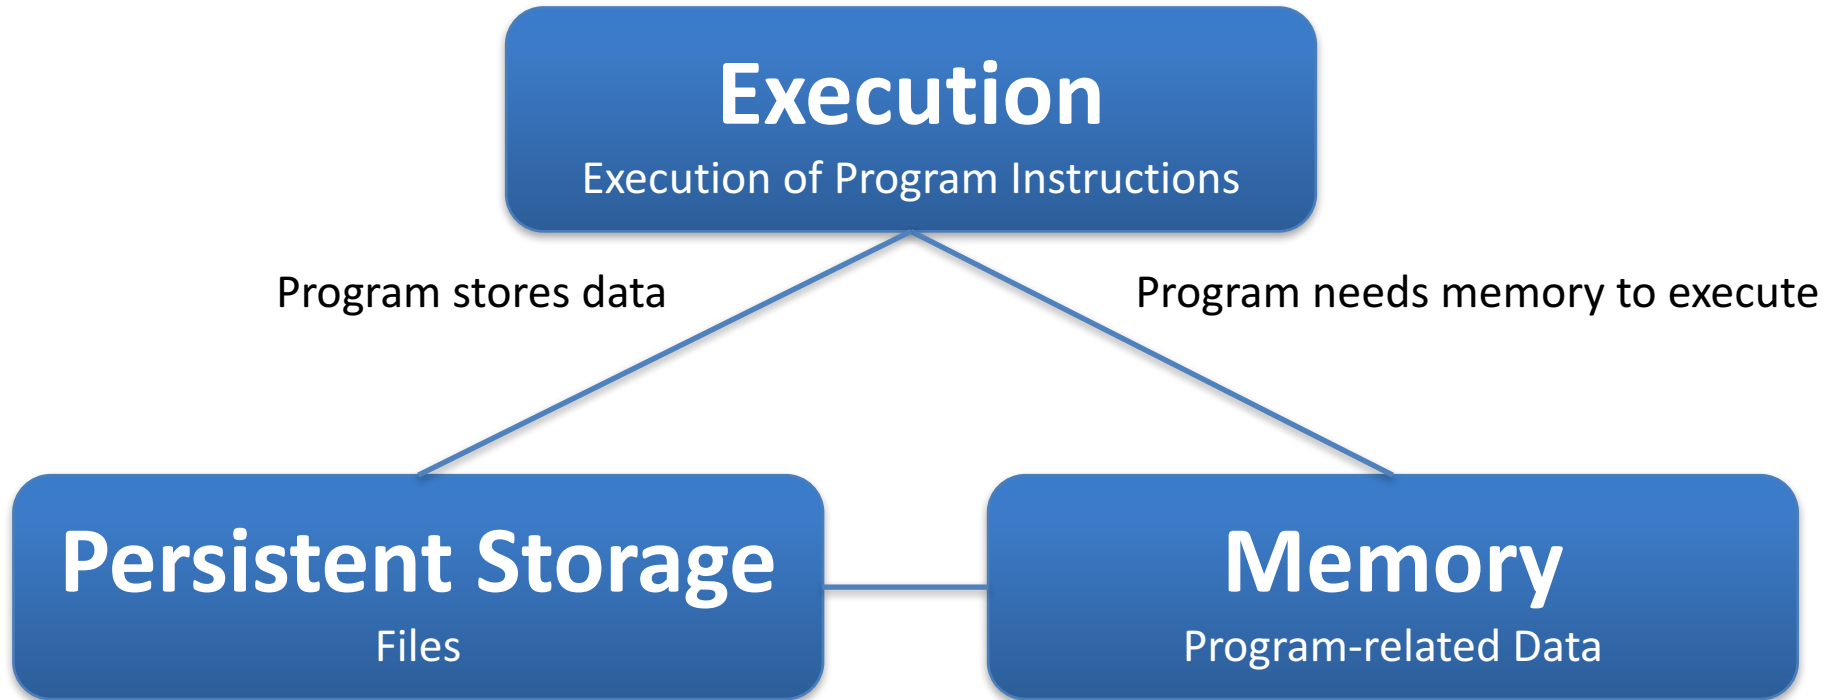
CS3026 Operating Systems

Lecture 03

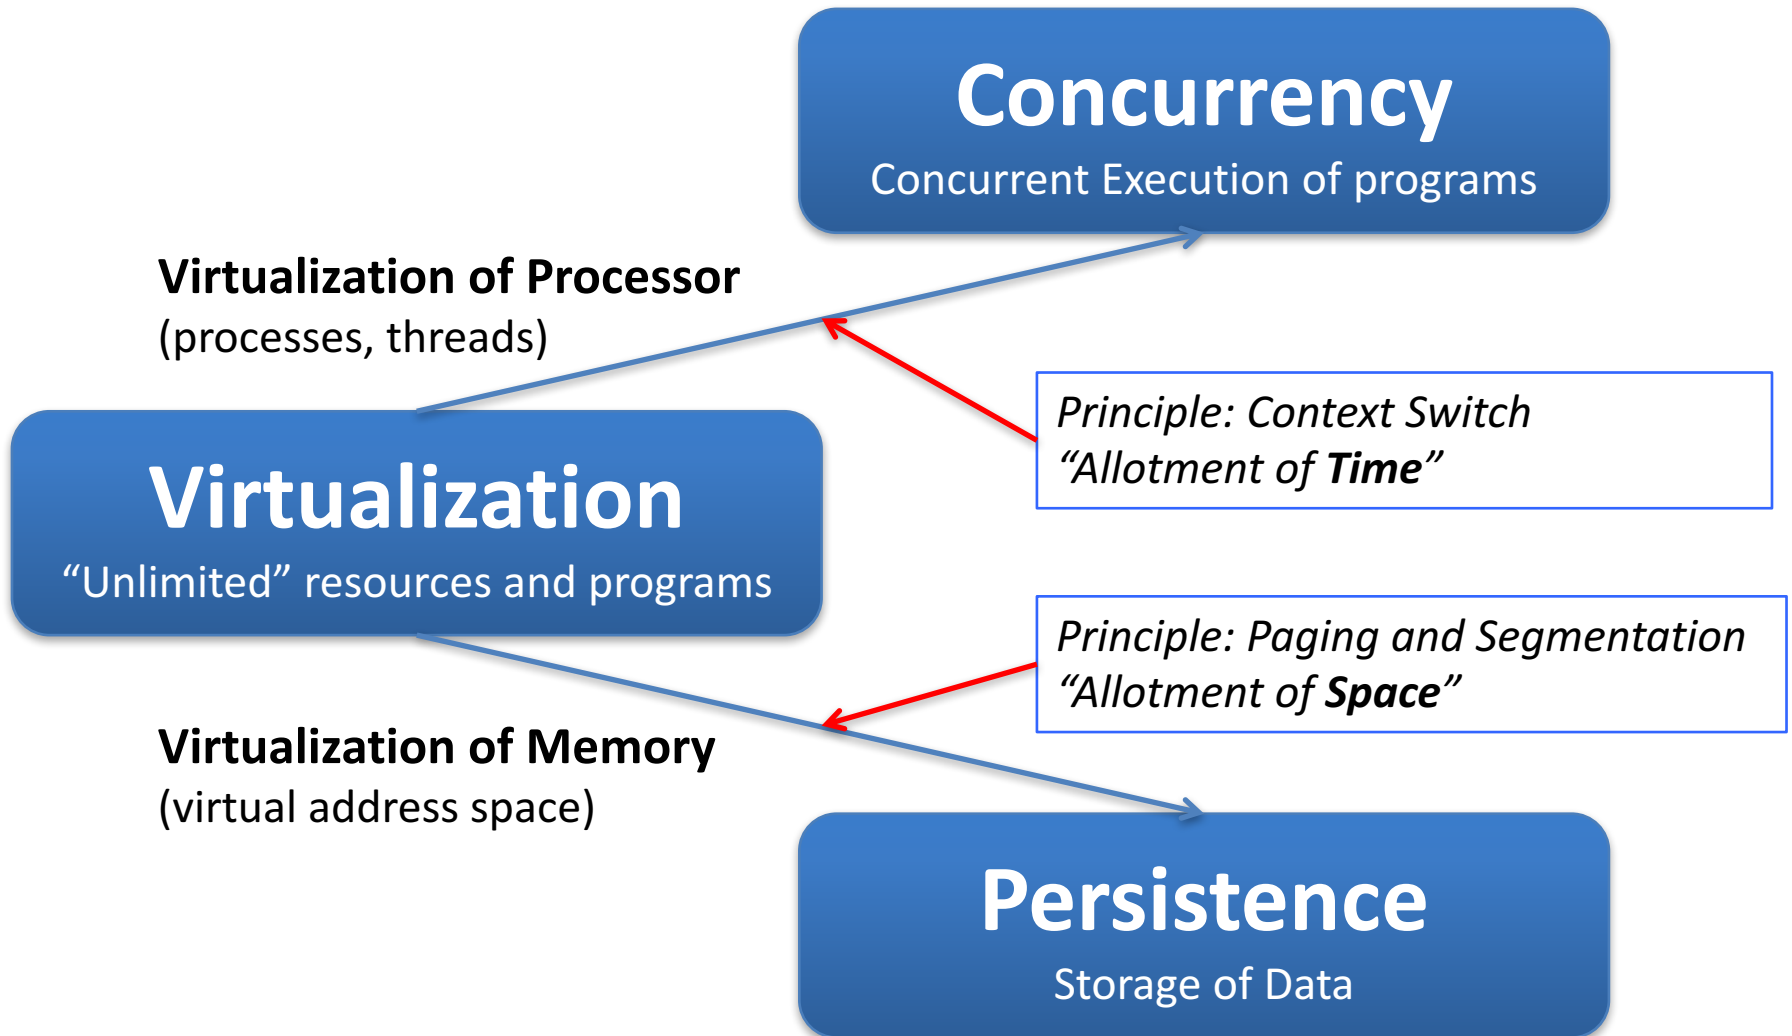
# The Role of an Operating System

- Service provider
  - Provide a set of services to system users
- Resource allocator
  - Exploit the hardware resources of one or more processors and allocate it to user programs
- Control program
  - Control the execution of programs and operations of I/O devices
    - interrupt them to send/receive data via I/O or to re-allocate hardware resources to other user programs
- Protection and Security
  - Protect multiple programs running from each other
  - Secure user access to data and define ownership of files and processes

# Operating System Functions



# Core Concepts



# Operating Systems

|                | Execution  | Memory  | Storage   |
|----------------|--|---|---|
| Virtualization | Process, thread<br>Context switch<br>Process Control<br>Block<br>Swapping    | Virtual memory<br>management<br>Segmentation<br>Free space<br>management<br>Paging<br>Page table<br>TLB | Storage volumes<br>File system                    |
| Concurrency    | Mutual exclusion<br>Locking<br>Condition variables<br>Semaphores<br>Deadlock | Shared memory   | Locking   |
| Persistence    |  |   | Files and directories<br>I/O Devices<br>Hard disk |

# Operating System Structure

The Kernel

# Protecting the Operating System

## Modes of Operation

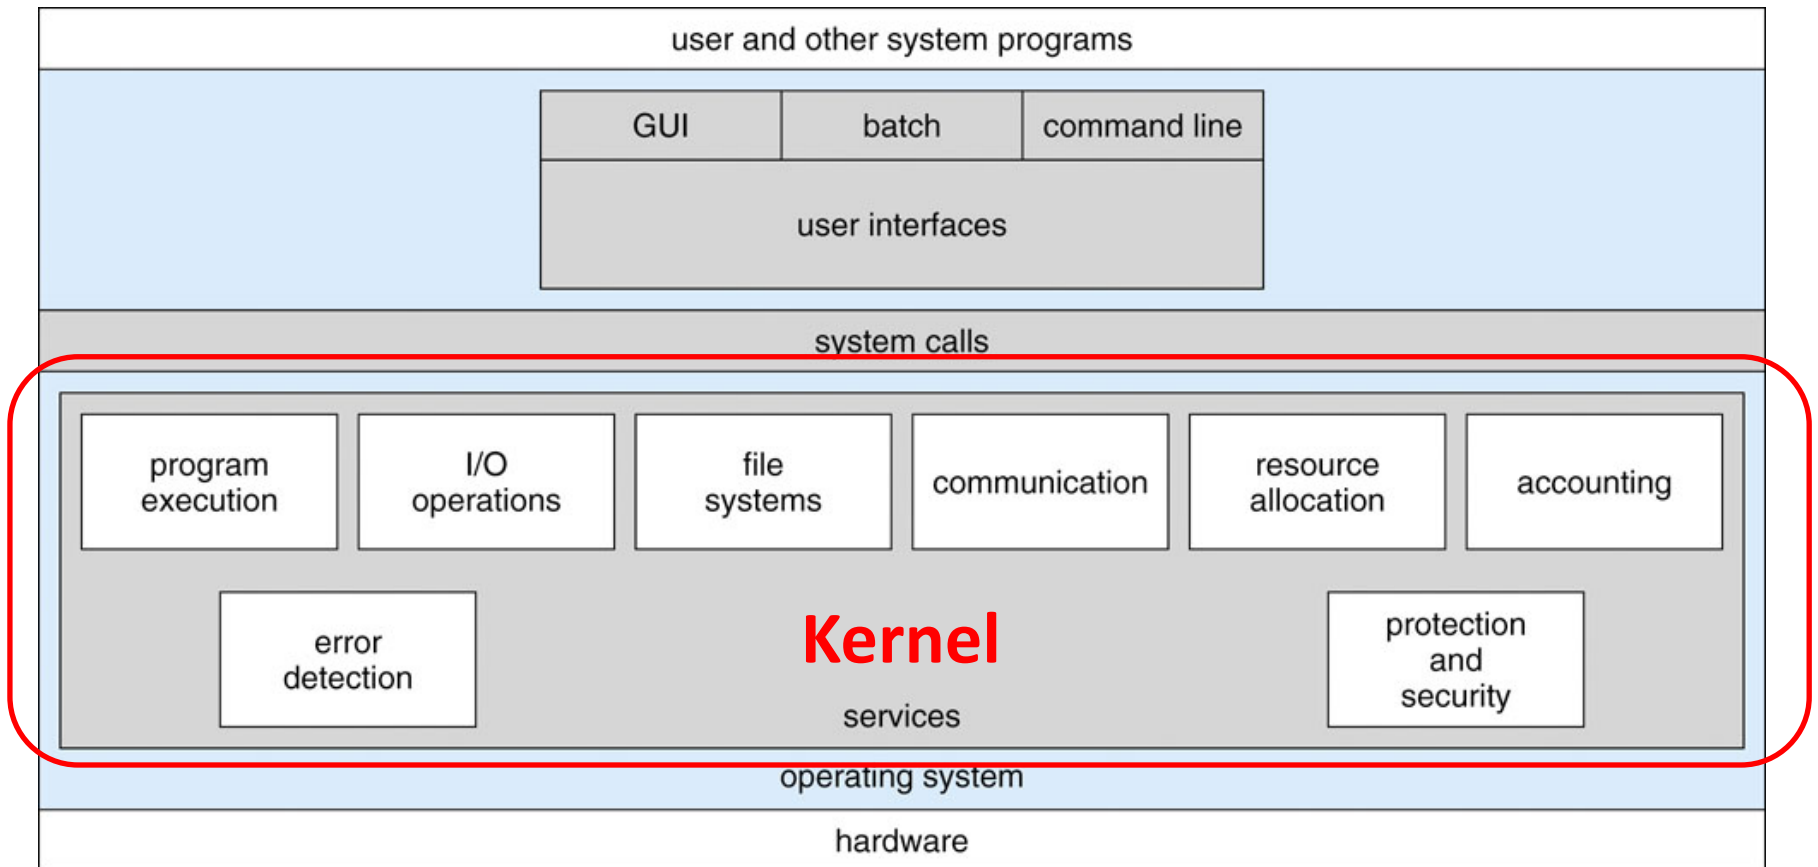
- User Mode

- User programs execute in *user mode*
- Certain areas are protected from user access
- Certain instructions may not be executed

- Kernel Mode

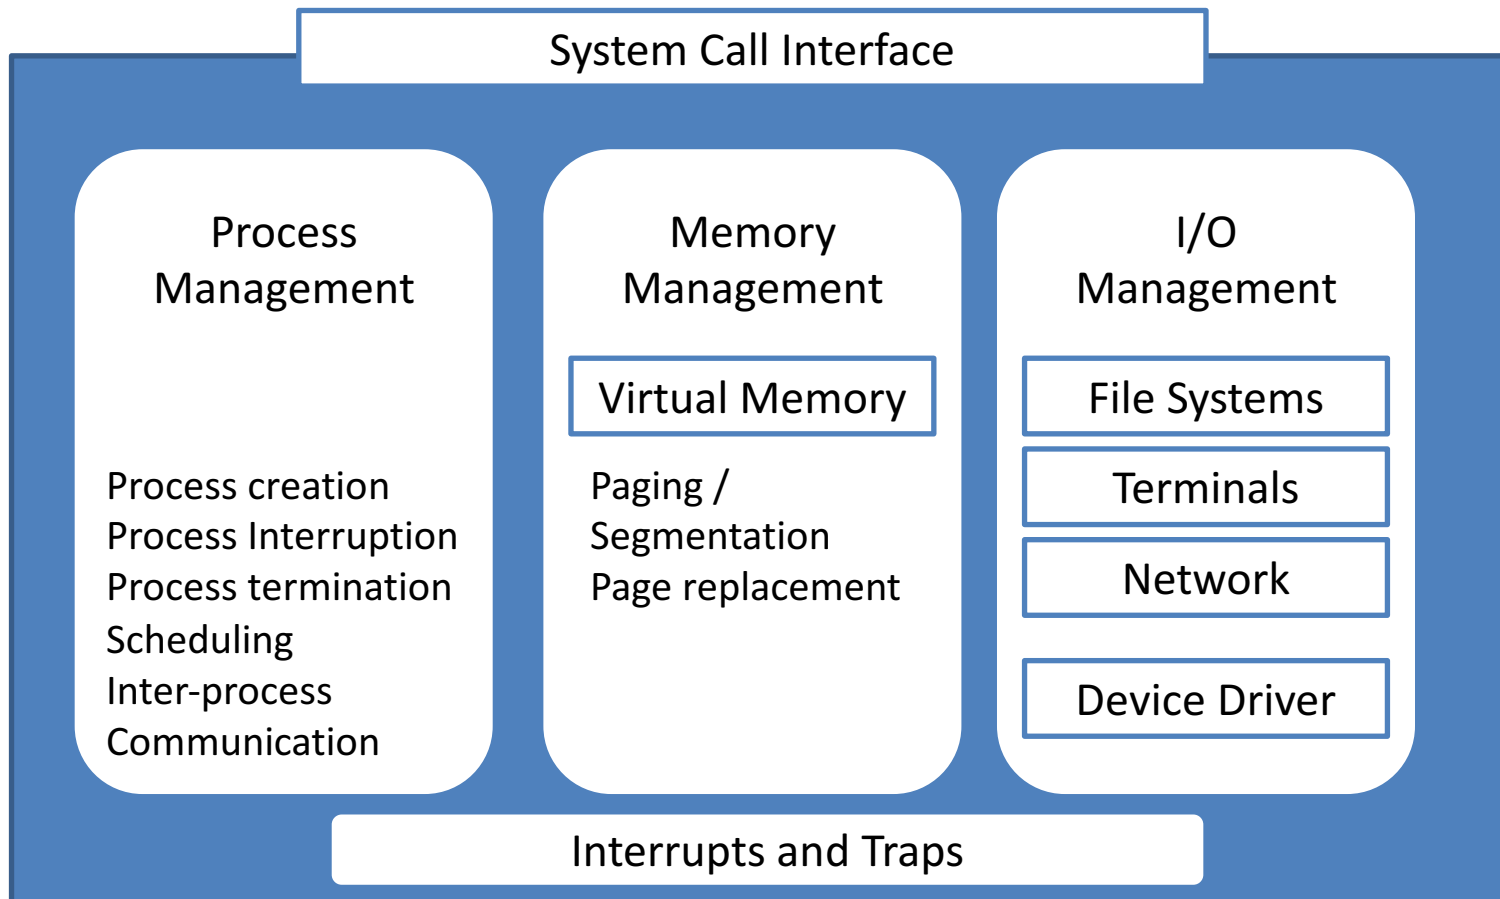
- Operating system executes in *kernel mode*
- Privileged instructions may be executed
- Protected areas of memory may be accessed

# Operating System





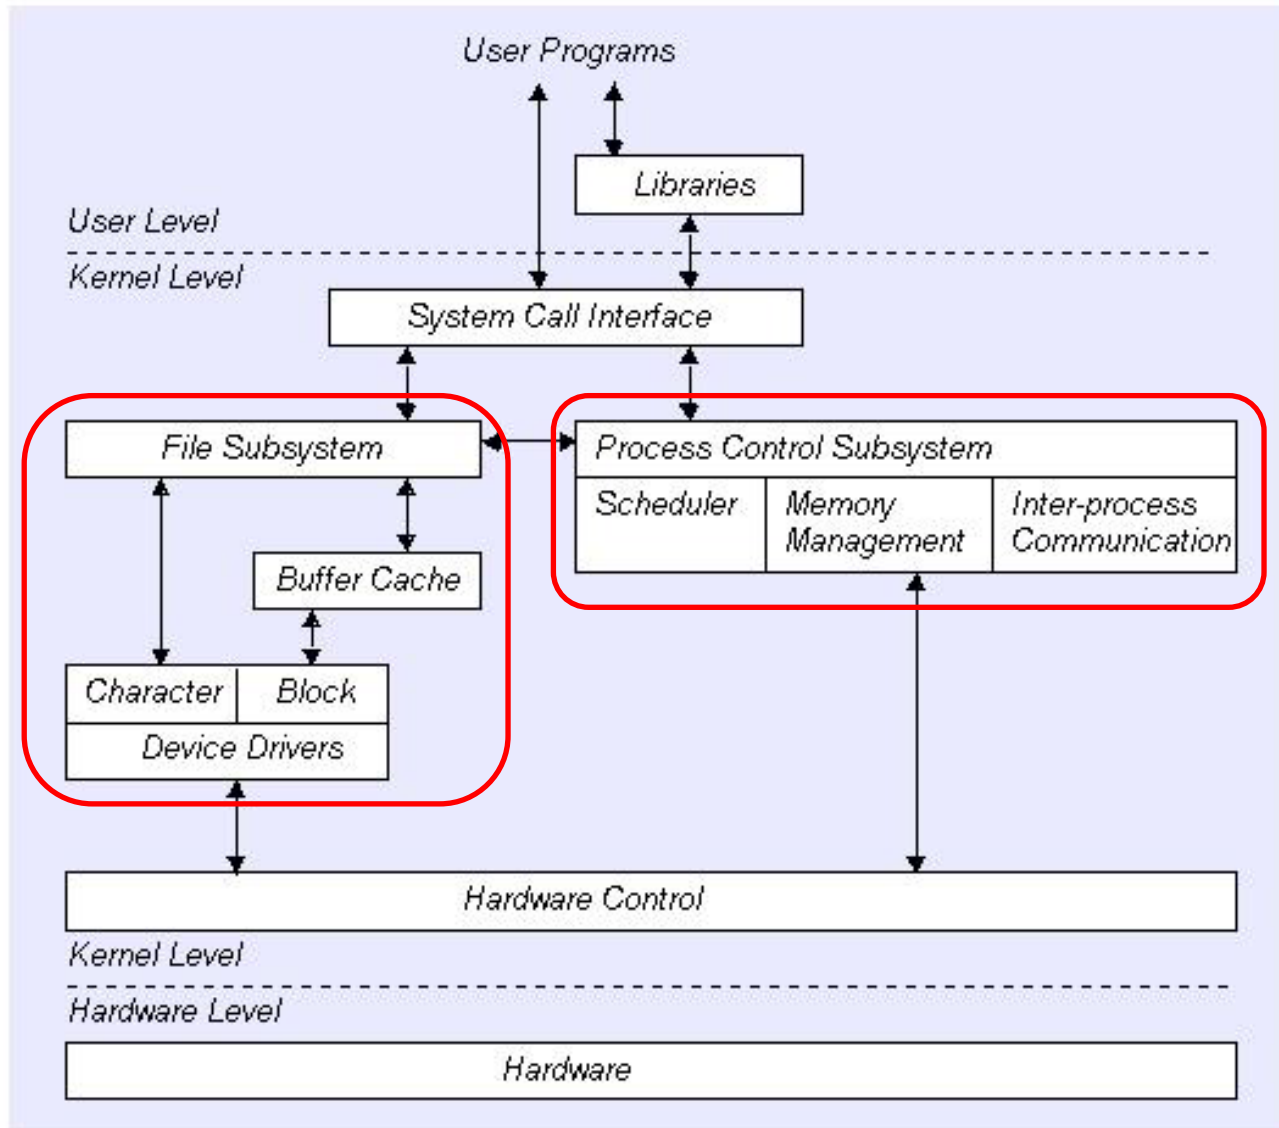
# Basic Components of Unix Kernel



# Utilising Hardware Resources

- Hardware supports
  - Basic instruction execution
  - Interrupt handling
  - Basic memory addressing mechanisms
  - User / kernel mode operation for protecting resources
- Operating system manages software constructs based on these hardware services
  - Process management
  - Virtual memory management
  - File storage and communication management

# Traditional Unix Kernel



# System Calls

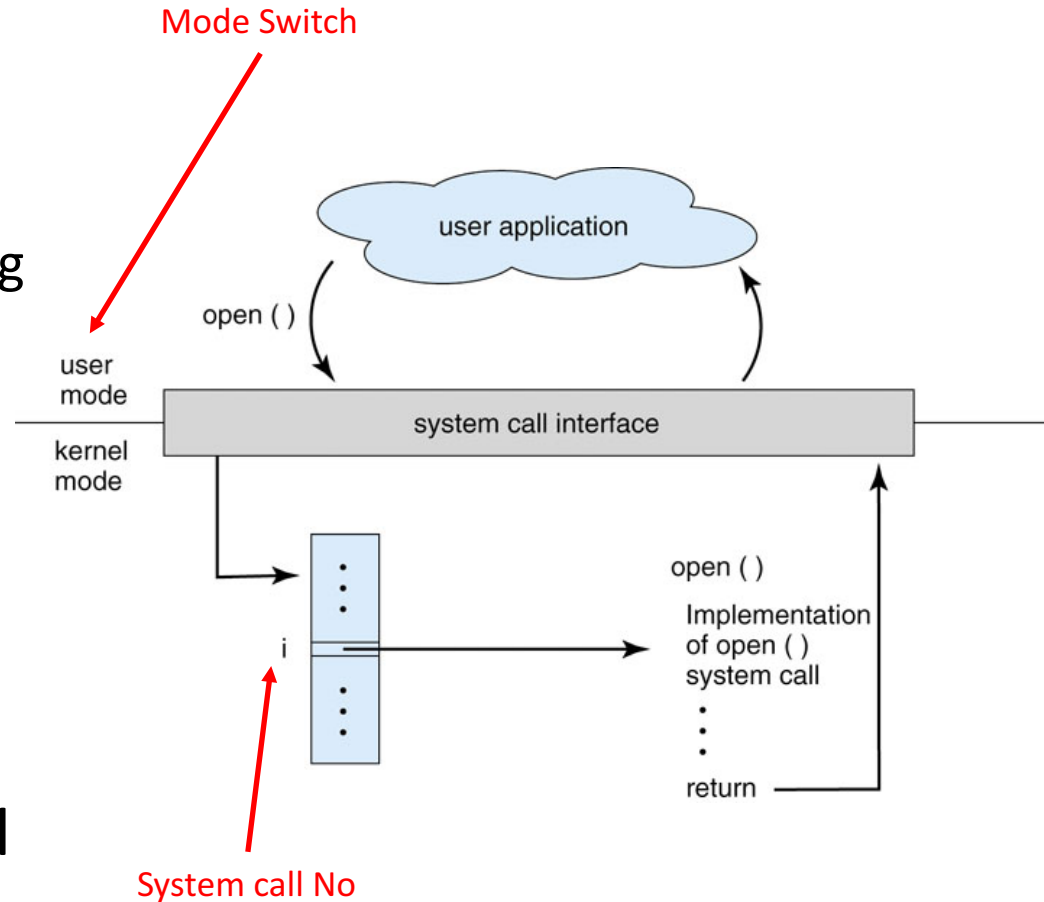
Calling Operating System Functions

# System Calls

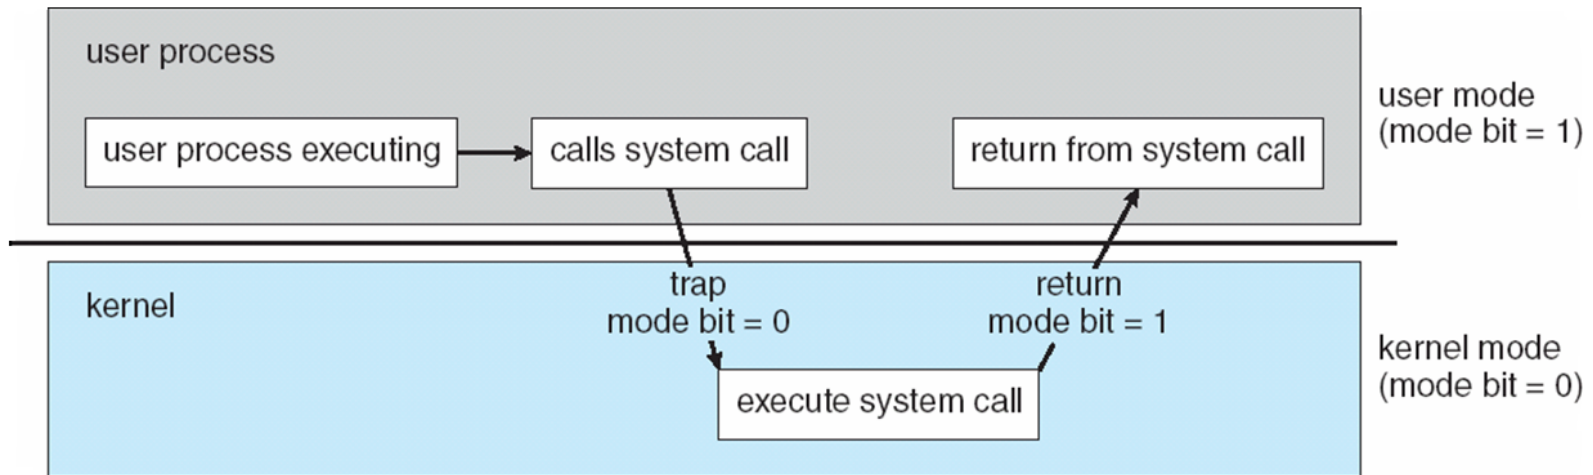
- System calls are the only entry point into the kernel
- Categories
  - Process management
  - Memory management
  - File management
  - Device management
  - Communication
- System calls are executed in *kernel mode*

# System Calls

- Interface between a program and the operating system kernel
  - Provide access to operating system services
- Is an explicit request to the kernel made via a software interrupt
- Executed in ***kernel mode***
  - Requires a mode switch
- Each system call is identified by a system call number



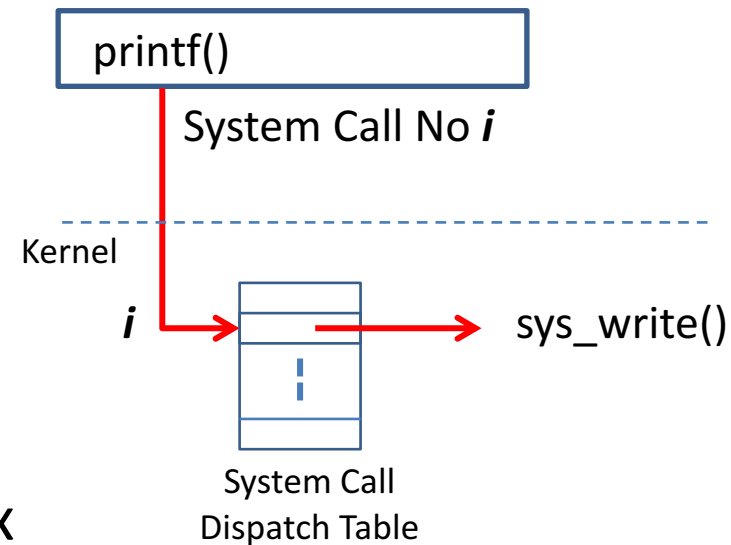
# Software Interrupt for System Calls



- Software interrupt due to system calls
  - Mode switch of processor hardware: System calls only allowed to execute in kernel mode
  - Mode bit managed by the processor
    - Provides the ability to distinguish between user and kernel mode
    - Privileged instructions only in kernel mode

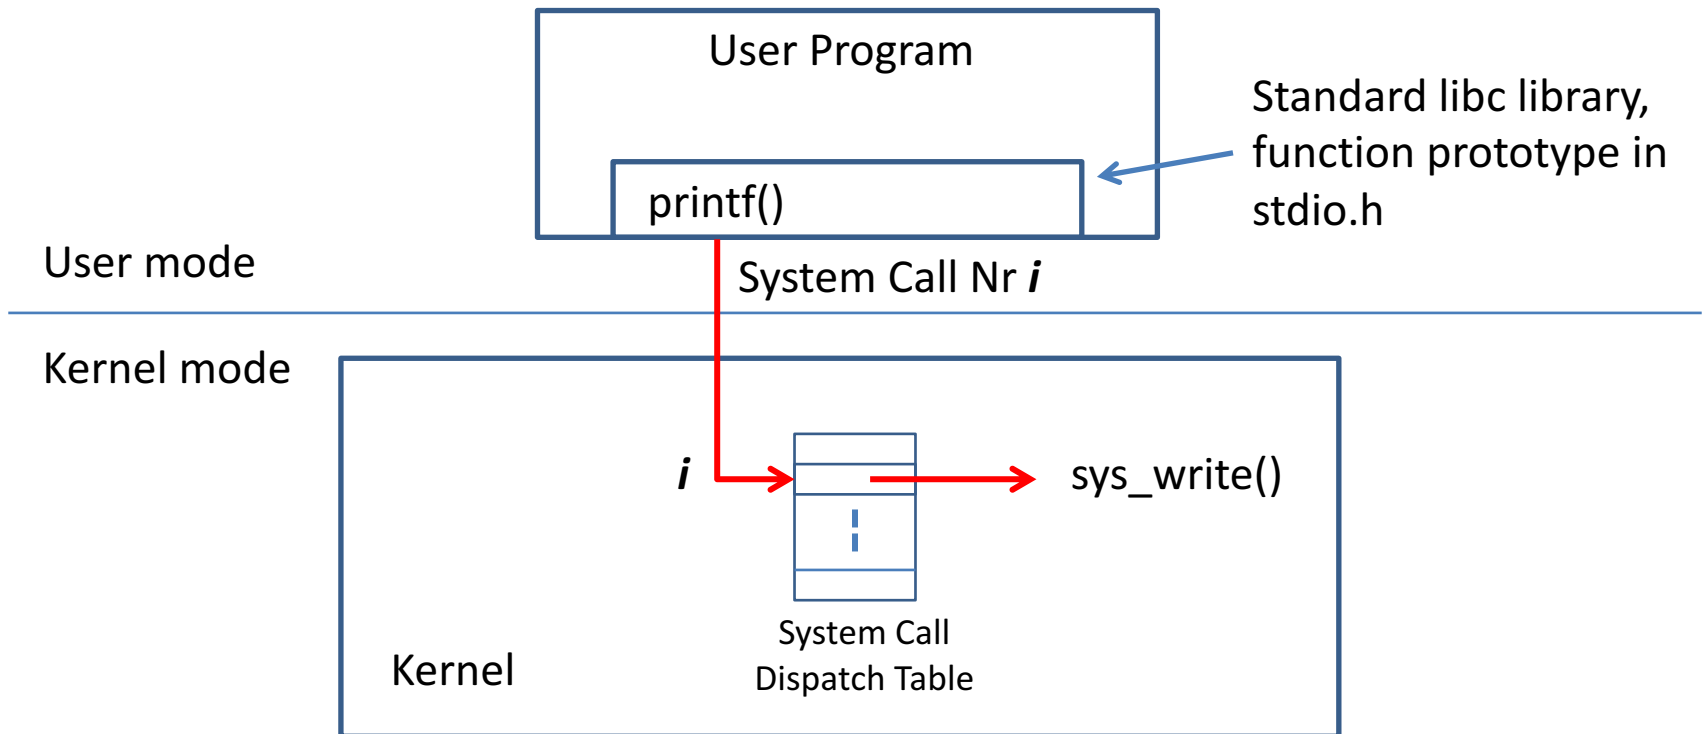
# Invoking a System Call

- Typically, a number is associated with each system call
  - The process invoking the system call must pass the *system call number* to the kernel to identify the corresponding system call service routine
- Operating system maintains a table of pointers to system call service routines, system call number is index for this table
- Operating system handles the invocation of the service routine and any return status / values



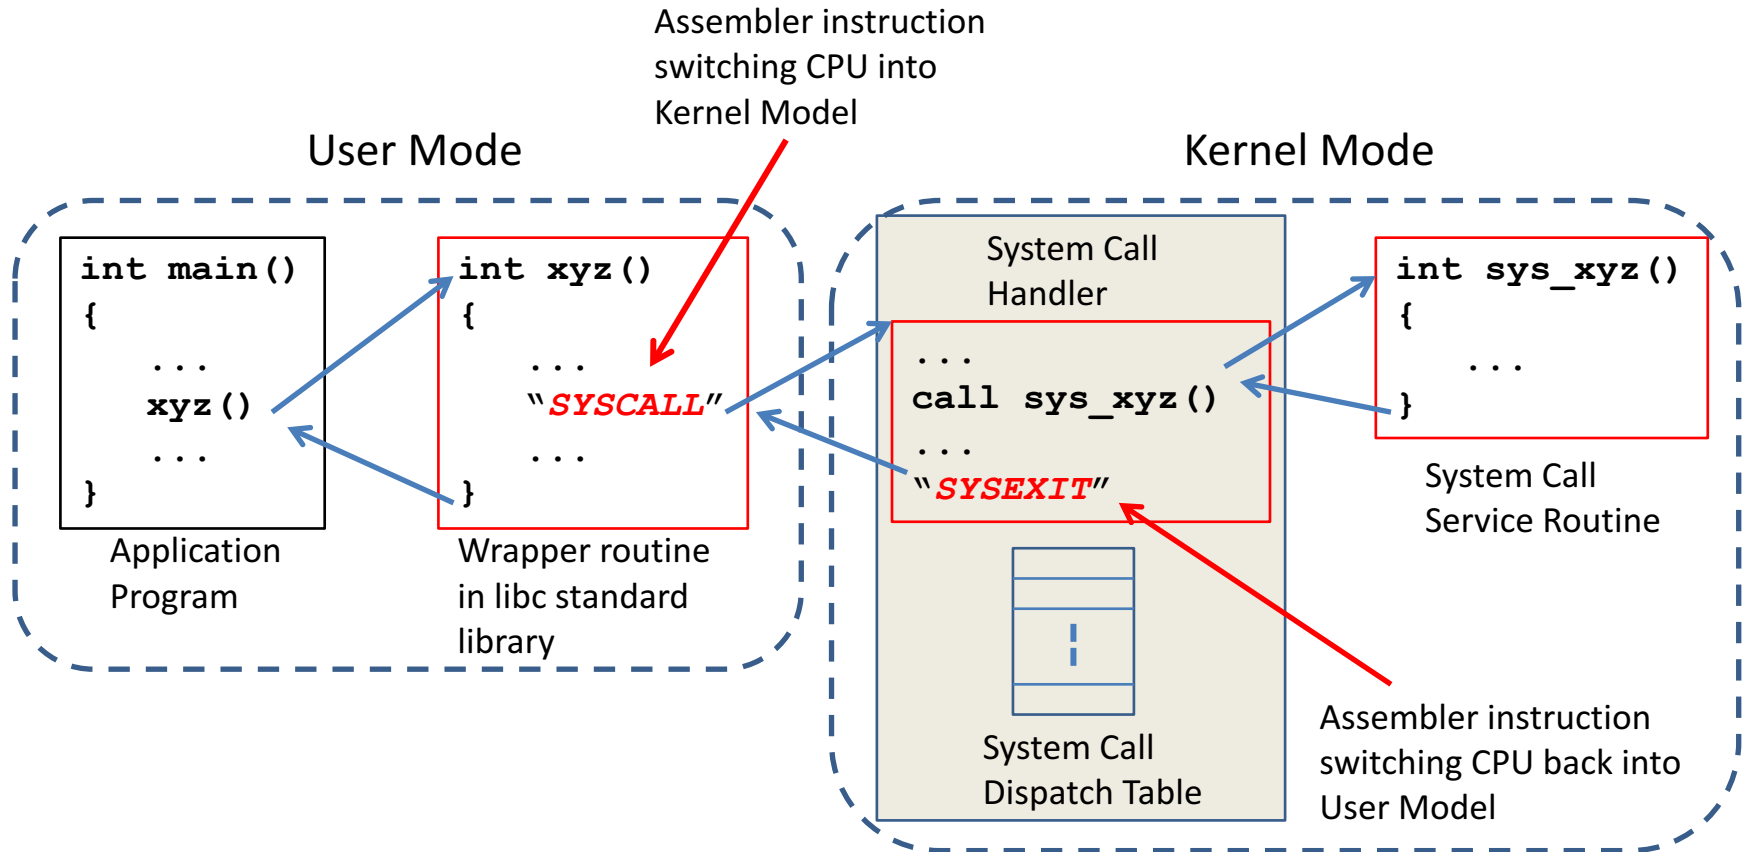


# System Call Handling



- When a process running in user mode invokes a system call, the CPU switches to kernel mode and starts the execution of a kernel function

# Invoking a System Call in Linux




- Linux implements a system call handler to manage the invocation of system call service routines
- System Dispatch Table holds all the service routine addresses
  - System Call number is index into this table

# Passing Parameters

- Three general methods:
  - Pass via CPU registers
  - Use a memory block:
    - Store parameters in memory in a table or memory block
    - Pass address of this memory block via CPU register to service routine
    - This approach is taken by Linux and Solaris
  - Use a stack
    - User program pushes parameters onto stack
    - System service routine pops parameters from stack

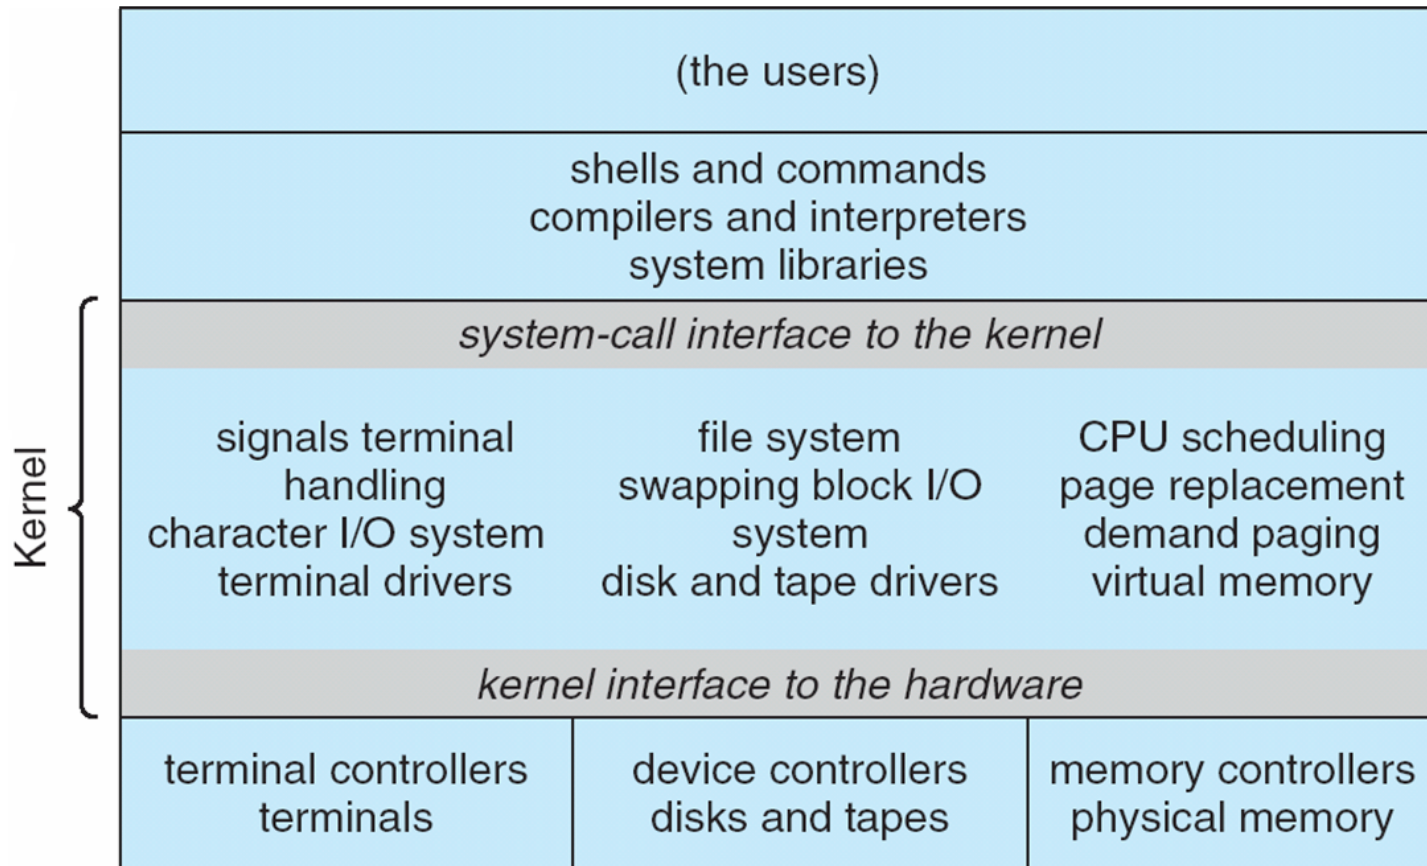
# API's and System Calls

- Operating systems usually come with a library that implements an API of functions wrapping these system calls:
  - Typically written in a high-level language (C or C++)
    - Standard C Library
    - Unix / Linux: libc or glibc
  - Usually, each system call has a corresponding wrapper routine, which an application programmer can use in their programs
    - E.g.: printf()
- POSIX is a standard API implemented by many kernel architectures:
  - Many Unix kernels, Linux, Mac OSX, Windows NT
- Win32 is another important API



# Kernel Architectures

# Unix Kernel



# Kernel Architectures

- Kernel the core element of operating system
- Various design and implementation approaches
  - Monolithic kernels
  - Layered approach
  - Microkernel
  - Kernel modules

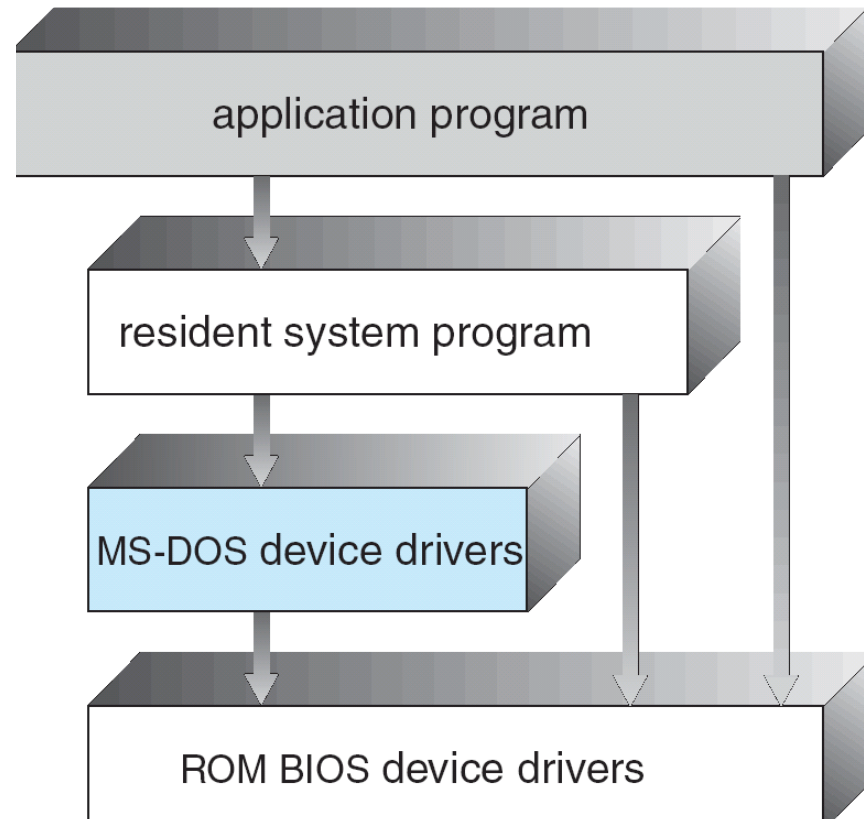


# Monolithic Kernel

- Most operating systems, until recently, featured a large monolithic kernel (most Unix systems, Linux)
- Provide
  - Scheduling
  - File system management
  - Networking
  - Device drivers
  - Memory management
  - Etc.
- Implemented as a single process
  - All functional components share same address space
- Benefit
  - Performance
- Problem
  - Vulnerability to failure in components

# Simple Monolithic Structure

- Early operating systems were monolithic
- No well defined structure
- No layering, not divided into modules
- Started as small and simple systems
- Example: MS-DOS
  - Developed to provide most functionality in the least space
  - Levels not well separated, programs can directly access I/O devices



# Introducing Layers

- Simple un-organised structures became infeasible
- Introduction of a layered approach
- Operating system is divided into a number of layers (levels), each built on top of lower layers
  - The bottom layer (layer 0) is the hardware
  - The highest layer (layer N) is the user interface
- Each layer uses only functions and services provided by a lower layer
- All or most of the layers operate in kernel mode
- Examples
  - MULTICS, VAX/VMS

# Layered Approach

- Approach used by original Unix kernel
  - Minimal layering, thick monolithic layers, no clear separation – circular dependencies, difficult to debug and extend
- Better approach - strict layering
- Difficulty
  - How to define layers appropriately?
  - Layering is only possible if there is a strict calling hierarchy among system calls and no circular dependencies
- Example
  - The TCP/IP networking stack is a strictly layered architecture

# Layered Approach - Problems

- Circular dependencies
  - Example disk device driver
    - Device driver may have to wait for I/O completion, invokes the CPU scheduling layer
    - CPU may need to call the device driver to swap processes in and out to hard disk
- The more layers the more indirections from function to function and the bigger the overhead in function calls
- Backlash against strict layering: return to fewer layers with more functionality

# Microkernel

- A microkernel is a reduced operating system core that contains only essential OS functions
- Idea: minimise kernel by executing as much functionality as possible in user mode
  - Run them as conventional user processes
  - Processes interact only via message passing (IPC)
- Many services are now external processes
  - Device drivers
  - File systems
  - Virtual memory manager
  - Windowing systems
  - Security services etc.
- Example: Mach operating system

# Mach Kernel

- Developed at Carnegie Mellon University 1985
- Research kernel
- Various versions of it were developed further
  - Microkernel as well as non-microkernel versions
- Notably
  - NeXTSTEP / Mac OSX, FreeBSD (not a microkernel, but provides microkernel IPC to applications)
- Problem of Mach kernel: IPC (Inter-Process Communication) overhead

# Microkernel System Structure

- Operating system components external to the microkernel are implemented as server processes
  - These processes interact via message passing (IPC)
- Microkernel facilitates the message exchange
  - Validates messages
  - Passes messages between components
  - Checks whether message passing is permitted
- Grants access to hardware
- Microkernel effectively implements a client-server infrastructure on a single computer



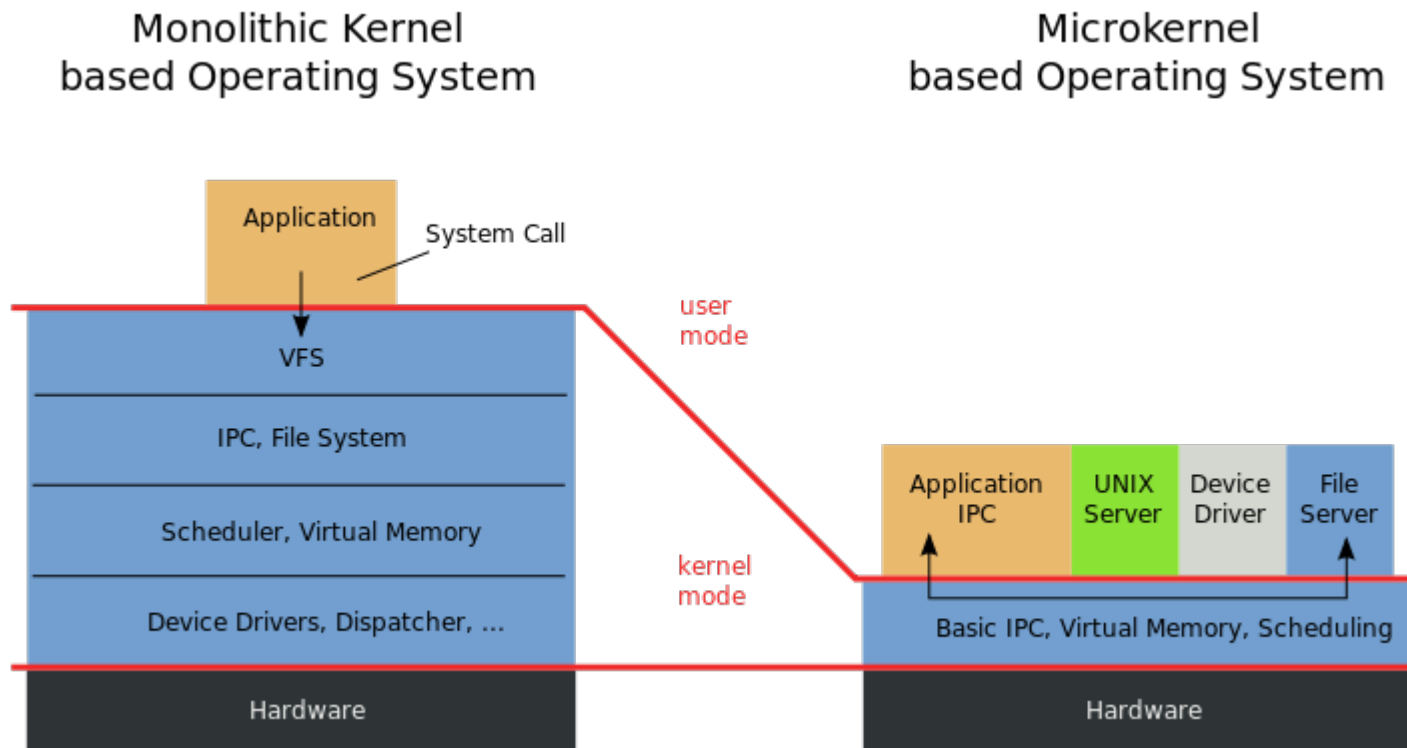
# Microkernel

- Benefits
  - Uniform interfaces
    - Processes pass messages, no distinction between user-mode and kernel-mode services, all services are provided via message passing as in a client-server infrastructure
  - Extendibility
    - Easier to extend, new services introduced as new applications
  - Portability
    - Only the microkernel has to be adapted to a new hardware
  - Reliability and security
    - much less code runs in kernel mode, program failures occurring in user mode execution does not affect the rest of the system

# Microkernel Design

- Minimal functionality that has to be included into a microkernel
  - Low-level memory management
    - Mapping of memory pages to physical memory locations
    - All other mechanisms of memory management are provided by services running in user mode
      - Address space protection
      - Page replacement algorithms
      - Virtual memory management
  - Interprocess communication (IPC)
  - I/O and interrupt management

# System Call Monolithic vs Micro Kernel



# Microkernel

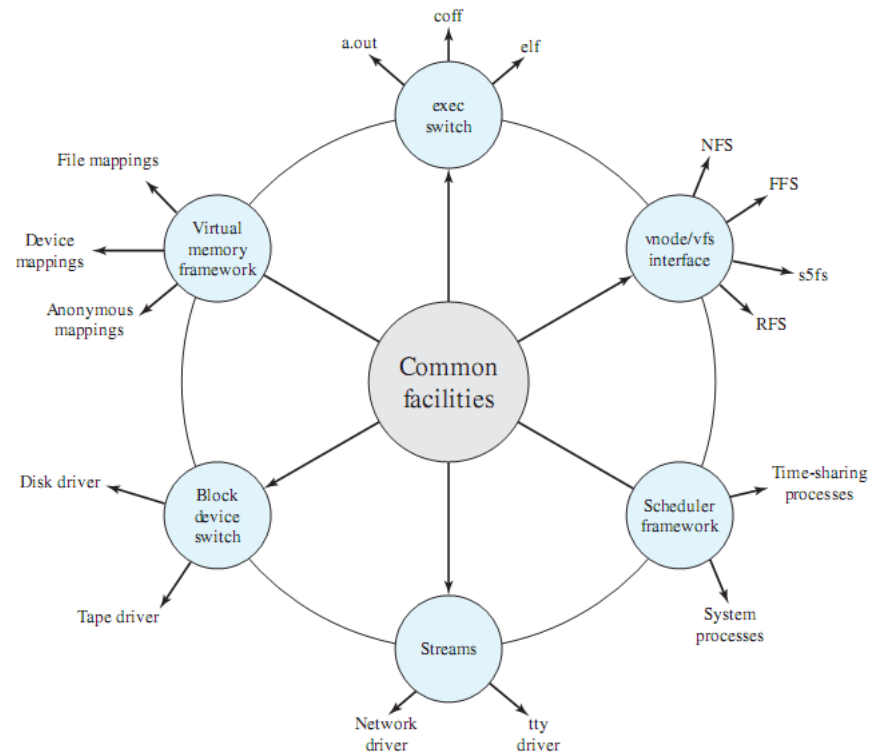
- Problems
  - Performance overhead of communication between system services
    - Each interaction involves the kernel and a user mode / kernel mode switch
    - System services running in user mode are processes, operating system has to switch between them
  - Solution: reintegration of services running in user mode back into the kernel
    - Improves performance: less mode switches, services integrated in kernel share one address space (one process)
    - This was done with the Mach kernel
  - Solution: make kernel even smaller – experimental kernel architectures (Nano kernels, pico kernels)

# Modular Kernel Design

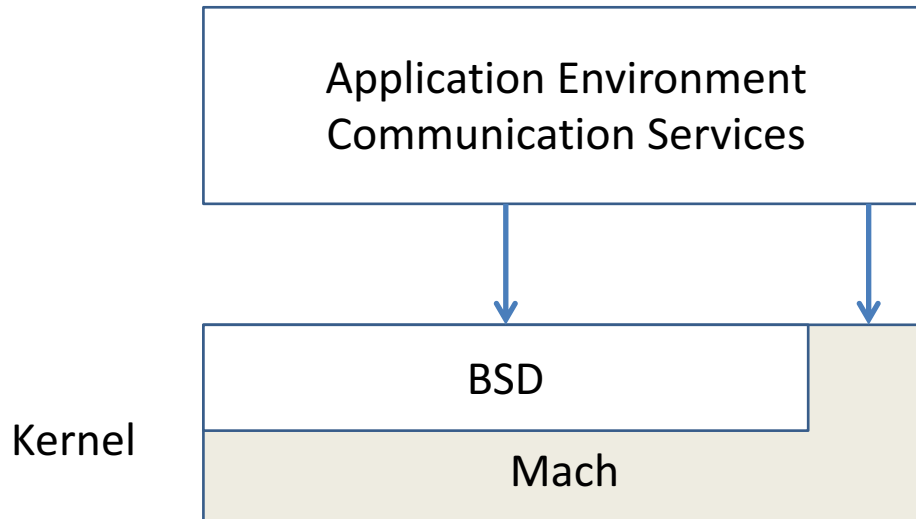
- Many operating systems implement kernel modules
  - E.g.: Linux
    - Each core component is separate
    - Communication via defined interfaces
    - Loadable on demand
- Modules are somehow a hybrid between the layered and microkernel approach
  - Clean software engineering approach
  - But: modules are inside the kernel space, they don't require the overhead of message passing
  - Compromise with performance benefits

# Modern Unix Kernel

- Modern Unix kernels have a modular architecture
- Common facilities as the inner core of the kernel
- Rest of system services added as modules
- See Linux



# Modular Approach



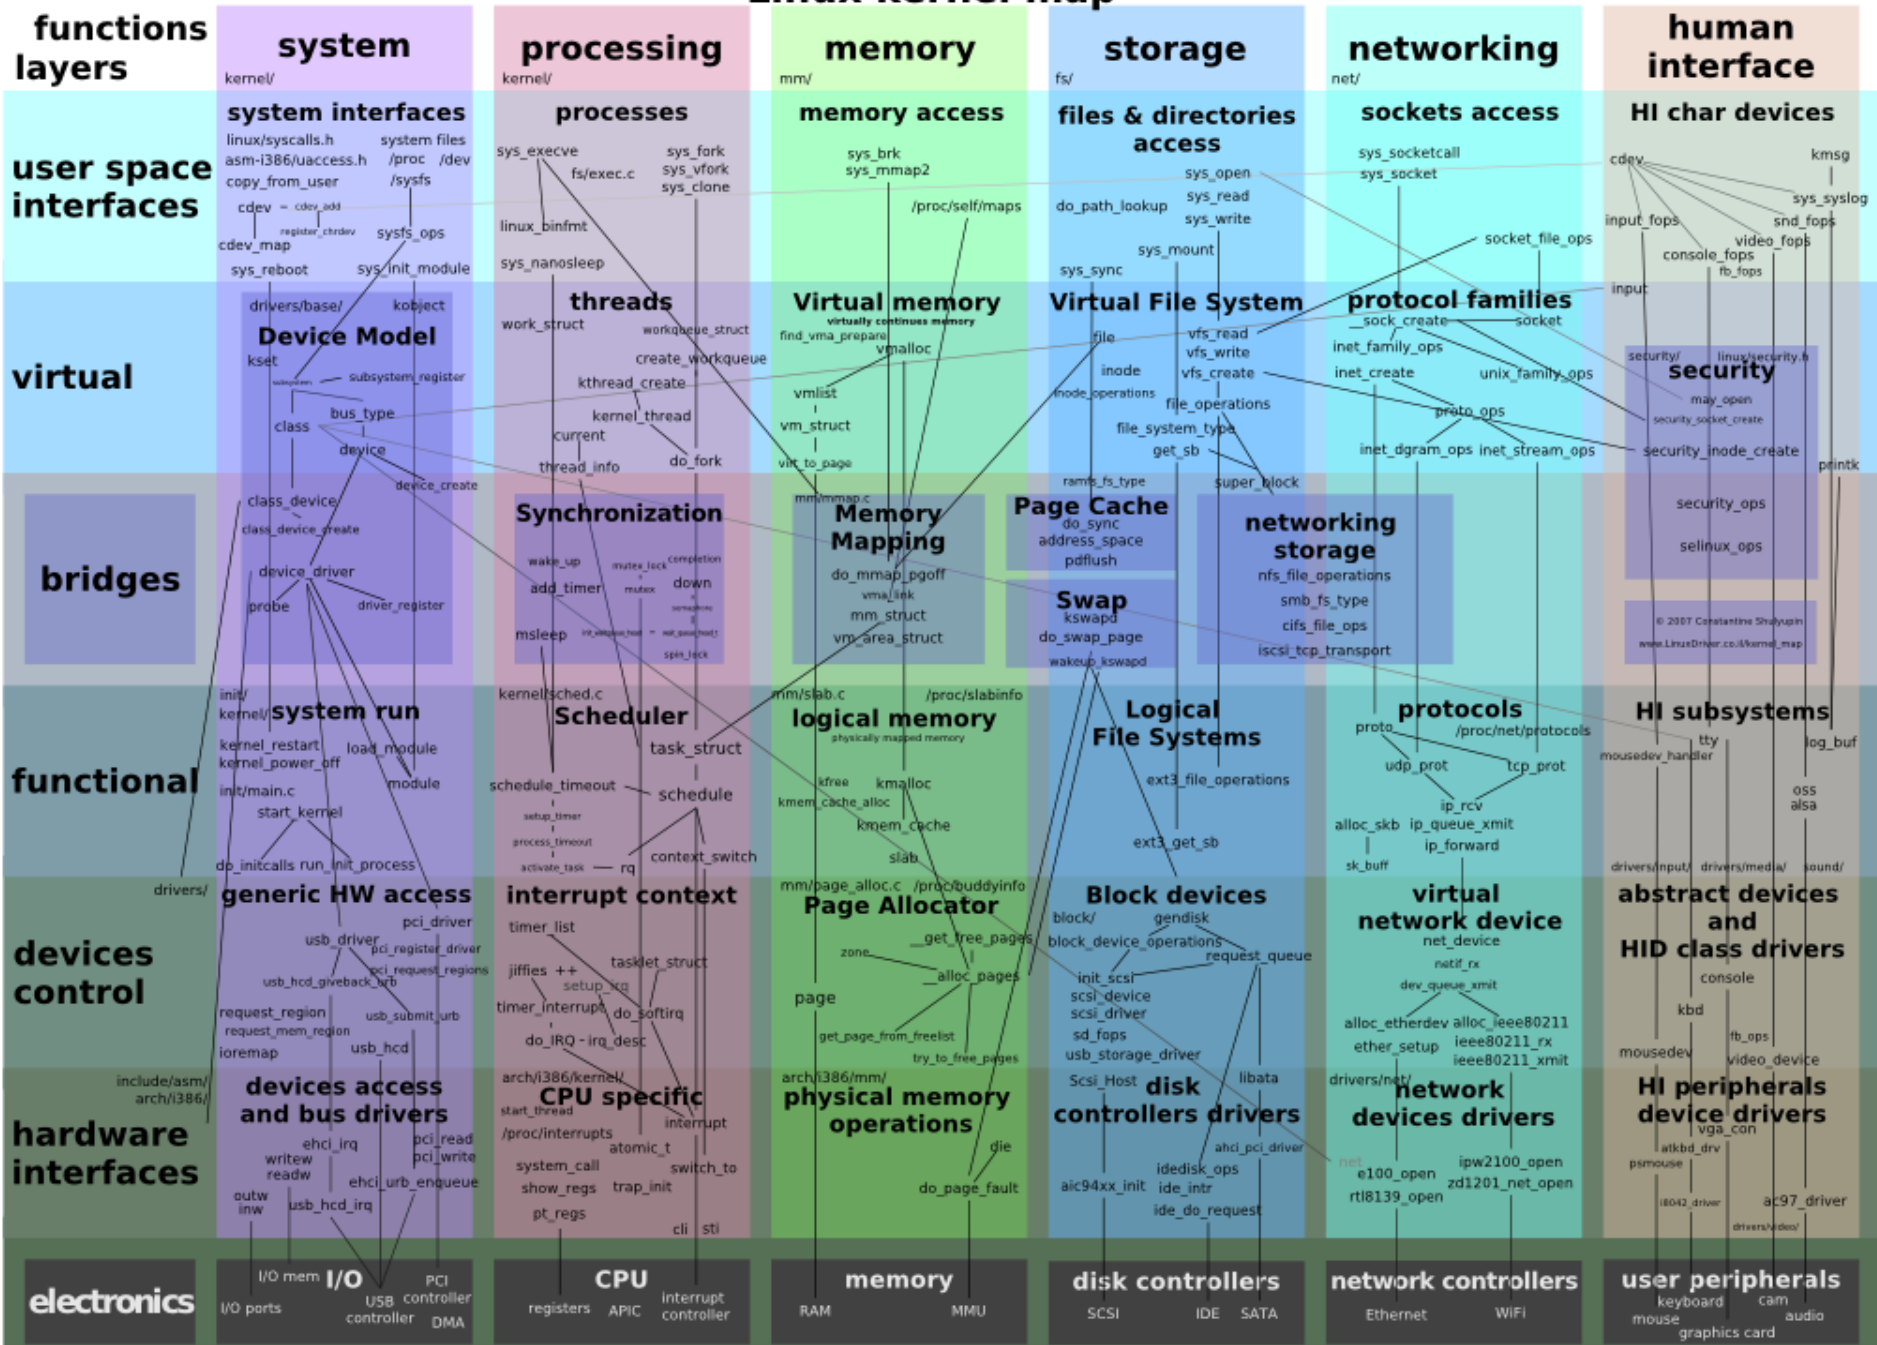
- Mac OSX
  - takes a hybrid approach, has a Mach kernel (microkernel) combined with a BSD interface
  - BSD (**Berkeley Software Distribution**, sometimes called **Berkeley Unix**): provides support for command line interface, networking, file system, file system, POSIX API and threads
  - Mach: memory management, Remote procedure Call (RPC), Interprocess communication (IPC), message passing

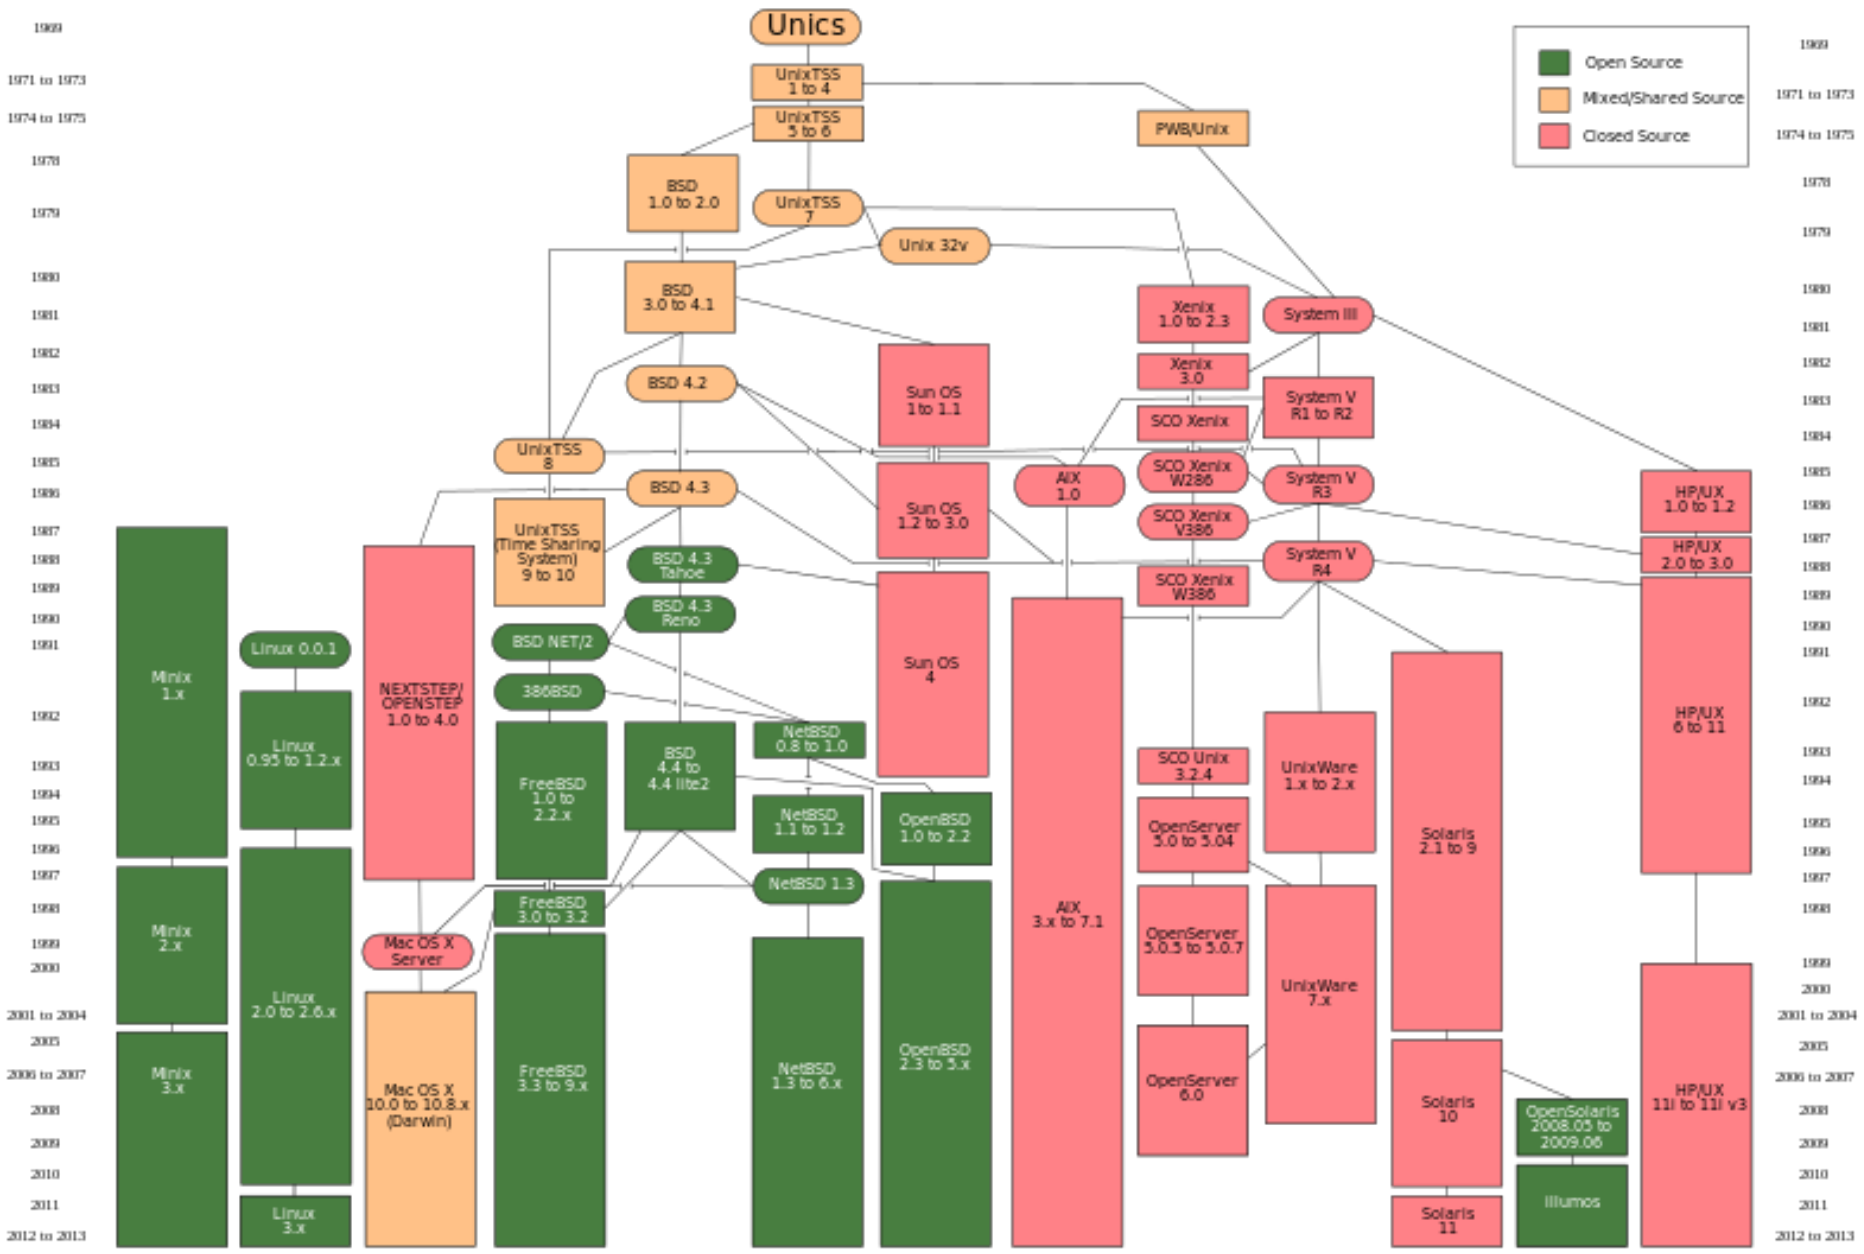
# OS Functionality

- Linux Kernel Map
  - [http://www.makelinux.net/kernel\\_map/](http://www.makelinux.net/kernel_map/)
- Also available at:
  - [http://upload.wikimedia.org/wikipedia/commons/5/5b/Linux\\_kernel\\_map.png](http://upload.wikimedia.org/wikipedia/commons/5/5b/Linux_kernel_map.png)
  - <http://i.imgur.com/4sftcoo.jpg>



# Linux kernel map





[http://en.wikipedia.org/wiki/File:Unix\\_history-simple.svg](http://en.wikipedia.org/wiki/File:Unix_history-simple.svg)